

Implementation of a multithreaded compile-time ECS in C++14



<http://vittorioromeo.info>

vittorio.romeo@outlook.com



Implementation of a multithreaded compile-time ECS in C++14



<http://vittorioromeo.info>

vittorio.romeo@outlook.com



<http://github.com/SuperV1234/cppnow2016>

Overview of the talk

- “Entity-component-system” pattern in general.
- Overview of **ECST**.
 - Design and core values.
 - Features/limitations.
- Complete usage example.
- Architecture of **ECST**.
- Implementation of **ECST**.
- Future ideas.

Entity-component-system

Concepts, benefits and drawbacks of the ECS architectural pattern.



What is an entity?

- Something **tied to a concept**.
- Has related **data** and/or **logic**.
- We deal with **many entities**.
 - We may want to track specific instances.
- Can be **created** and **destroyed**.
- Examples:
 - **Game objects**: *player, bullet, car*.
 - **GUI widgets**: *window, textbox, button*.



Encoding entities – OOP inheritance

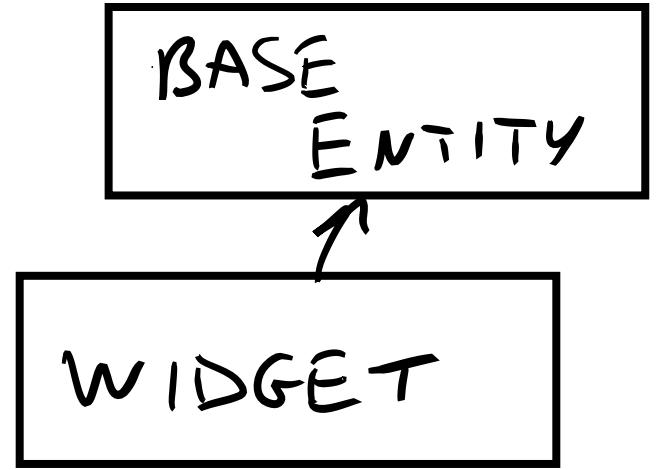
- An **entity type** is a **polymorphic class**.
- **Data** is stored inside the class.
- **Logic** is handled using **runtime polymorphism**.
- **Very easy to implement**.
- **Cache-unfriendly**.
- **Runtime overhead**.
- **Lack of flexibility**.

BASE
ENTITY



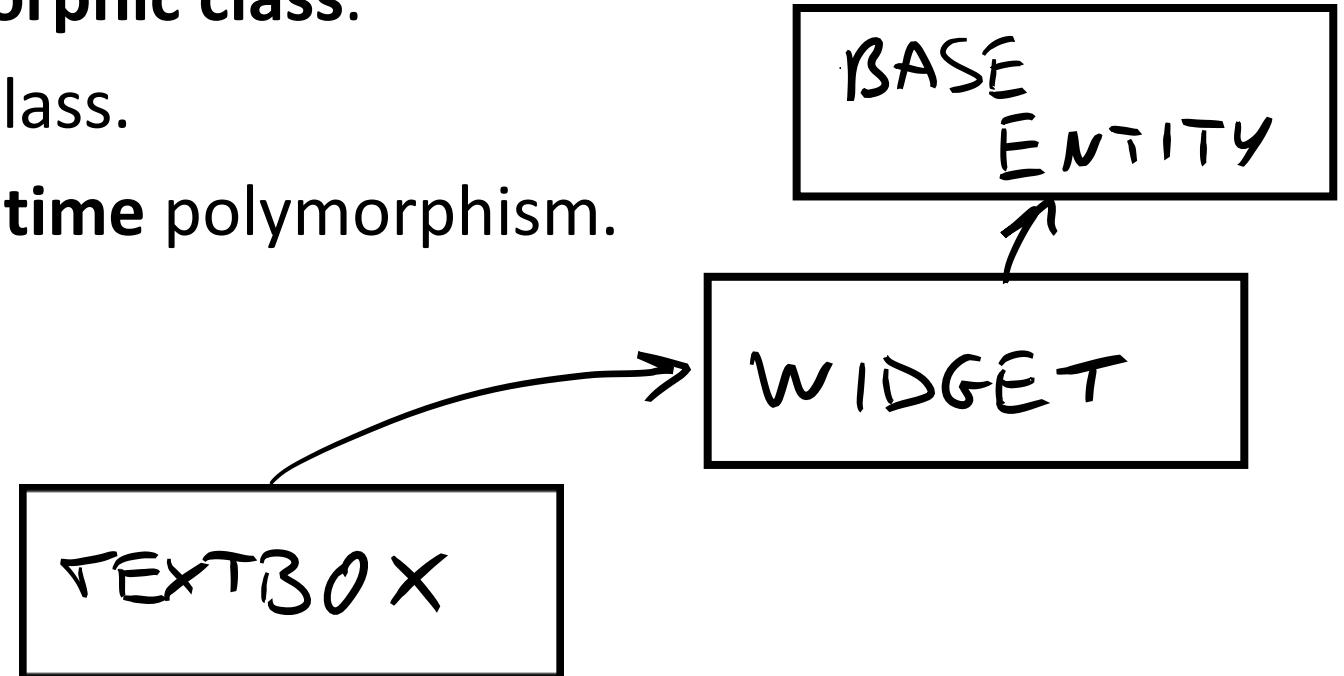
Encoding entities – OOP inheritance

- An **entity type** is a **polymorphic class**.
- **Data** is stored inside the class.
- **Logic** is handled using **runtime polymorphism**.
- **Very easy to implement**.
- **Cache-unfriendly**.
- **Runtime overhead**.
- **Lack of flexibility**.



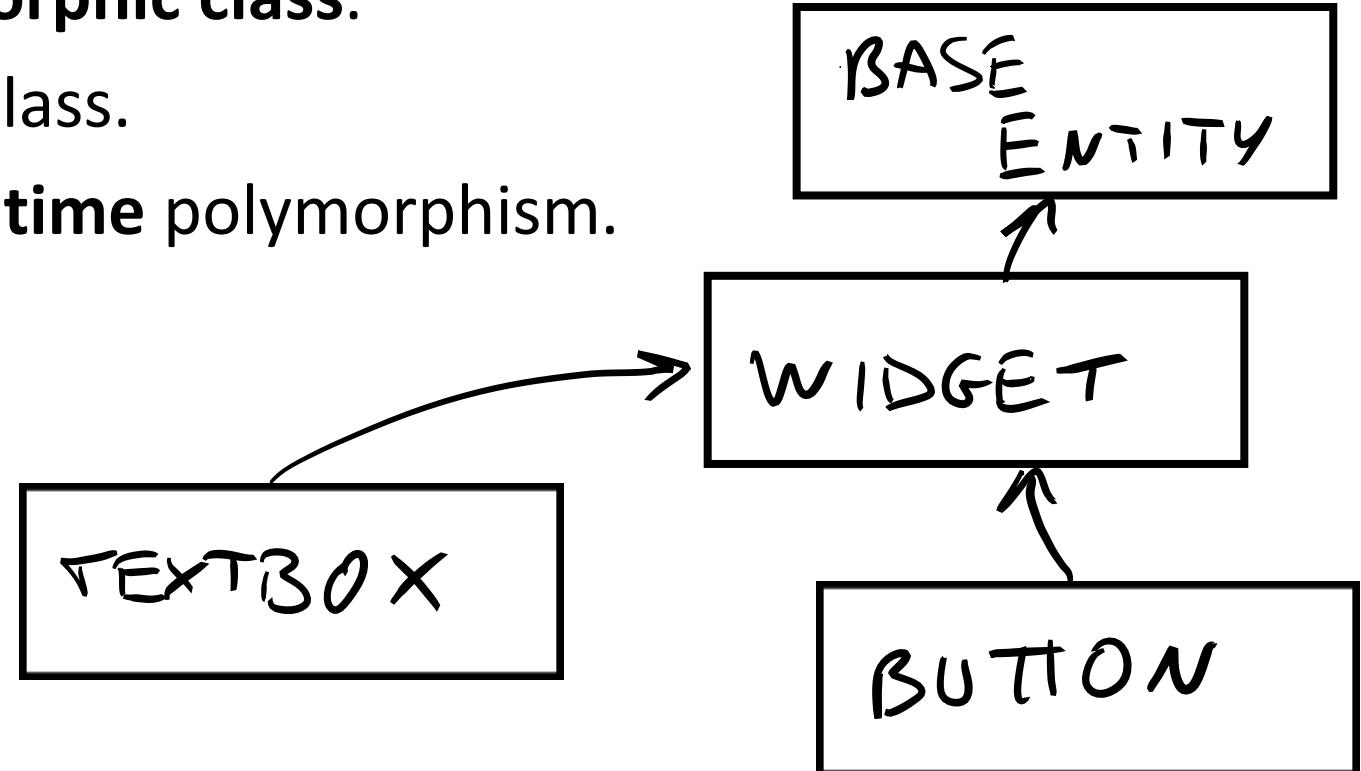
Encoding entities – OOP inheritance

- An **entity type** is a **polymorphic class**.
- **Data** is stored inside the class.
- **Logic** is handled using **runtime polymorphism**.
- **Very easy to implement.**
- **Cache-unfriendly.**
- **Runtime overhead.**
- **Lack of flexibility.**

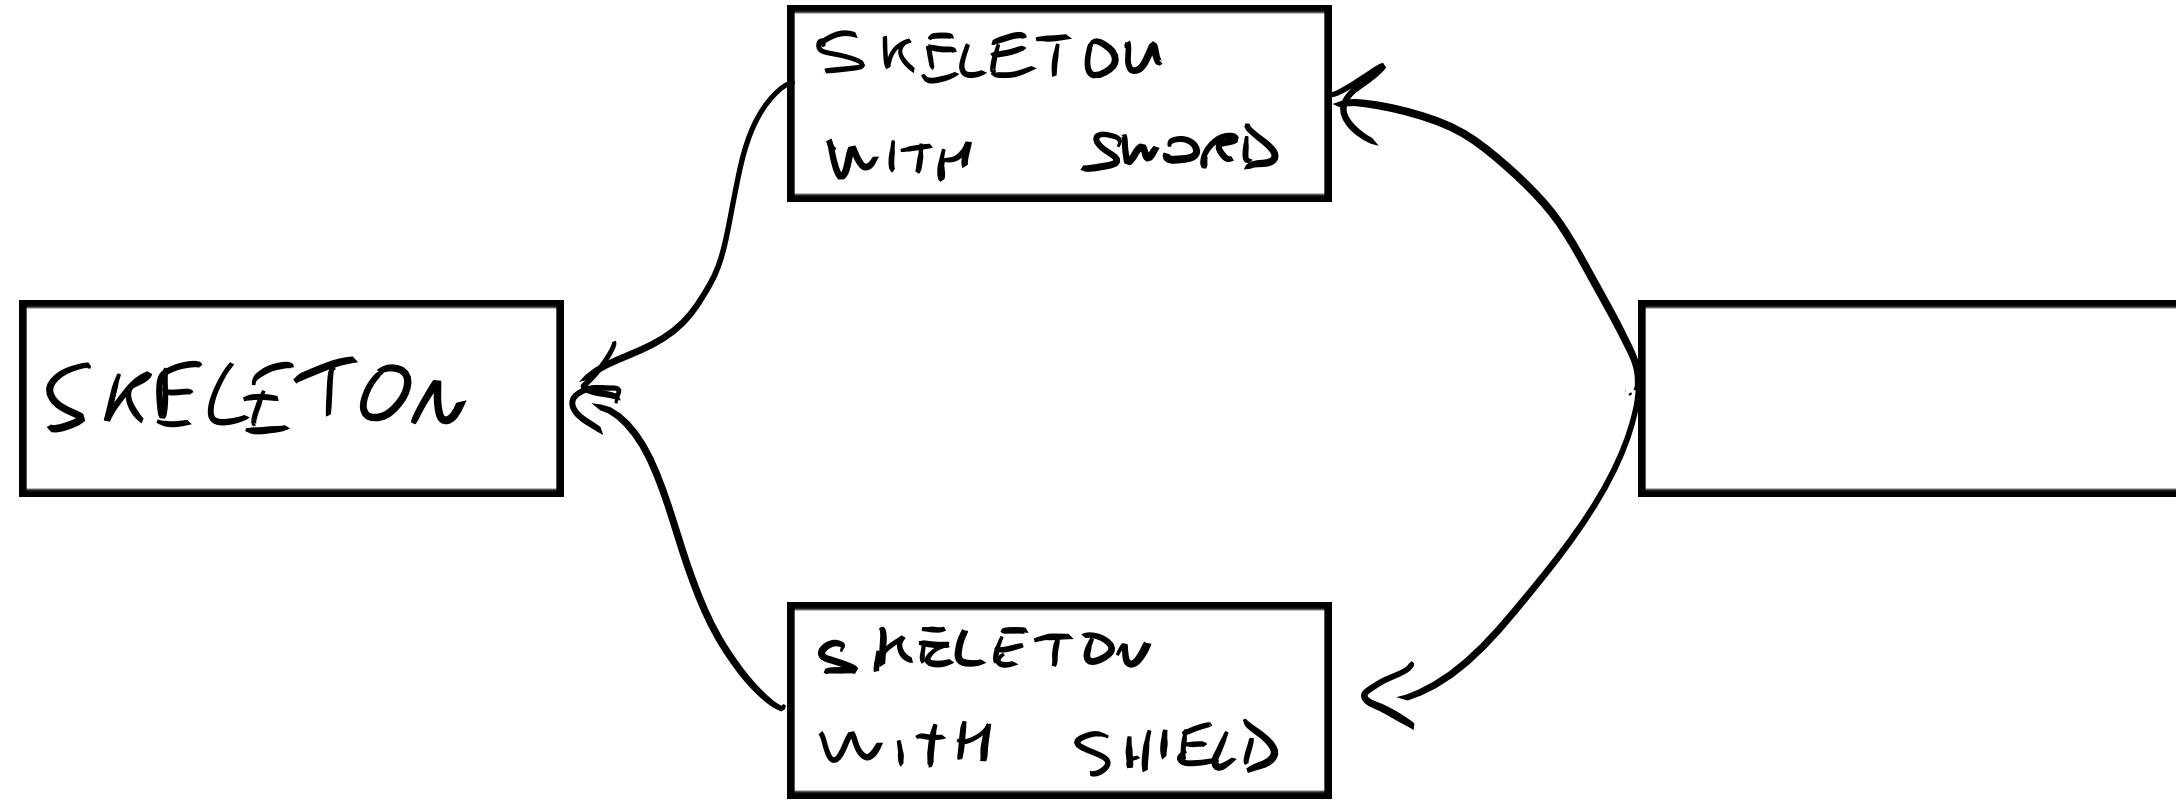


Encoding entities – OOP inheritance

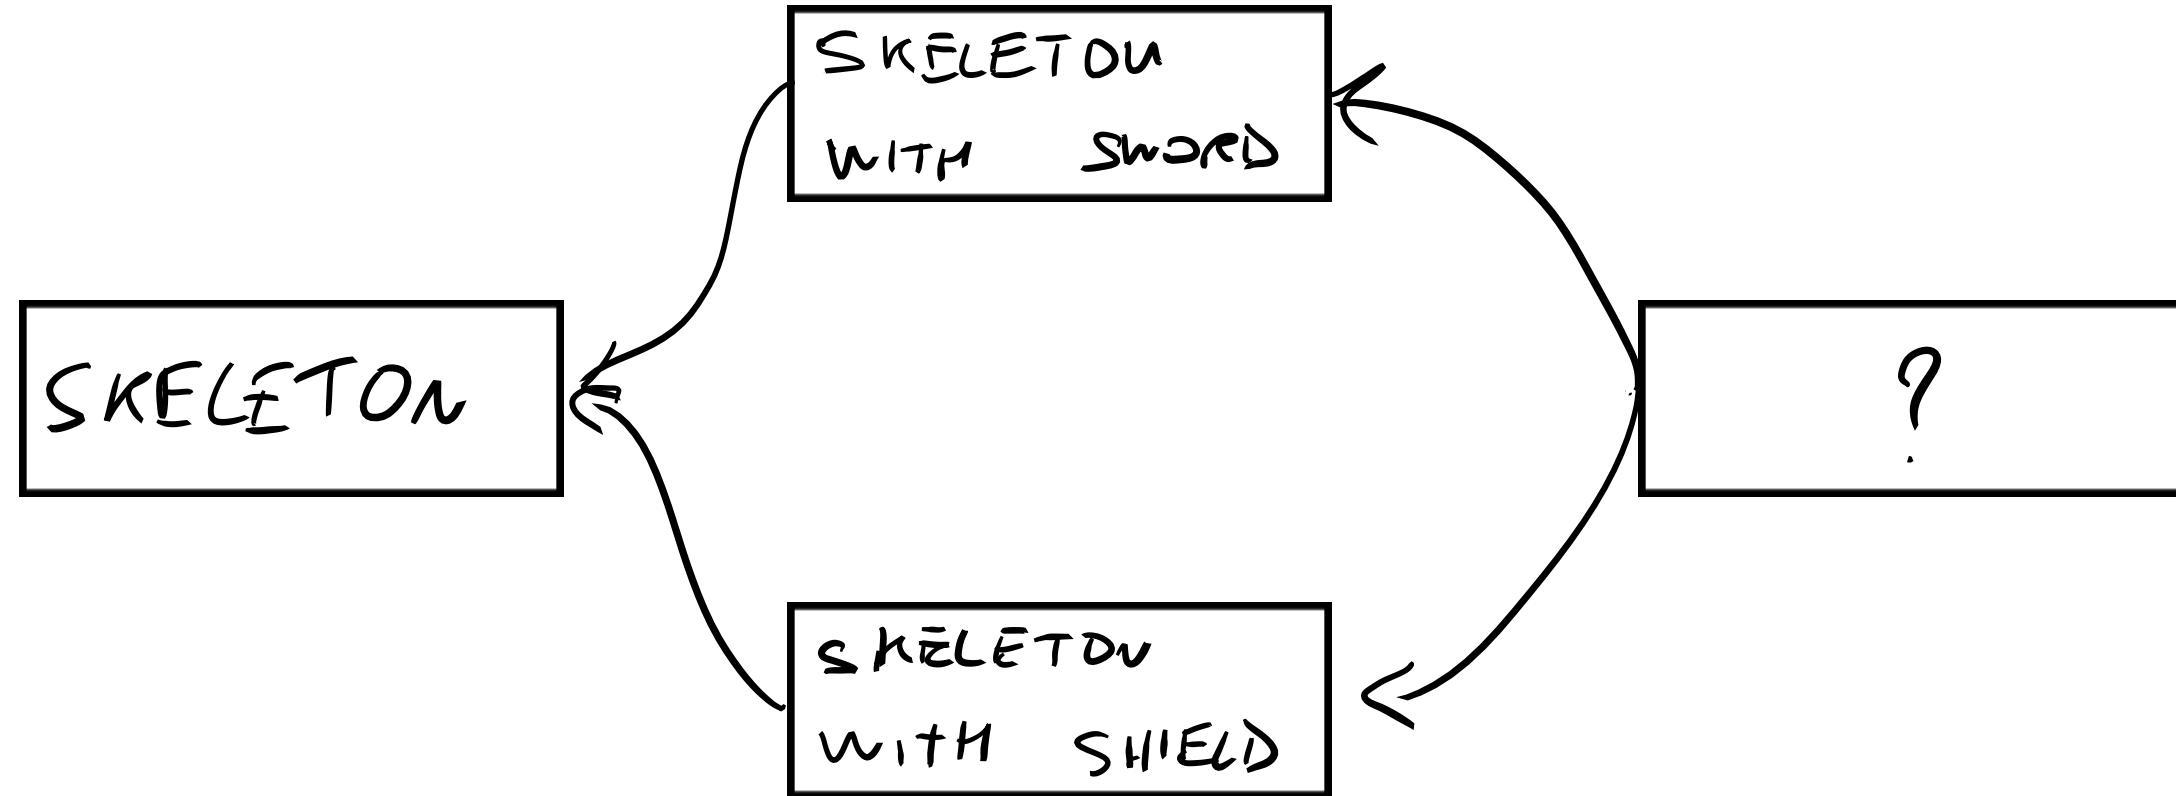
- An **entity type** is a **polymorphic class**.
- **Data** is stored inside the class.
- **Logic** is handled using **runtime polymorphism**.
- **Very easy to implement.**
- **Cache-unfriendly.**
- **Runtime overhead.**
- **Lack of flexibility.**



Encoding entities – OOP inheritance



Encoding entities – OOP inheritance



Encoding entities – OOP inheritance

```
struct Entity
{
    virtual ~Entity() { }
    virtual void update() { }
    virtual void draw() { }
};

struct Skeleton : Entity
{
    std::vector<Bone> bones;
    void update() override
    {
        // do things skeletons do
    }
};
```

Encoding entities – OOP inheritance

```
struct Entity
{
    virtual ~Entity() { }
    virtual void update() { }
    virtual void draw() { }
};

struct Skeleton : Entity
{
    std::vector<Bone> bones;
    void update() override
    {
        // do things skeletons do
    }
};
```

Encoding entities – OOP composition

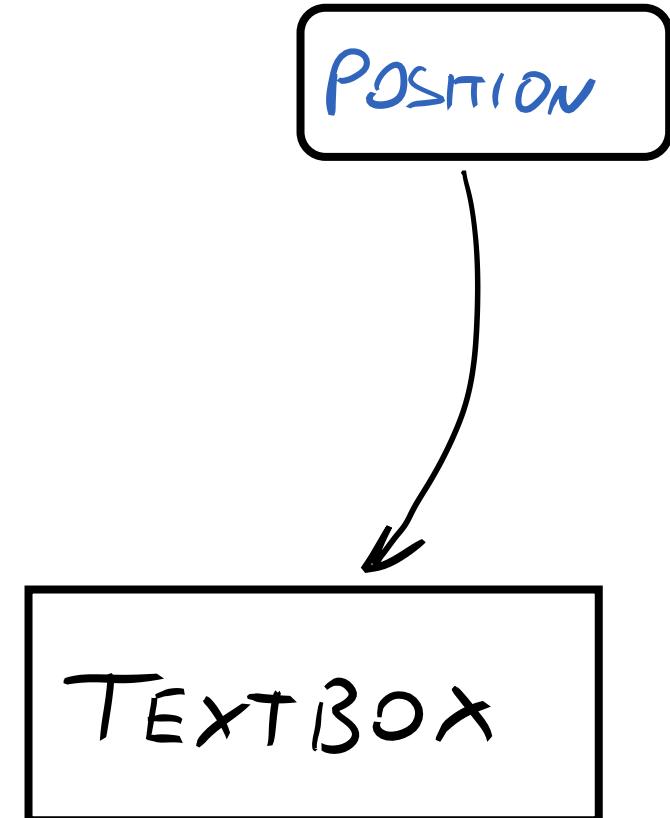
- An **entity** is an aggregate of **components**.
- Components **store data** and have **logic**.
- **Logic** is handled using **runtime polymorphism**.
- **Easy to implement.**
- **More flexible.**
- **Cache-unfriendly.**
- **Runtime overhead.**

TEXTBOX



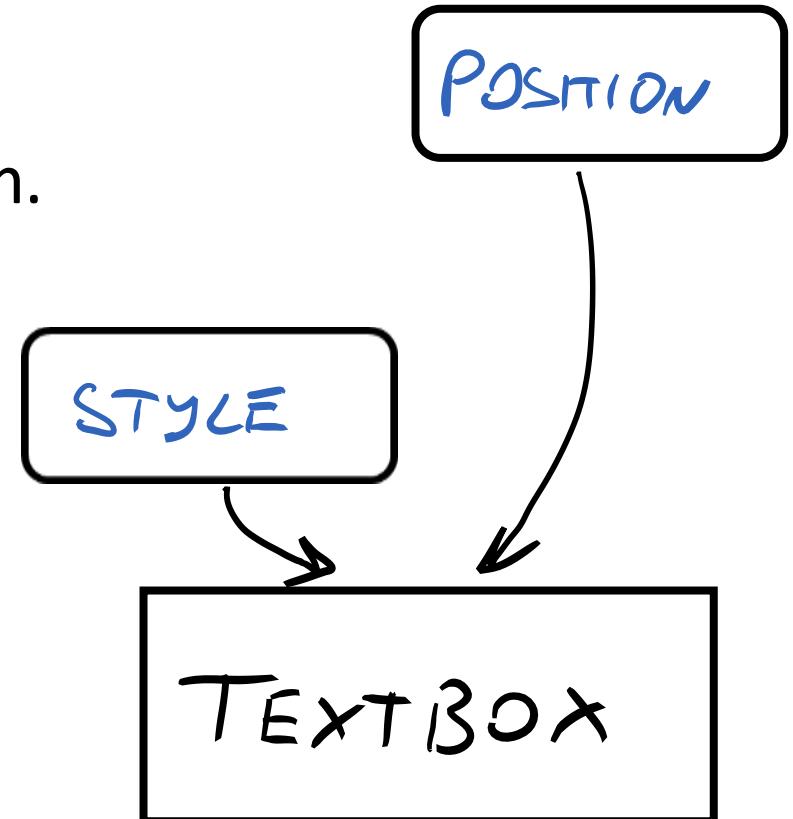
Encoding entities – OOP composition

- An **entity** is an aggregate of **components**.
- Components **store data** and have **logic**.
- **Logic** is handled using **runtime polymorphism**.
- **Easy to implement.**
- **More flexible.**
- **Cache-unfriendly.**
- **Runtime overhead.**



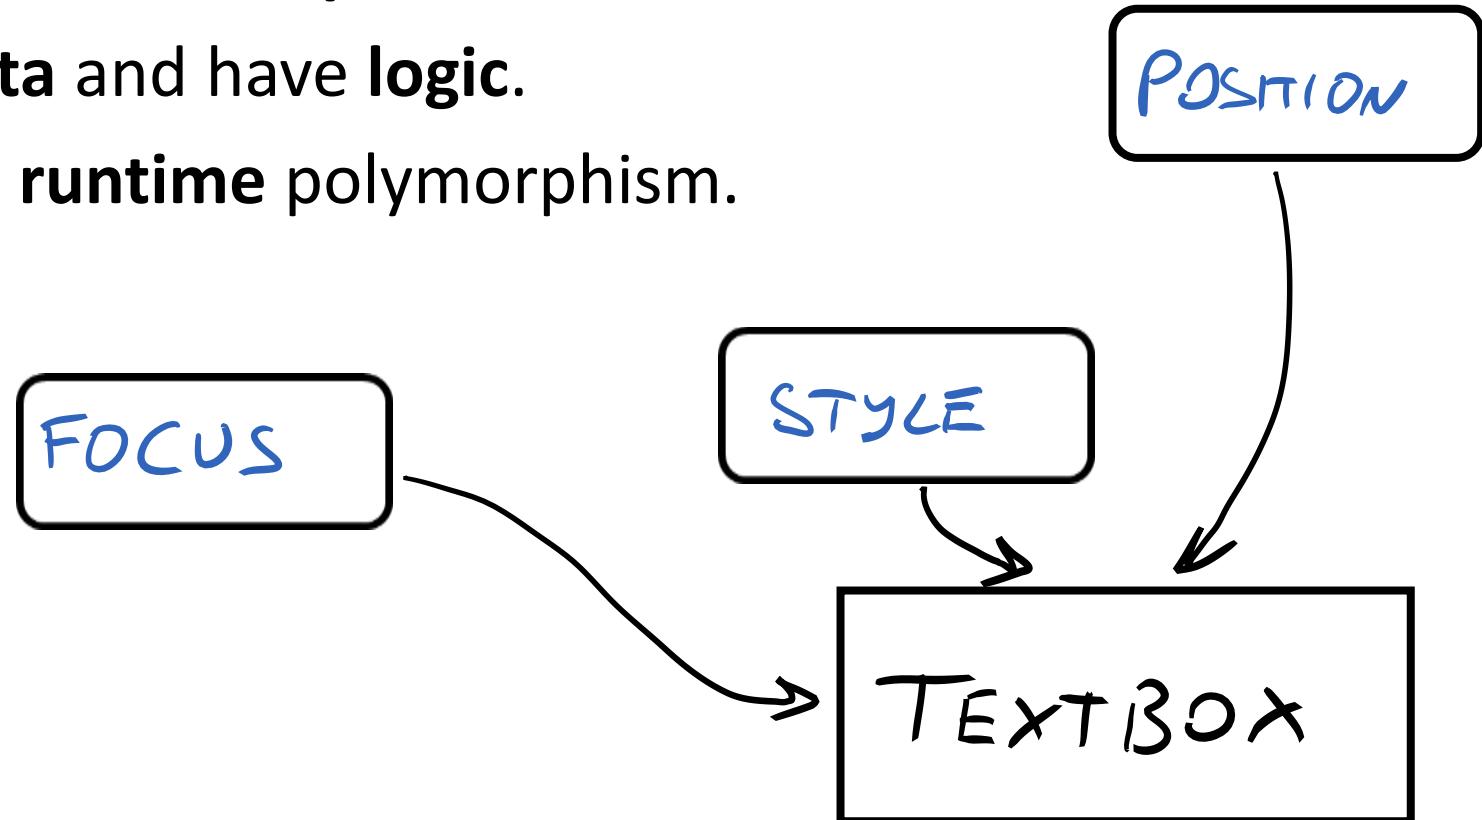
Encoding entities – OOP composition

- An **entity** is an aggregate of **components**.
- Components **store data** and have **logic**.
- **Logic** is handled using **runtime polymorphism**.
- **Easy to implement.**
- **More flexible.**
- **Cache-unfriendly.**
- **Runtime overhead.**



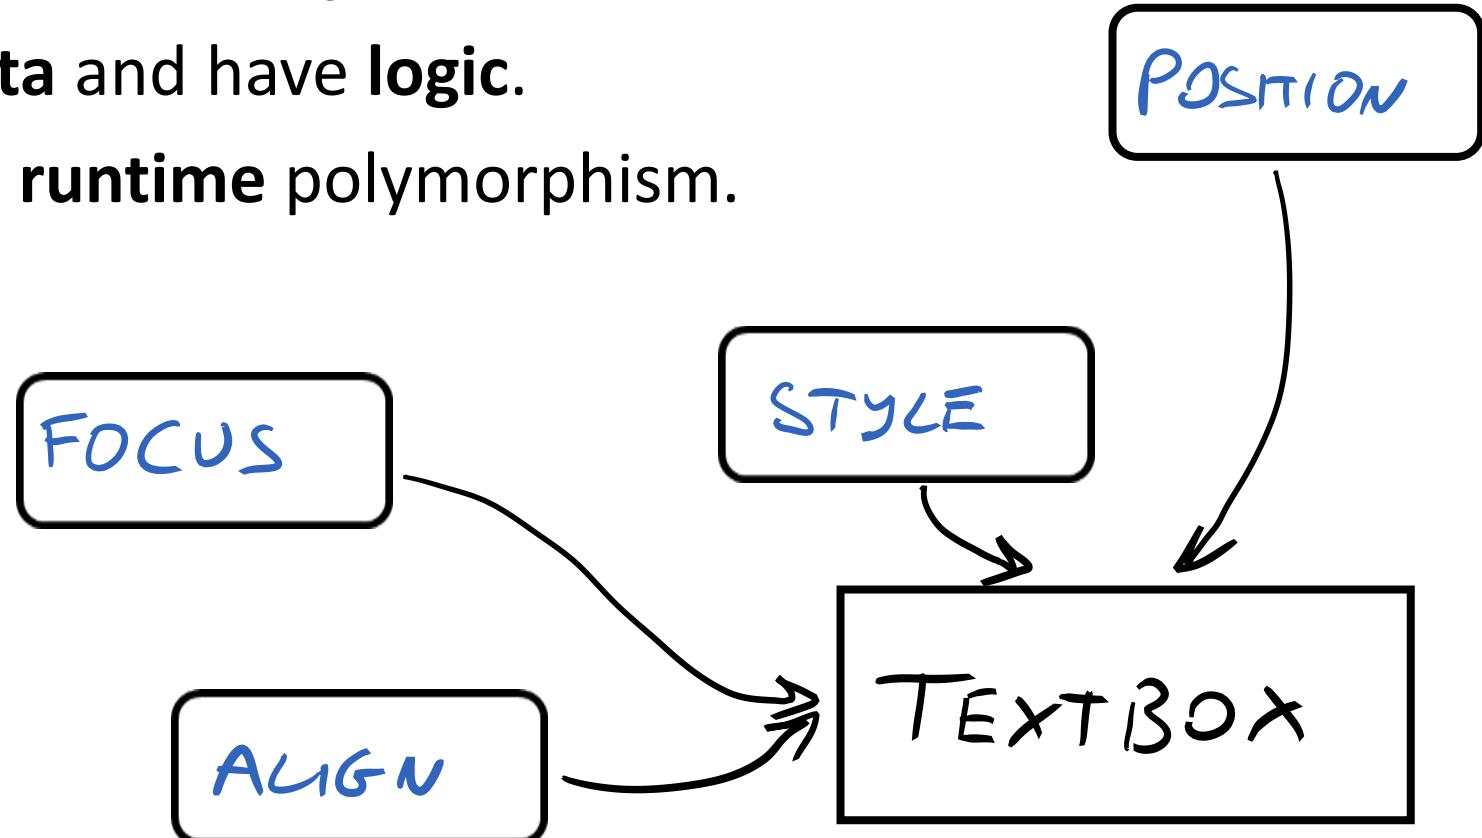
Encoding entities – OOP composition

- An **entity** is an aggregate of **components**.
- Components **store data** and have **logic**.
- **Logic** is handled using **runtime polymorphism**.
- **Easy to implement.**
- **More flexible.**
- **Cache-unfriendly.**
- **Runtime overhead.**



Encoding entities – OOP composition

- An **entity** is an aggregate of **components**.
- Components **store data** and have **logic**.
- **Logic** is handled using **runtime polymorphism**.
- **Easy to implement.**
- **More flexible.**
- **Cache-unfriendly.**
- **Runtime overhead.**



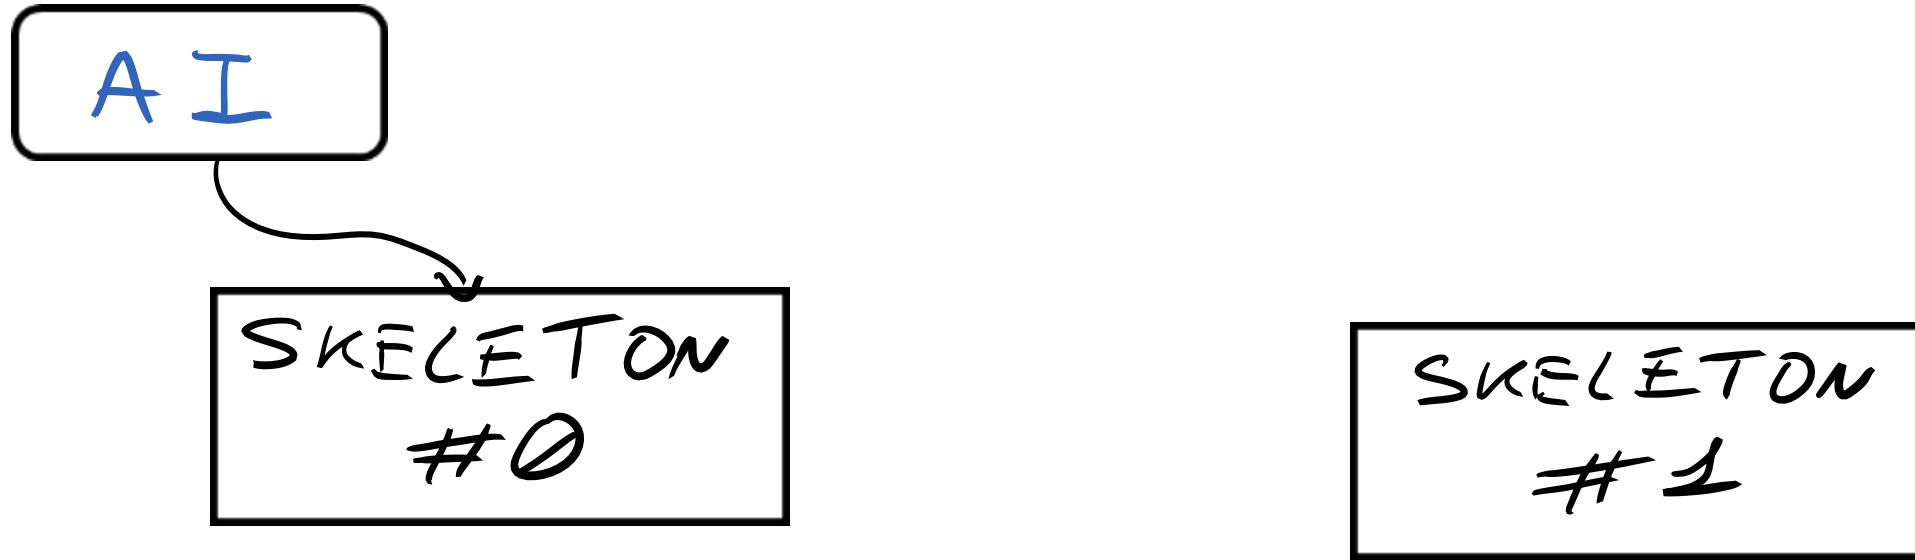
Encoding entities – OOP composition

SKELETON
#0

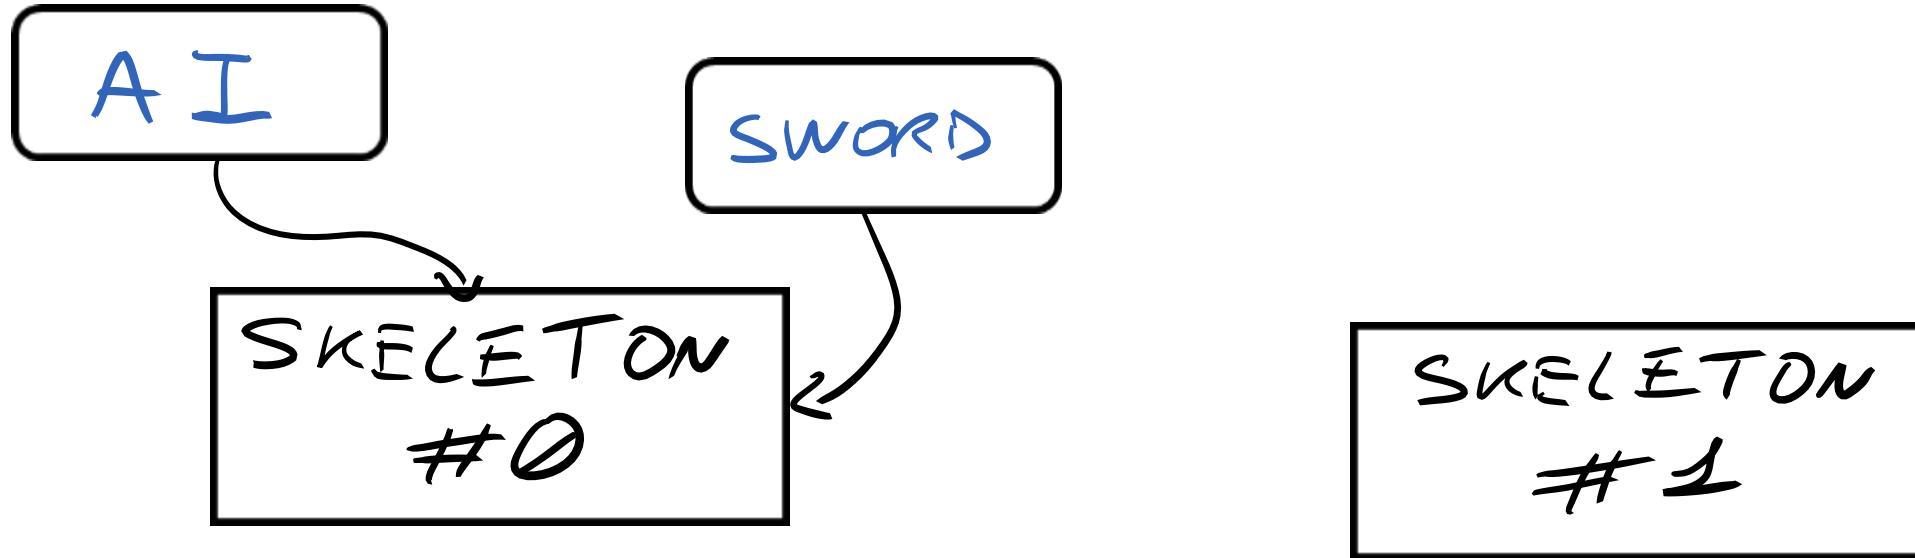
SKELETON
#1



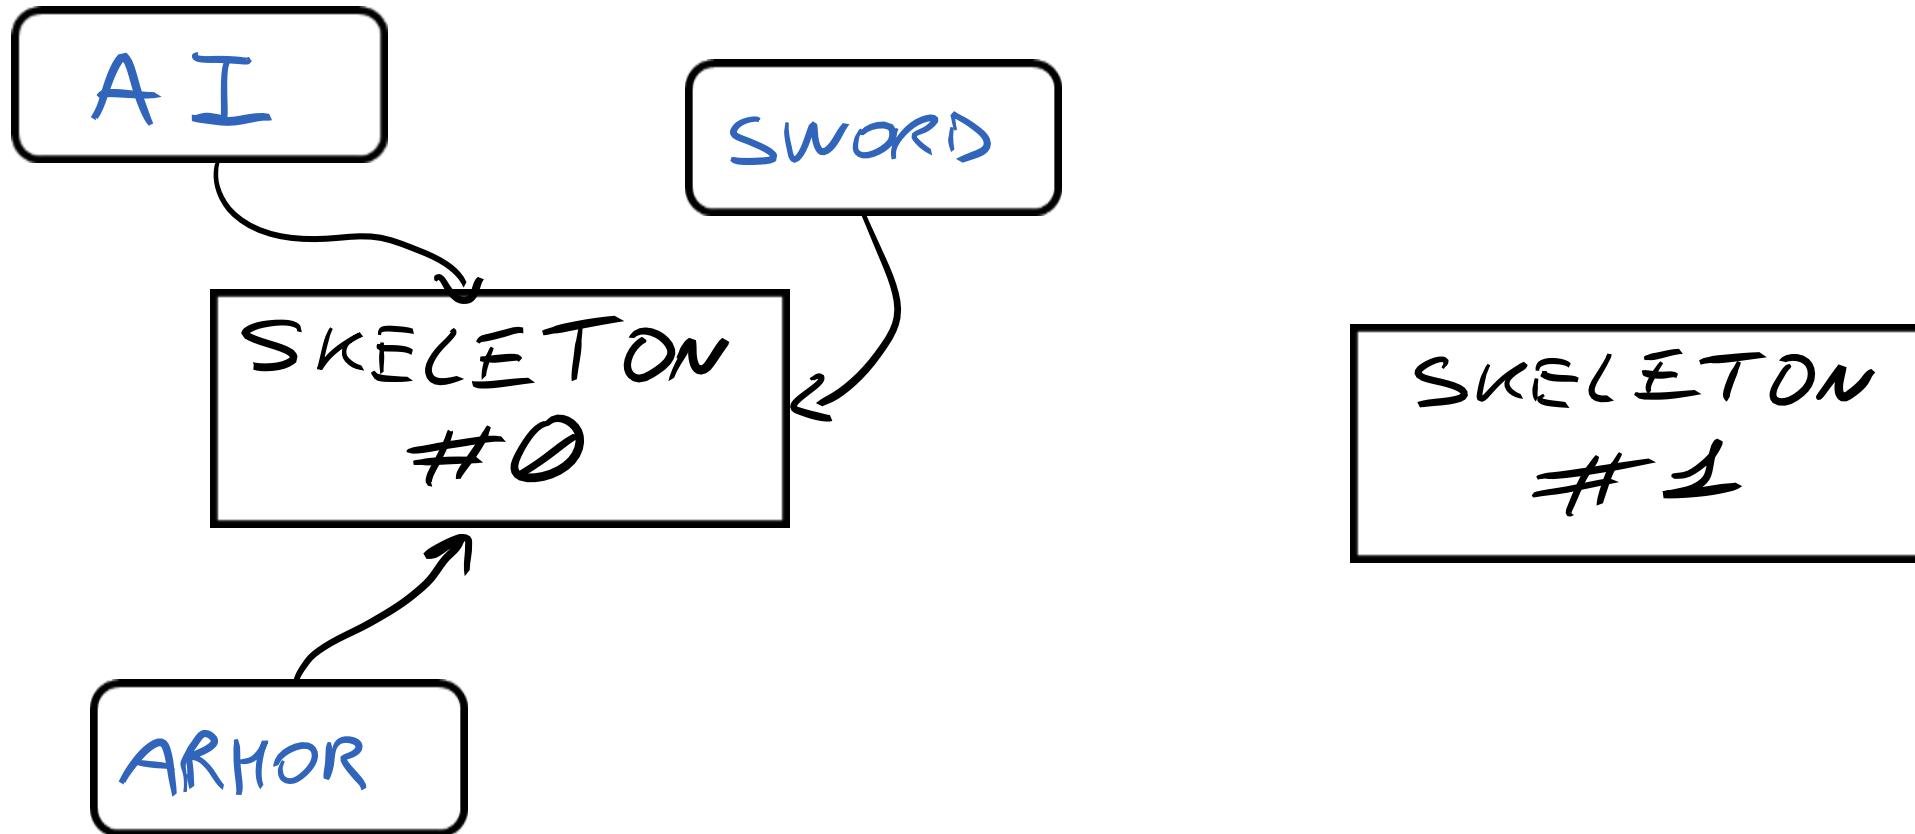
Encoding entities – OOP composition



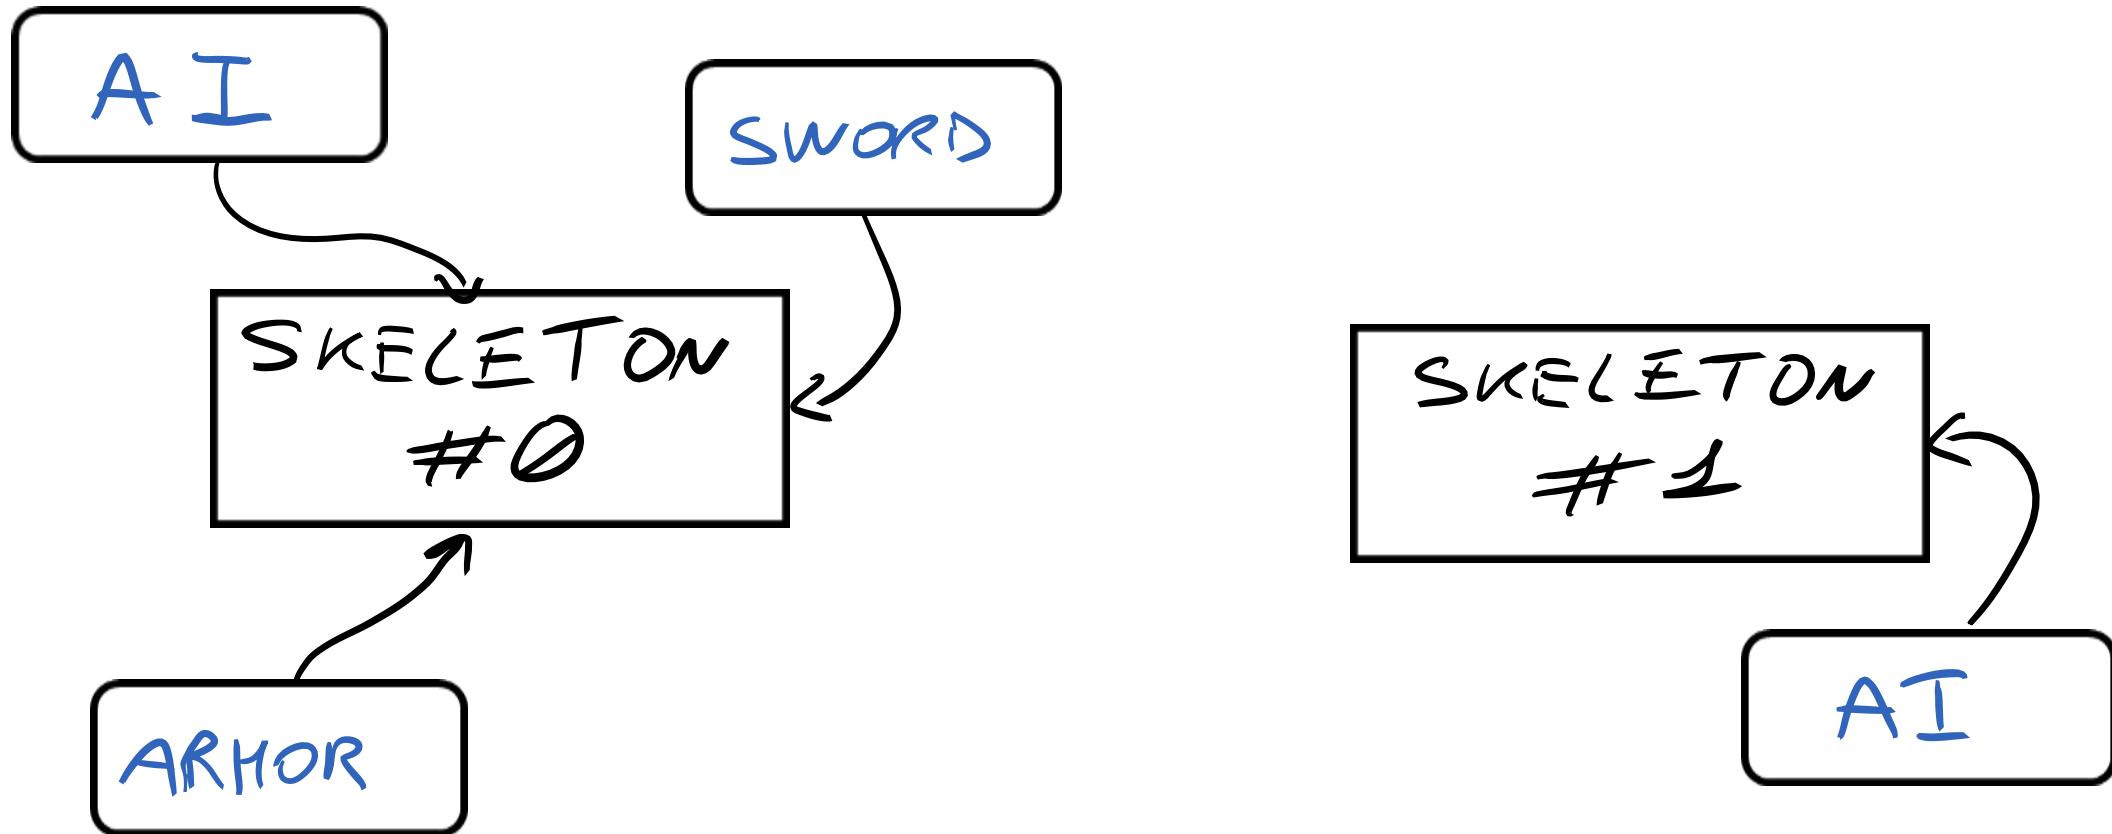
Encoding entities – OOP composition



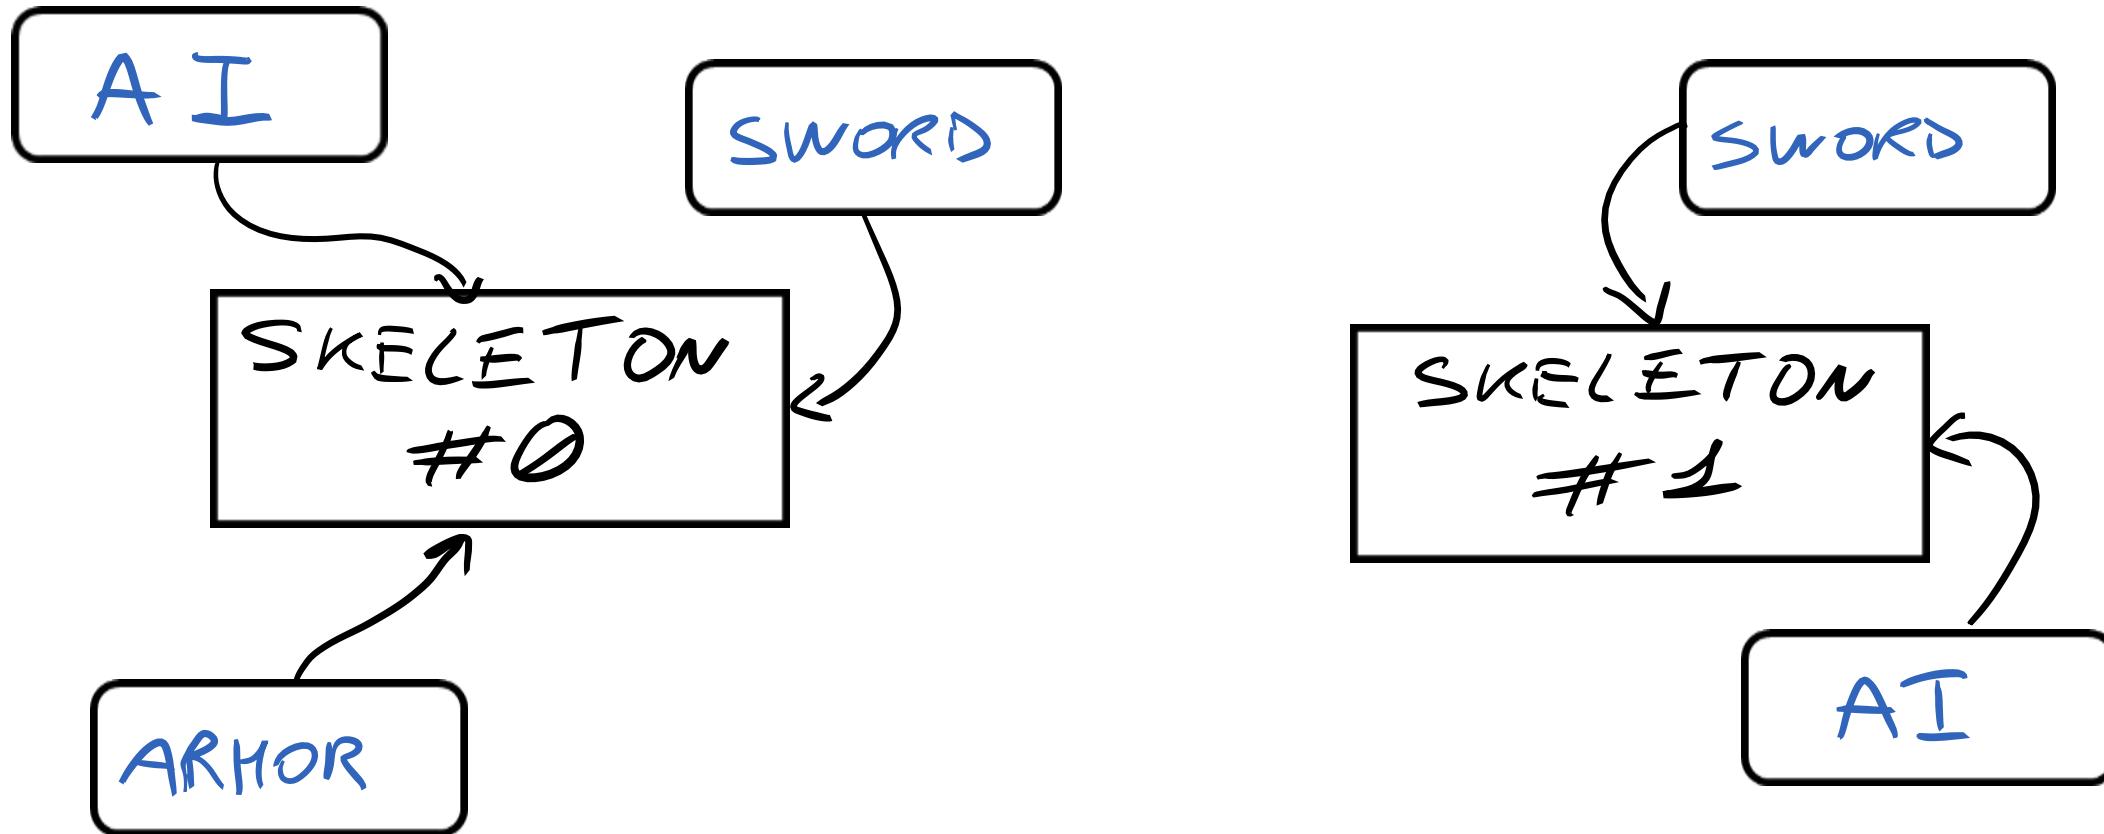
Encoding entities – OOP composition



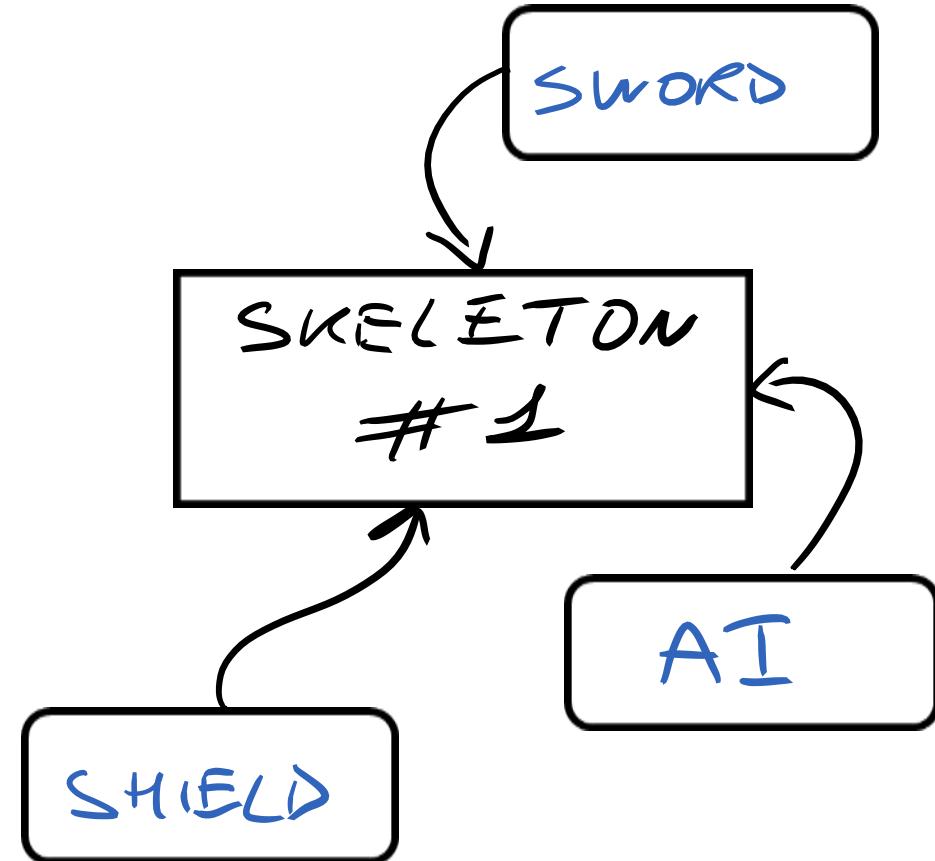
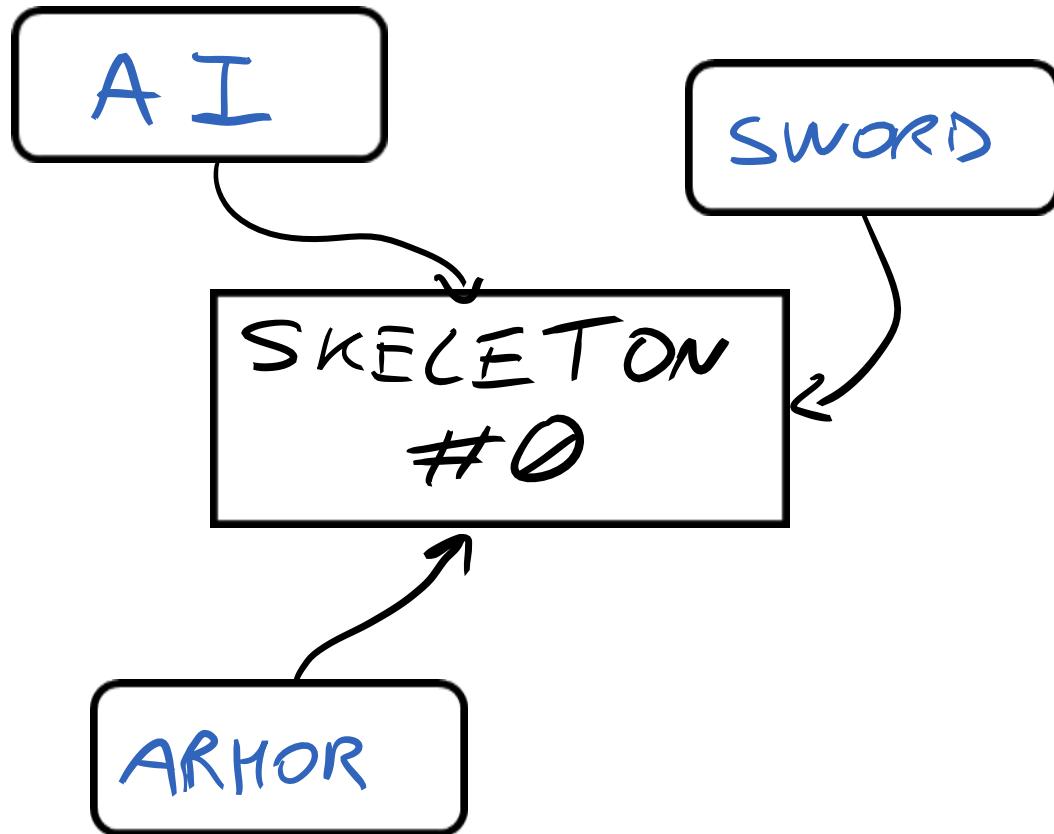
Encoding entities – OOP composition



Encoding entities – OOP composition



Encoding entities – OOP composition



Encoding entities – OOP composition

```
struct Component
{
    virtual ~Component() { }
    virtual void update() { }
    virtual void draw() { }
};

struct Entity
{
    std::vector<std::unique_ptr<Component>> components;

    void update() { for(auto& c : components) c->update(); }
    void draw() { for(auto& c : components) c->draw(); }
};
```

Encoding entities – OOP composition

```
struct Component
{
    virtual ~Component() { }
    virtual void update() { }
    virtual void draw() { }
};

struct Entity
{
    std::vector<std::unique_ptr<Component>> components;

    void update() { for(auto& c : components) c->update(); }
    void draw() { for(auto& c : components) c->draw(); }
};
```

Encoding entities – OOP composition

```
struct BonesComponent : Component
{
    std::vector<Bone> bones;
    void update() override
    {
        // do things skeletons do
    }
};

auto makeSkeleton()
{
    Entity e;
    e.components.emplace_back(std::make_unique<BonesComponent>());
    return e;
}
```

Encoding entities – DOD composition

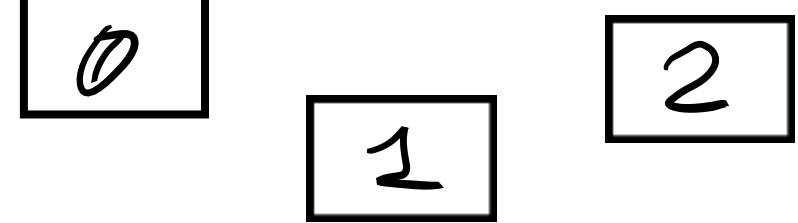
- An **entity** is a numerical **ID**.
- Components only store **data** (logicless).
- **Logic** is handled using **systems**.
- Potentially **cache-friendly**.
- Minimal **runtime overhead**.
- Great **flexibility**.
- Hard to implement.

ID	NAME	KEYBOARD	STYLE	Focus	Mouse
0	TextBox	✓	✓	✓	
1	Button		✓	✓	✓
2	Browser	✓		✓	✓

Encoding entities – DOD composition

	AI	SWORD	SHIELD
SKELETON # 0			
SKELETON # 1			
SKELETON # 2			

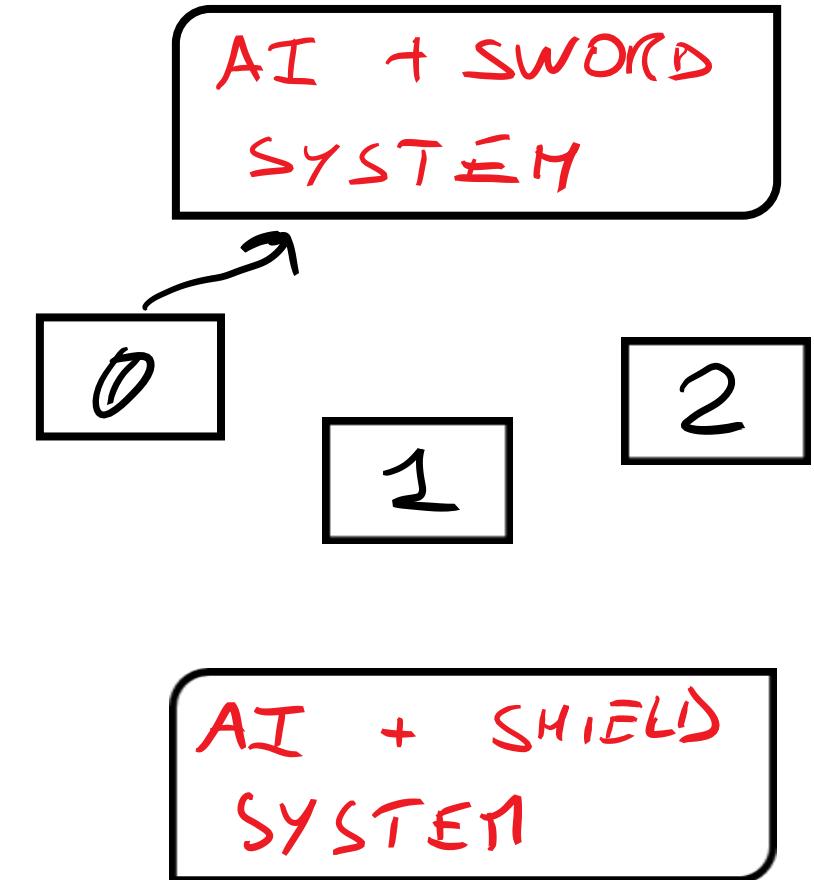
AI + SWORD
SYSTEM



AI + SHIELD
SYSTEM

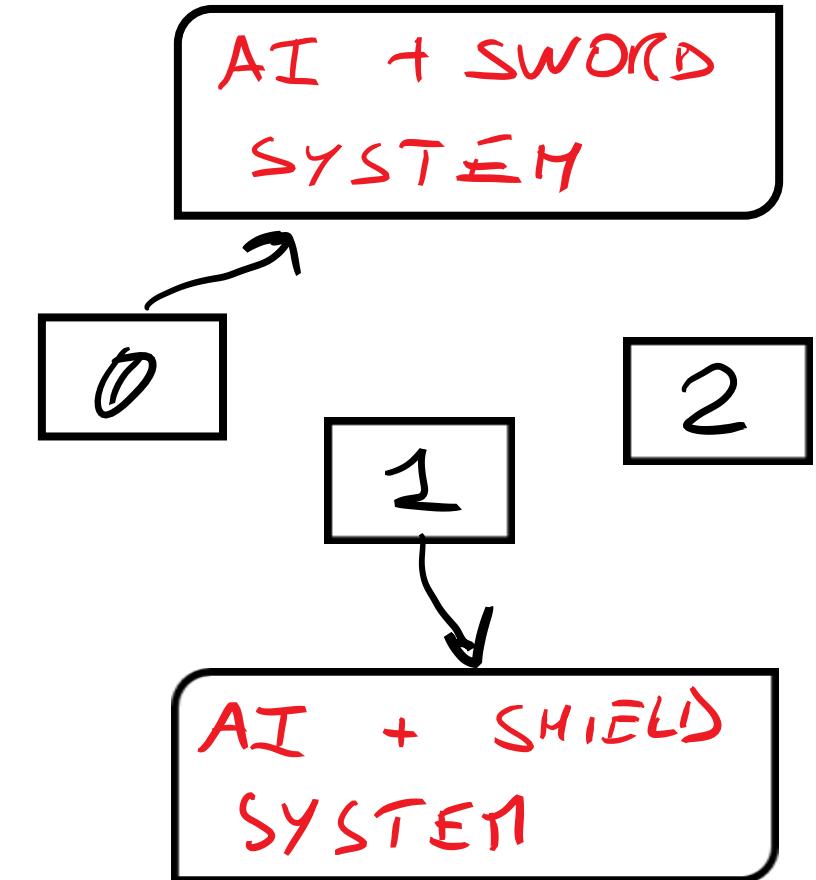
Encoding entities – DOD composition

	AI	SWORD	SHIELD
SKELETON # 0	✓	✓	
SKELETON # 1			
SKELETON # 2			



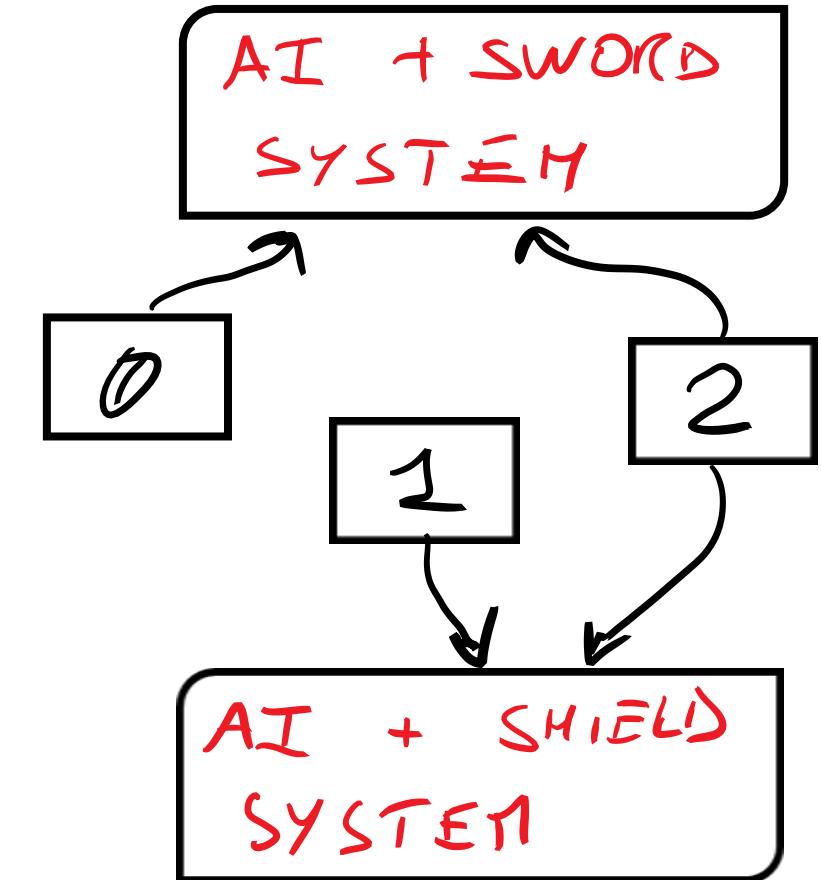
Encoding entities – DOD composition

	AI	SWORD	SHIELD
SKELETON # 0	✓	✓	
SKELETON # 1	✓		✓
SKELETON # 2			

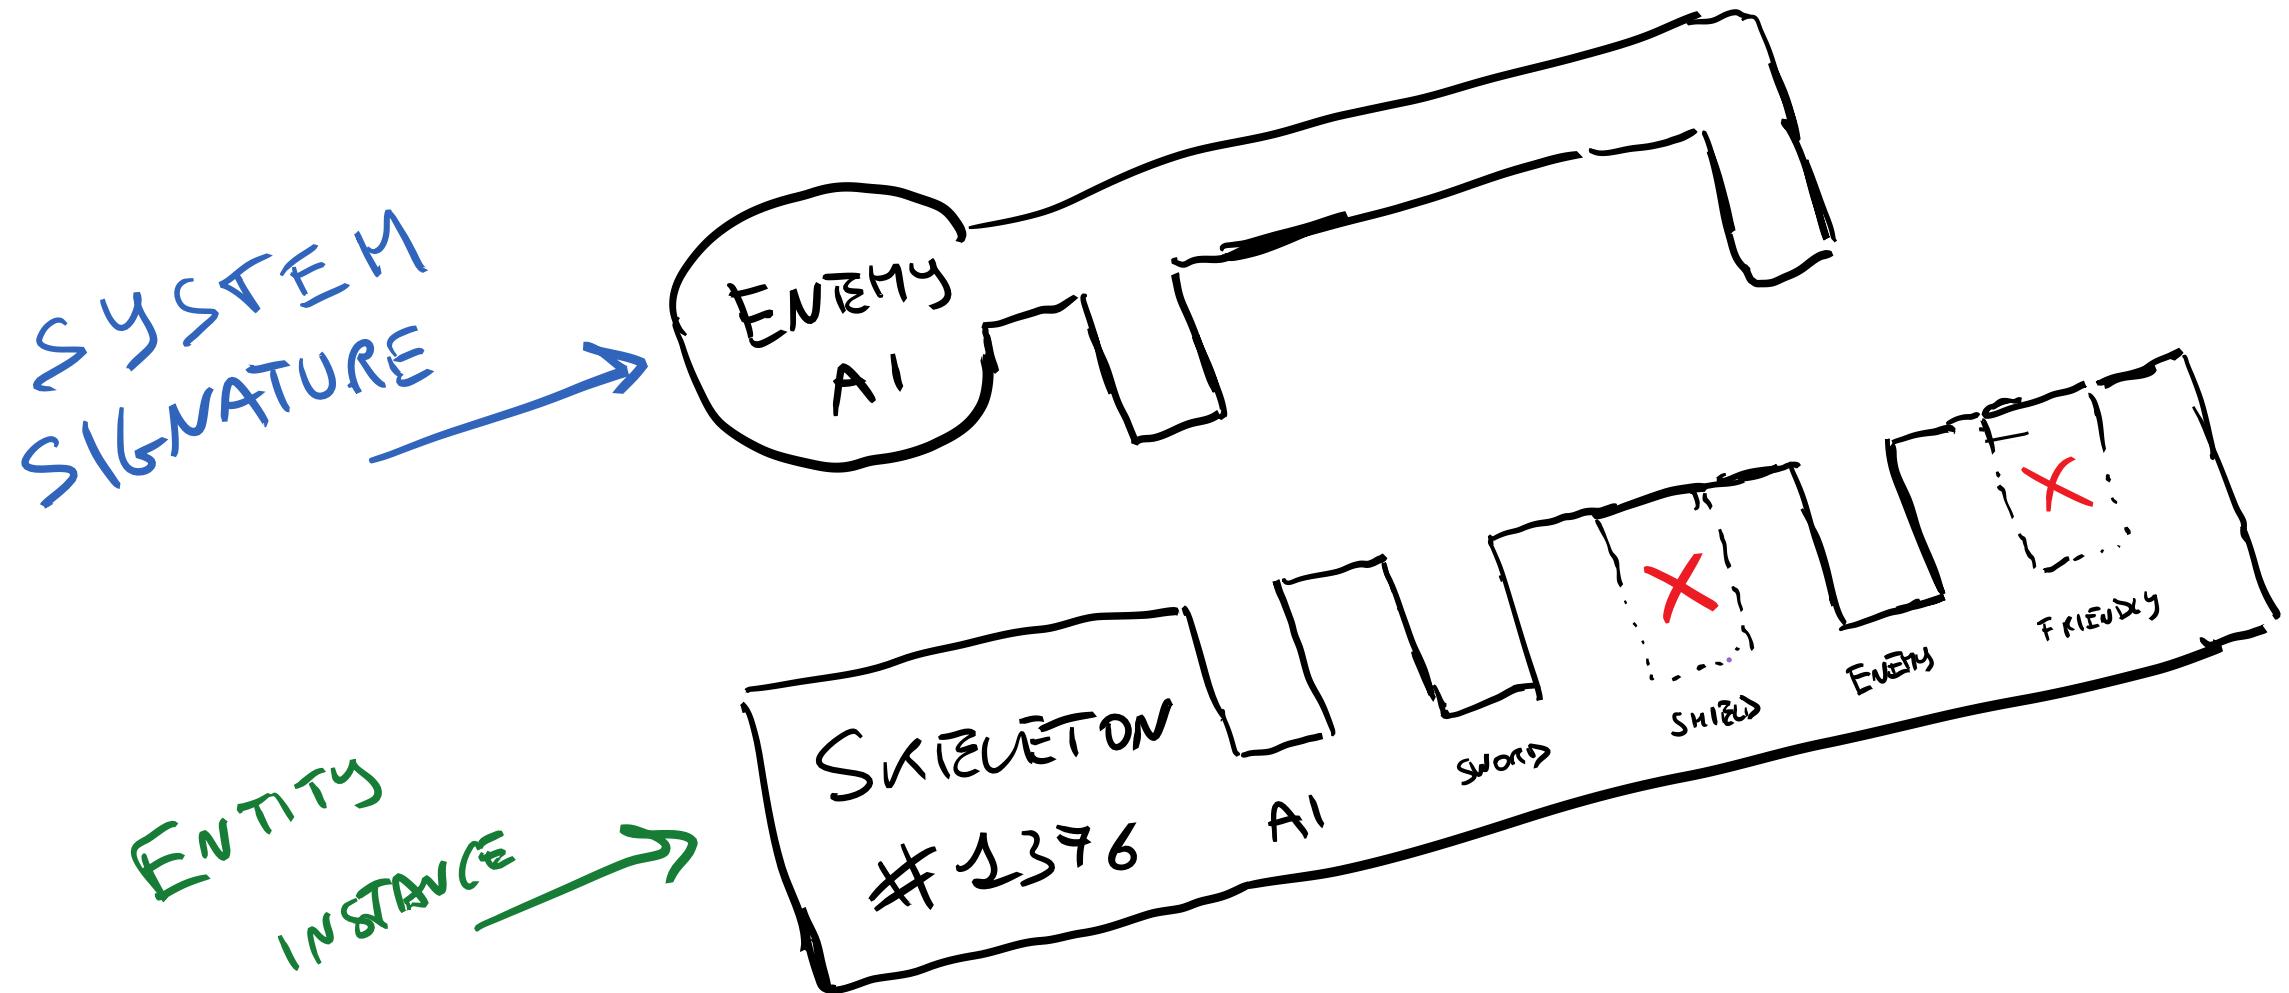


Encoding entities – DOD composition

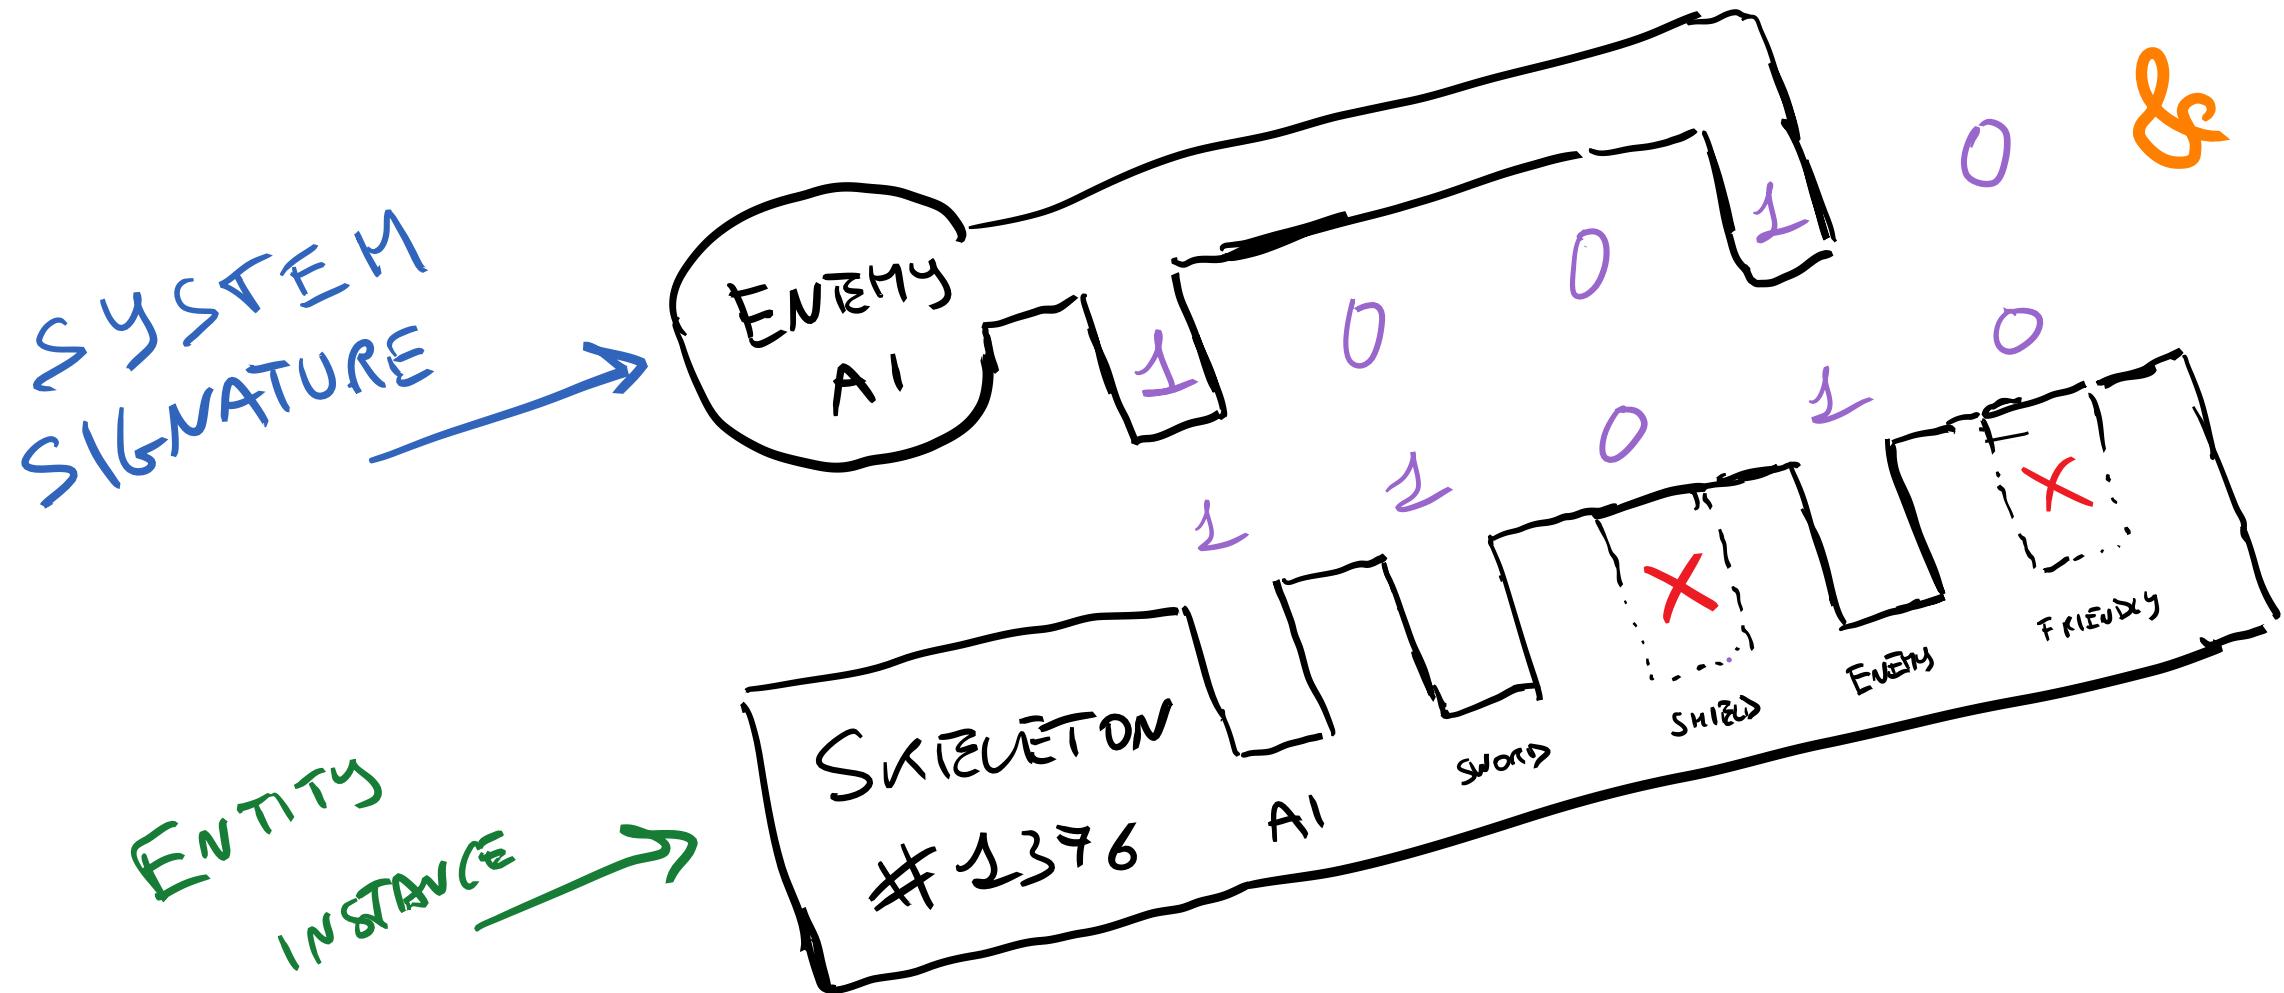
	AI	SWORD	SHIELD
SKELETON # 0	✓	✓	
SKELETON # 1	✓		✓
SKELETON # 2	✓	✓	✓



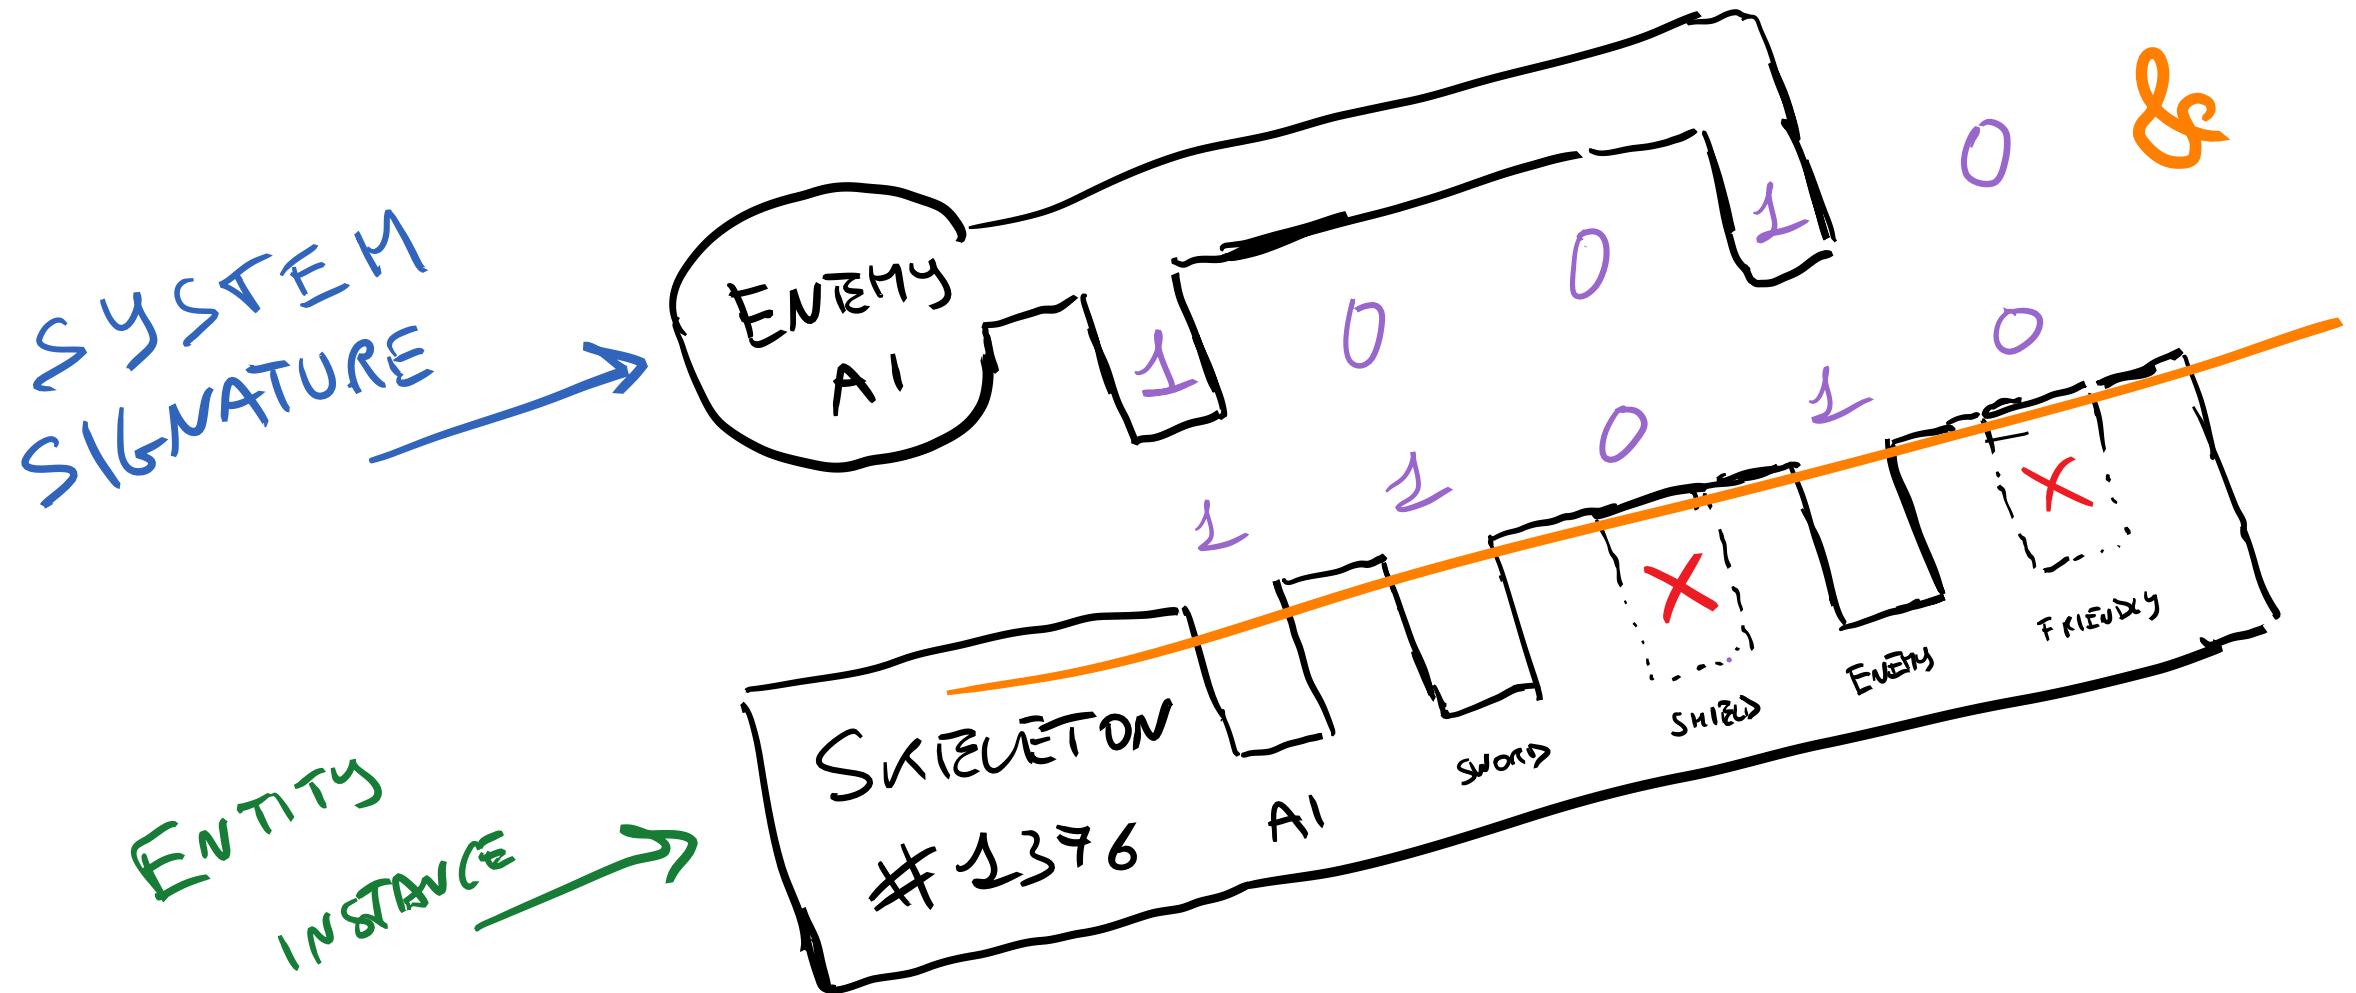
Encoding entities – DOD composition



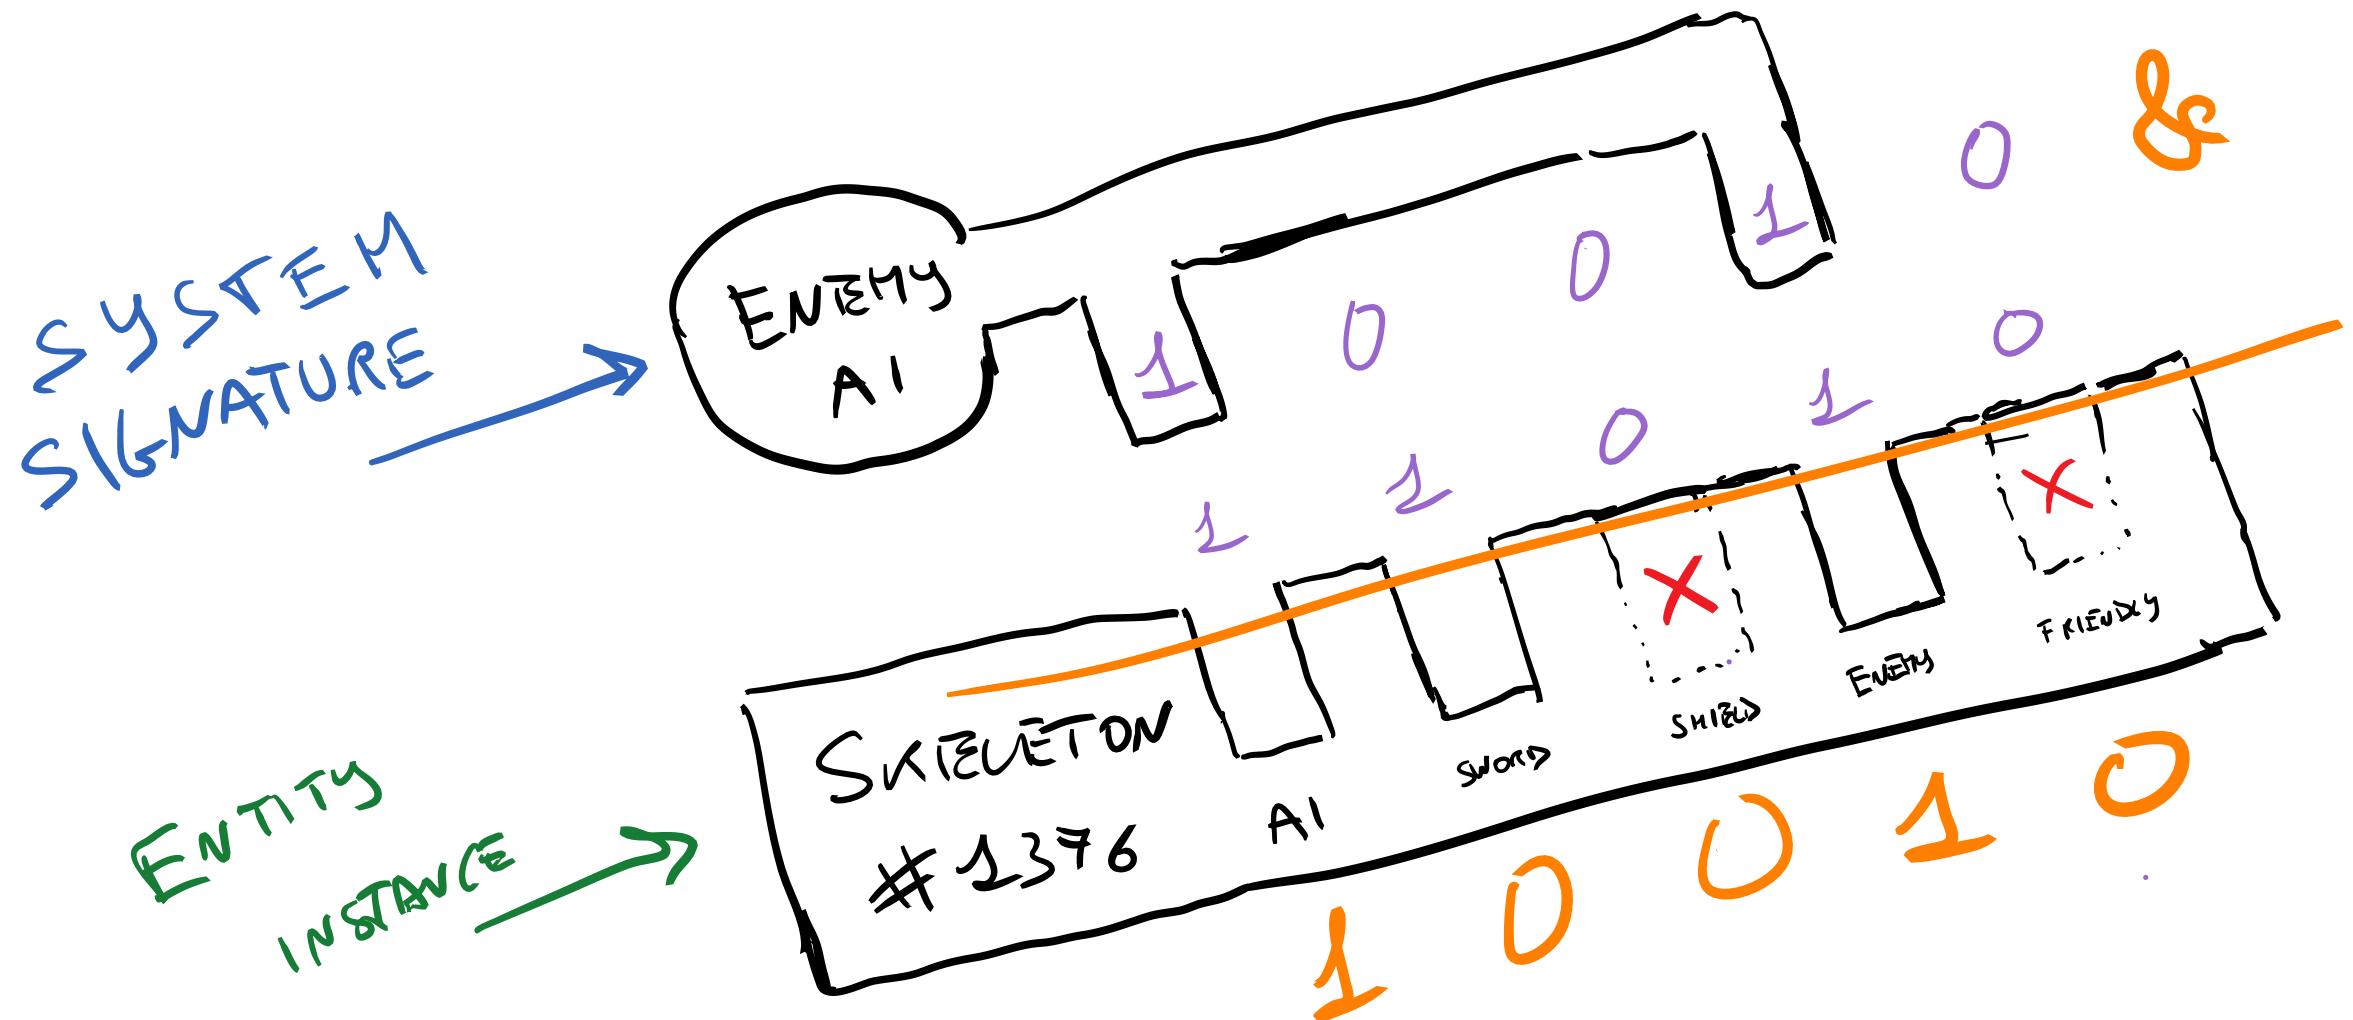
Encoding entities – DOD composition



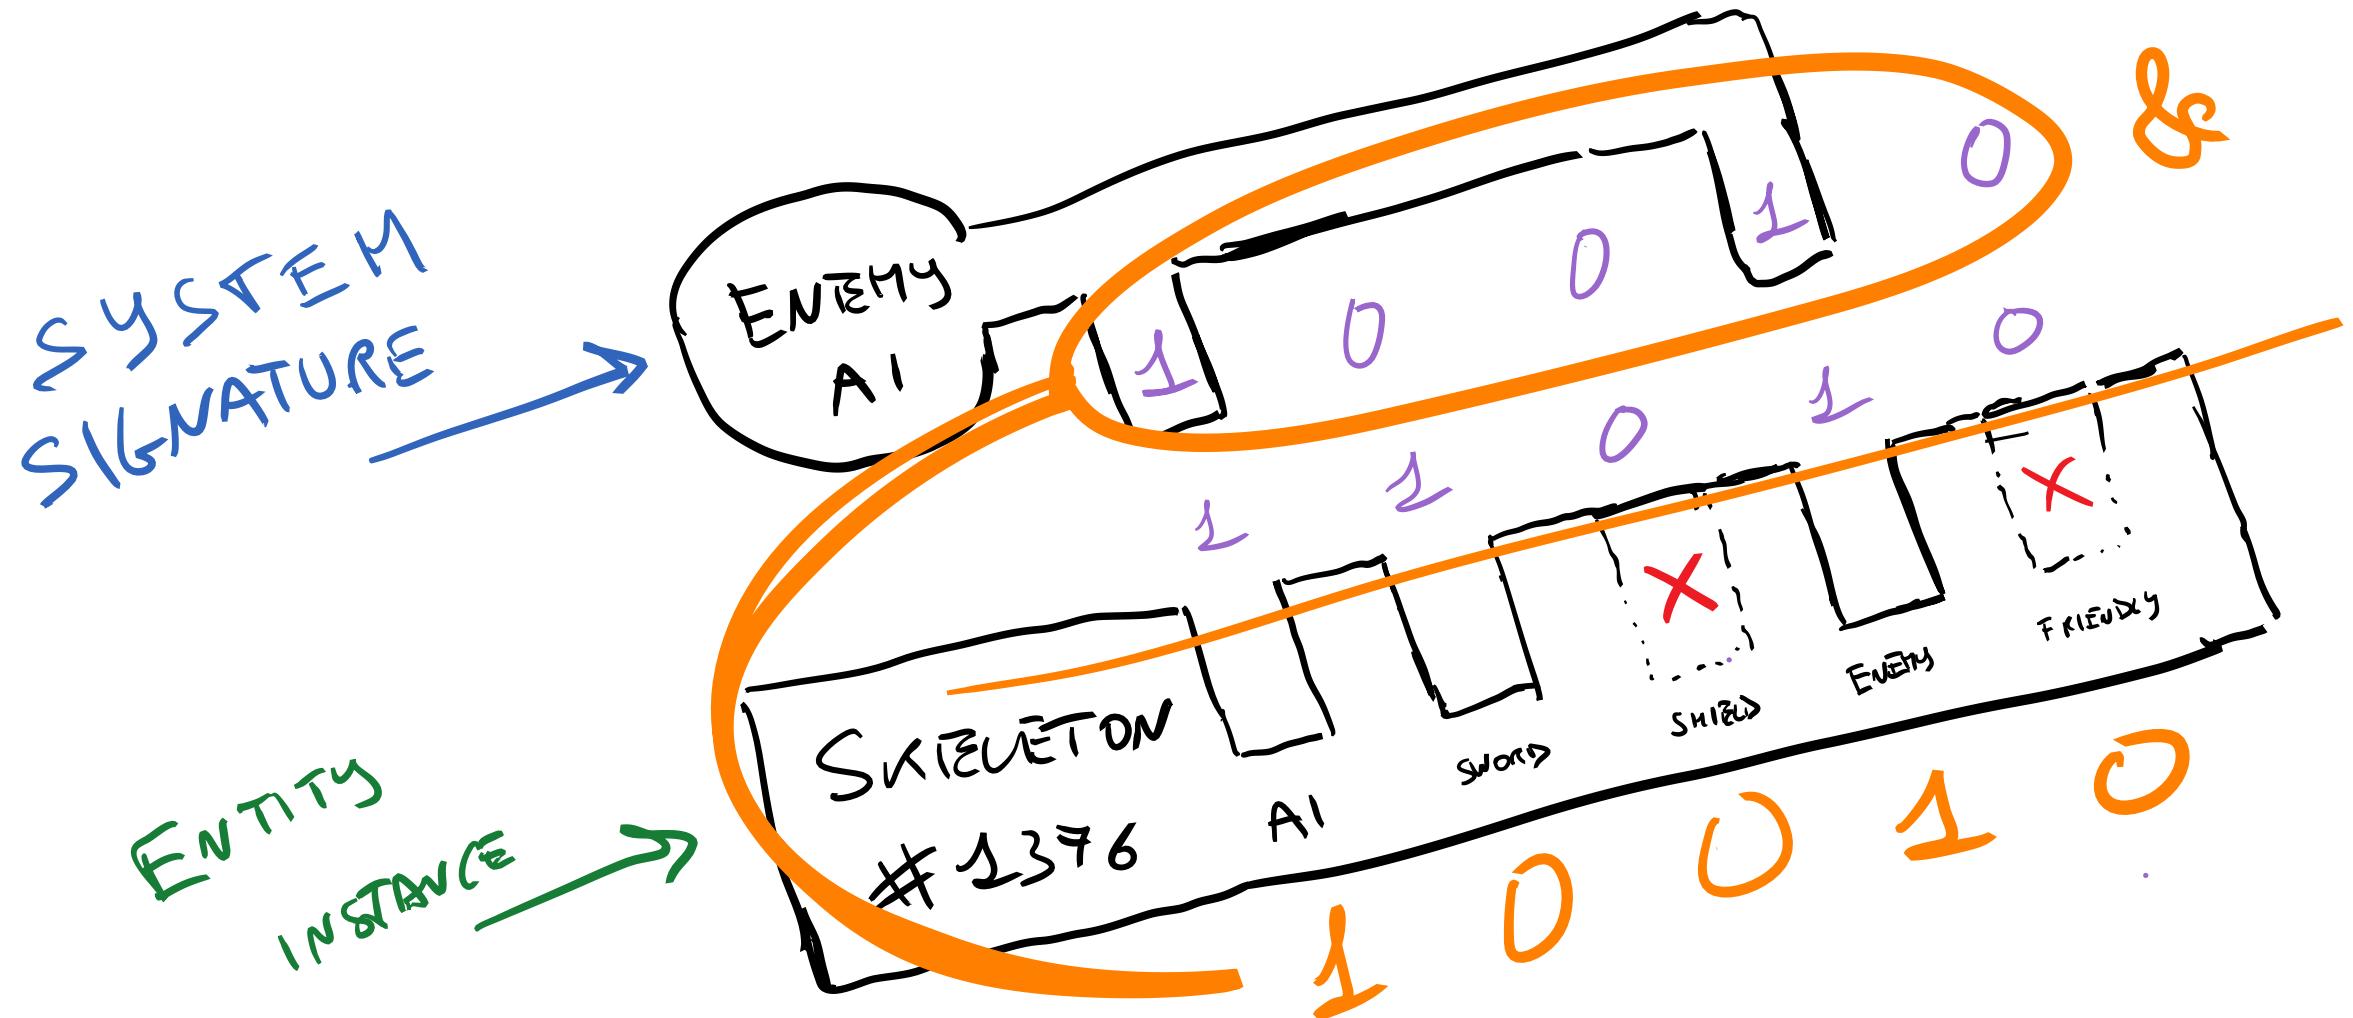
Encoding entities – DOD composition



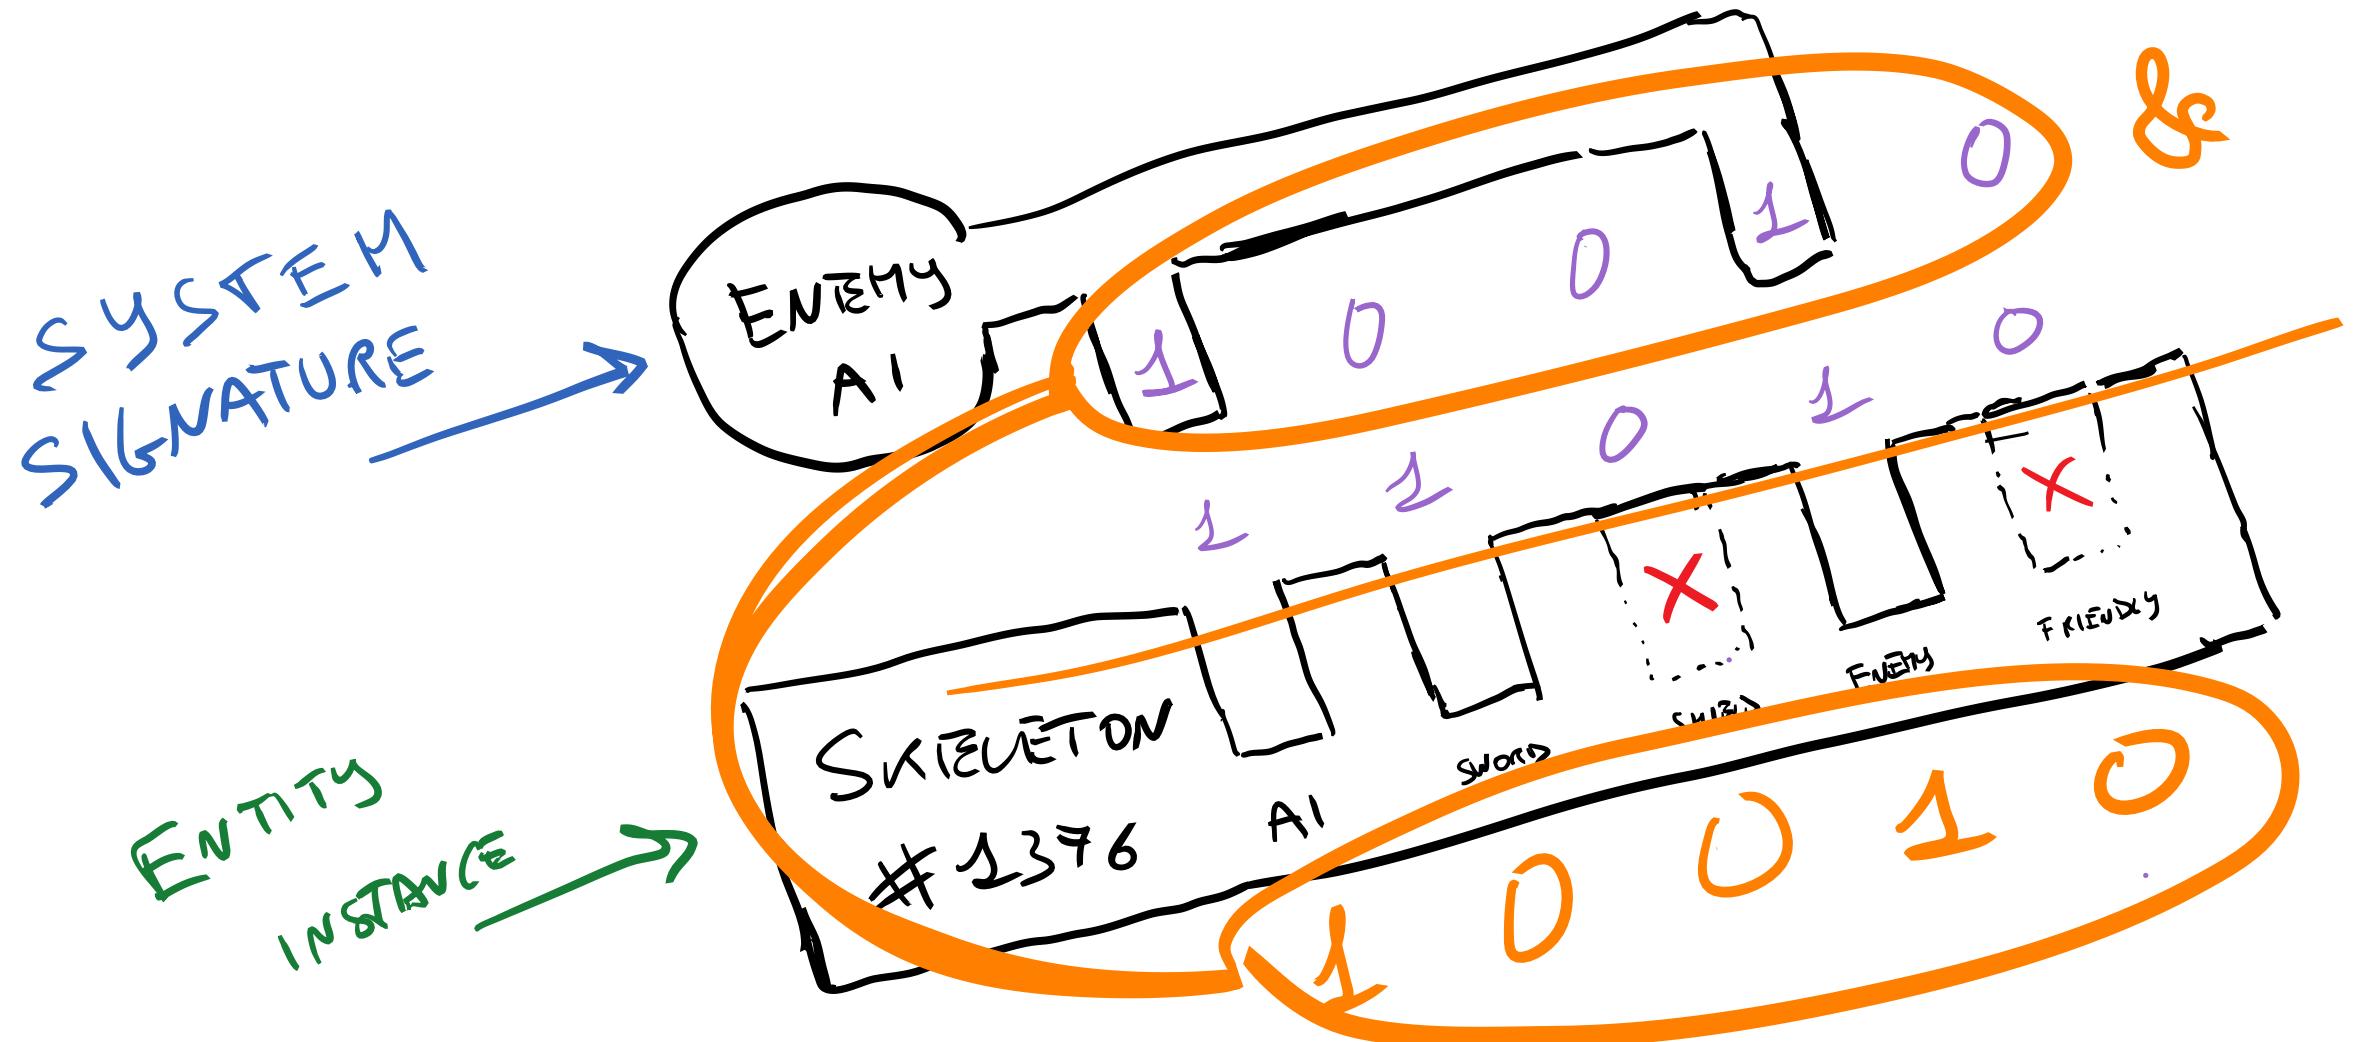
Encoding entities – DOD composition



Encoding entities – DOD composition



Encoding entities – DOD composition



Encoding entities – DOD composition

```
using Entity = std::size_t;
constexpr std::size_t maxEntities{1000};

struct BonesComponent { /* ... */ };
struct AIComponent { /* ... */ };
struct SpriteComponent { /* ... */ };

struct Manager
{
    std::array<BonesComponent, maxEntities> bonesComponents;
    std::array<AIComponent, maxEntities> aiComponents;
    std::array<SpriteComponent, maxEntities> spritesComponents;

    // ...
};
```

Encoding entities – DOD composition

```
using Entity = std::size_t;
constexpr std::size_t maxEntities{1000};

struct BonesComponent { /* ... */ };
struct AIComponent { /* ... */ };
struct SpriteComponent { /* ... */ };

struct Manager
{
    std::array<BonesComponent, maxEntities> bonesComponents;
    std::array<AIComponent, maxEntities> aiComponents;
    std::array<SpriteComponent, maxEntities> spritesComponents;

    // ...
};
```

Encoding entities – DOD composition

```
using Entity = std::size_t;
constexpr std::size_t maxEntities{1000};

struct BonesComponent { /* ... */ };
struct AIComponent { /* ... */ };
struct SpriteComponent { /* ... */ } } POD

struct Manager
{
    std::array<BonesComponent, maxEntities> bonesComponents;
    std::array<AIComponent, maxEntities> aiComponents;
    std::array<SpriteComponent, maxEntities> spritesComponents;

    // ...
};
```



Encoding entities – DOD composition

```
using Entity = std::size_t;
constexpr std::size_t maxEntities{1000};

struct BonesComponent { /* ... */ };
struct AIComponent { /* ... */ };
struct SpriteComponent { /* ... */ } } POD

struct Manager
{
    std::array<BonesComponent, maxEntities> bonesComponents;
    std::array<AIComponent, maxEntities> aiComponents;
    std::array<SpriteComponent, maxEntities> spritesComponents;

    // ...
};
```



Encoding entities – DOD composition

```
struct SkeletonSystem
{
    Manager& manager;

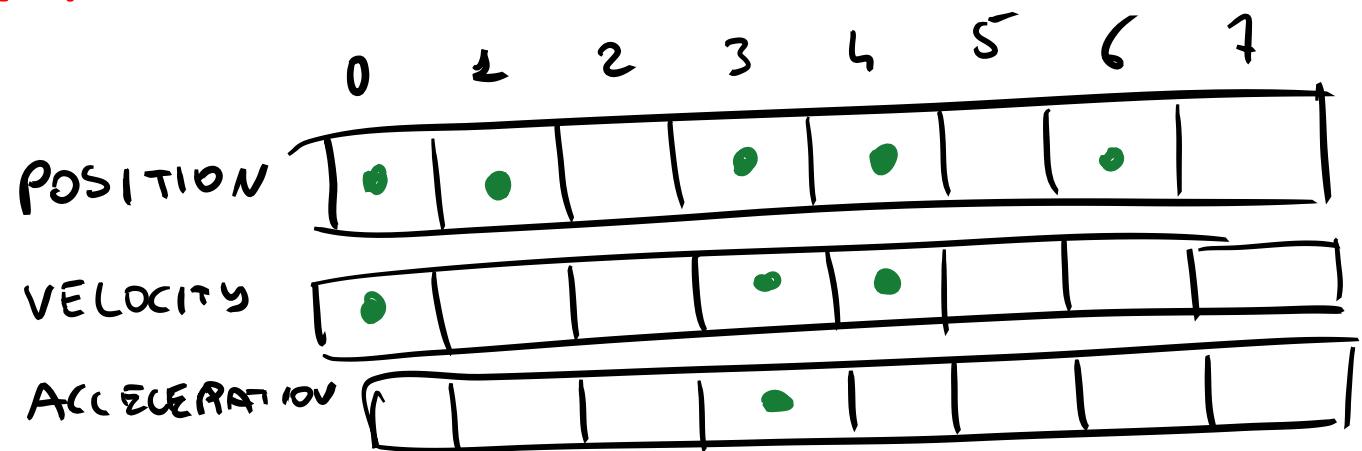
    void process()
    {
        manager.forEntitiesWith<AIComponent, BonesComponent>(
            [](auto& data, auto id)
            {
                auto& ai = data.get<AIComponent>(id);
                auto& bones = data.get<BonesComponent>(id);

                // ...
            });
    }
};
```



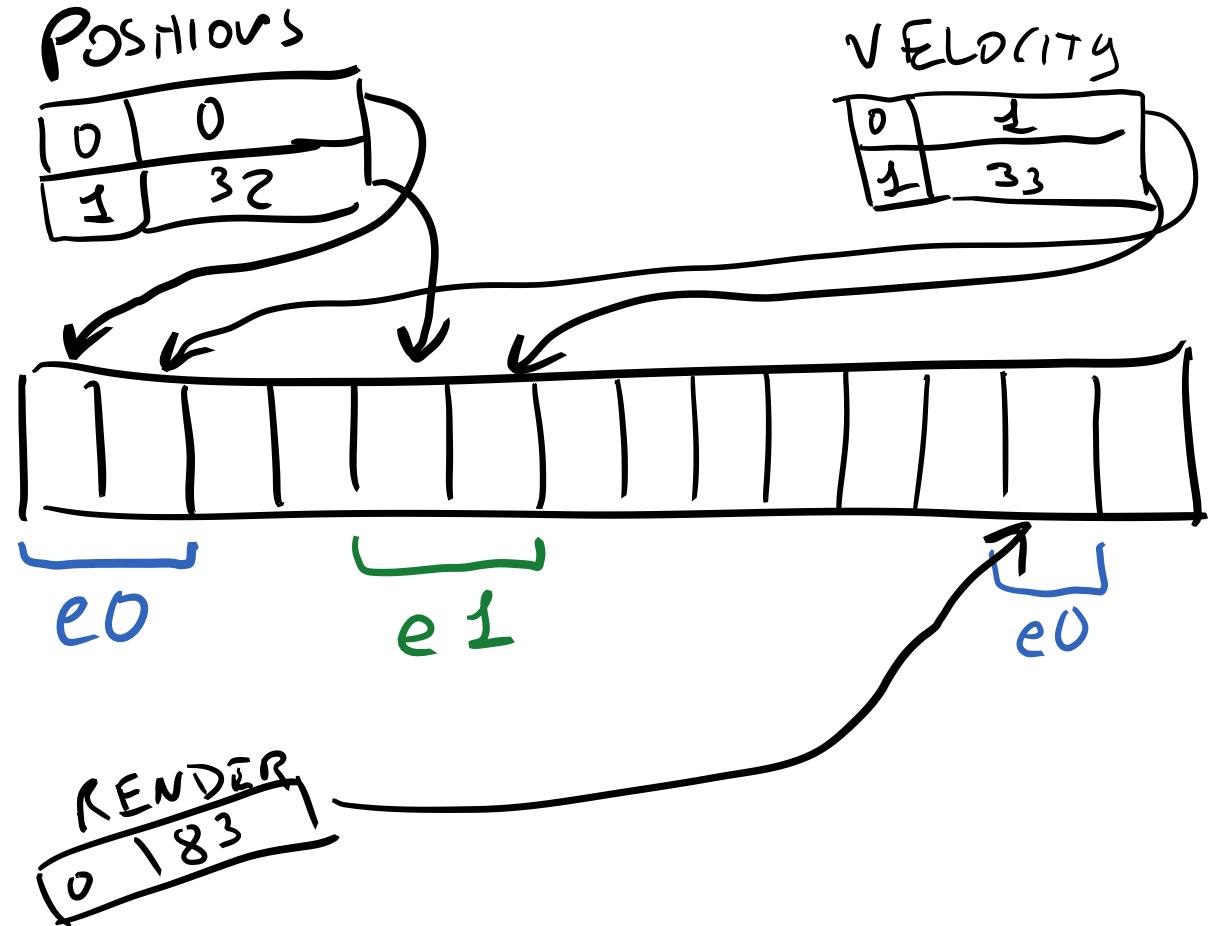
Storing components - one array per type

- Very easy to implement.
- Suitable for most projects.
- Easy to add/remove components at runtime.
- Can be «cache-friendlier».
- Wasteful of memory.



Storing components – mega-array

- Potentially **cache-friendlier**.
- **Minimizes** memory waste.
- **Very hard** to implement.
 - Requires **indexing tables**.
 - Hard to deal with component addition/removal.
 - **Very hard** to keep track of free «slots» in the mega-array.



Overview of ECST

Design, features, limitations and examples.



<http://vittorioromeo.info>
vittorio.romeo@outlook.com

<http://github.com/SuperV1234/cppnow2016>

Core values and concepts

- «**Compile-time**» ECS.
- Customizable **settings/policy-based** design.
- User-friendly syntax, independent from settings.
- **Multithreading** support.
 - Avoid explicit locking.
 - «**Dataflow**» programming.
- «**Type-value encoding**» modern metaprogramming.
- Clean, modern and safe C++14.

«Compile-time» ECS

- Component and system **types** are known at **compile-time**.
- Component **instances** can be created and destroyed at **run-time**.
- **Entities** can be created, destroyed and tracked at **run-time**.
- Combining it with run-time solutions is possible.
- Minimal run-time overhead.
- Better optimization opportunities.
- Longer compilation times.
- Unwieldy error messages.

«Compile-time» ECS

```
struct position
{
    float _x;
    float _y;
};

struct velocity
{
    float _x;
    float _y;
};

struct acceleration
{
    float _x;
    float _y;
};

constexpr auto ssig_acceleration =
    ss::make<s::acceleration>(
        ipc::none_below_threshold::v(sz_v<10000>,
            ips::split_evenly_fn::v_cores()
        ),
        ss::no_dependencies,
        ss::component_use(
            ss::mutate<c::velocity>,
            ss::read<c::acceleration>
        ),
        ss::output::none
    );
```

```
struct acceleration
{
    template <typename TData>
    void process(TData& data)
    {
        data.for_entities([&](auto eid)
        {
            auto& v = ecst::get<c::velocity>(data, eid);
            const auto& a = ecst::get<c::acceleration>(data, eid);
            actions::accelerate(v, a);
        });
    }
};
```

«Compile-time» ECS

```
struct position
{
    float _x;
    float _y;
};

struct velocity
{
    float _x;
    float _y;
};

struct acceleration
{
    float _x;
    float _y;
};

constexpr auto ssig_acceleration =
    ss::make<s::acceleration>(
        ipc::none_below_threshold::v(sz_v<10000>,
            ips::split_evenly_fn::v_cores()
        ),
        ss::no_dependencies,
        ss::component_use(
            ss::mutate<c::velocity>,
            ss::read<c::acceleration>
        ),
        ss::output::none
    );
```

```
struct acceleration
{
    template <typename TData>
    void process(TData& data)
    {
        data.for_entities([&](auto eid)
        {
            auto& v = ecst::get<c::velocity>(data, eid);
            const auto& a = ecst::get<c::acceleration>(data, eid);
            actions::accelerate(v, a);
        });
    }
};
```



«Compile-time» ECS

```
struct position
{
    float _x;
    float _y;
};

struct velocity
{
    float _x;
    float _y;
};

struct acceleration
{
    float _x;
    float _y;
};
```

```
struct acceleration
{
    template <typename TData>
    void process(TData& data)
    {
        data.for_entities([&](auto eid)
        {
            auto& v = ecst::get<c::velocity>(data, eid);
            const auto& a = ecst::get<c::acceleration>(data, eid);
            actions::accelerate(v, a);
        });
    }
};
```

```
constexpr auto ssig_acceleration =
ss::make<s::acceleration>(
    ipc::none_below_threshold::v(sz_v<10000>,
        ips::split_evenly_fn::v_cores()
    ),
    ss::no_dependencies,
    ss::component_use(
        ss::mutate<c::velocity>,
        ss::read<c::acceleration>
    ),
    ss::output::none
);
```

«Compile-time» ECS – g++ errors

```
strategy::none::impl::parameters, std::tuple<ecst::signature::system::impl::tag<impl<example::s::spatial_partition>>, std::tuple<ecst::signature::system::output::impl::option>, ecst::inner_parallelism::strategy::split_evenly_fn::impl::parameters<ecst::inner_parallelism::strategy::split_evenly_fn::v_cores_getter>, std::tuple<ecst::signature::system::impl::mutate<impl<example::c::position>>, ecst::signature::system::impl::read<impl<example::c::circle>>, ecst::signature::system::impl::tag<impl<example::s::solve_contacts>>, ecst::inner_parallelism::strategy::none::impl::parameters, std::tuple<ecst::signature::system::impl::mutate<impl<example::c::velocity>>, ecst::signature::system::impl::mutate<impl<example::c::position>>, ecst::signature::system::impl::read<impl<example::c::circle>>, ecst::signature::system::impl::tag<impl<example::s::render_colored_circle>>, ecst::inner_parallelism::strategy::split_evenly_fn::impl::parameters<ecst::inner_parallelism::strategy::split_contacts>, std::tuple<ecst::signature::system::impl::read<impl<example::c::circle>>, ecst::signature::system::impl::tag<impl<example::s::color>>, ecst::signature::system::impl::read<impl<example::c::color>>, ecst::signature::system::impl::tag<impl<example::c::velocity>>, std::tuple<sf::Vertex>>>>, ecst::settings::impl::scheduler_unwrapper<ecst::scheduler::s_atomic_counter>; TSystemSignature = ecst::signature::system::impl::isn::strategy::split_evenly_fn::impl::parameters<ecst::inner_parallelism::strategy::split_evenly_fn::v_cores_getter>, std::tuple<ecst::signature::system::impl::tag<impl<example::c::position>>, ecst::signature::system::impl::read<impl<example::c::color>>, ecst::signature::system::impl::tag<impl<example::c::position>>, ecst::signature::system::impl::read<impl<example::c::circle>>, ecst::signature::system::impl::tag<impl<example::s::acceleration>>, ecst::signature::component::impl::dataecst::signature::component::impl::tag<impl<example::c::acceleration>>, ecst::signature::component::impl::dataecst::signature::component::impl::tag<impl<example::c::mass>>>, std::tuple<ecst::signature::component::impl::tag<impl<example::c::velocity>>, ecst::inner_parallelism::strategy::split_evenly_fn::impl::parameters<ecst::inner_parallelism::strategy::split_evenly_fn::v_cores_getter>, std::tuple<ecst::signature::component::impl::tag<impl<example::c::circle>>, ecst::signature::component::impl::dataecst::signature::component::impl::tag<impl<example::c::mass>>>, std::tuple<ecst::signature::component::impl::tag<impl<example::c::velocity>>, ecst::inner_parallelism::strategy::split_evenly_fn::impl::parameters<ecst::inner_parallelism::strategy::split_evenly_fn::v_cores_getter>, std::tuple<ecst::signature::component::impl::tag<impl<example::c::circle>>, ecst::signature::component::impl::dataecst::signature::component::impl::tag<impl<example::c::position>>, ecst::signature::component::impl::tag<impl<example::s::keep_in_bounds>>, ecst::inner_parallelism::strategy::split_evenly_fn::impl::parameters<ecst::inner_parallelism::strategy::split_evenly_fn::v_cores_getter>, std::tuple<ecst::signature::component::impl::tag<impl<example::s::keep_in_bounds>>, ecst::inner_parallelism::strategy::split_evenly_fn::impl::parameters<ecst::inner_parallelism::strategy::split_evenly_fn::v_cores_getter>, std::tuple<ecst::signature::component::impl::tag<impl<example::s::spatial_partition>>, std::allocator<example::sp_data>>>, ecst::signature::system::impl::dataecst::signature::system::tuple<ecst::signature::system::impl::tag<impl<example::s::spatial_partition>>, std::tuple<ecst::signature::system::impl::tag<impl<example::c::velocity>>, ecst::signature::system::impl::mutate<impl<example::c::velocity>>, ecst::signature::component::impl::dataecst::signature::component::impl::tag<impl<example::c::circle>>, ecst::signature::system::impl::tag<impl<example::c::position>>, ecst::signature::system::impl::read<impl<example::c::circle>>, ecst::signature::system::impl::tag<impl<example::s::solve_contacts>>, ecst::inner_parallelism::strategy::none::impl::parameters, std::tuple<ecst::signature::system::impl::tag<impl<example::c::velocity>>, ecst::signature::system::impl::read<impl<example::c::circle>>, ecst::signature::system::impl::tag<impl<example::s::render_colored_circle>>, ecst::inner_parallelism::strategy::split_evenly_fn::impl::parameters<ecst::inner_parallelism::strategy::split_evenly_fn::v_cores_getter>, std::tuple<ecst::signature::system::impl::read<impl<example::c::circle>>, ecst::signature::system::impl::tag<impl<example::c::color>>>, ecst::settings::impl::scheduler_unwrapper<ecst::scheduler::s_atomic_counter>; TSystemSignature = ecst::signature::system::impl::dataecst::signature::system::impl::tag<impl<example::c::velocity>>, std::tuple<ecst::signature::system::impl::tag<impl<example::c::position>>, ecst::signature::system::impl::read<impl<example::c::color>>, ecst::signature::system::impl::tag<impl<example::s::solve_contacts>>, std::tuple<sf::Vertex>>>, v_settings::impl::dataecst::settings::impl::threading::multicst::settings::impl::system_parallelism::enabled, ecst::settings::impl::fixed_Impl<std::integral_constant<long::impl::tag<impl<example::c::position>>, ecst::signature::component::impl::dataecst::signature::component::impl::tag<impl<example::c::velocity>>, ecst::signature::component::impl::tag<impl<example::c::color>>, ecst::signature::component::impl::dataecst::signature::component::impl::tag<impl<example::c::mass>>>, std::tuple<ecst::signature::system::impl::tag<impl<example::c::acceleration>>, ecst::inner_parallelism::strategy::split_evenly_fn::impl::parameters<ecst::inner_parallelism::strategy::split_evenly_fn::v_cores_getter>, std::tuple<ecst::signature::system::impl::tag<impl<example::c::velocity>>, ecst::signature::system::impl::mutate<impl<example::c::velocity>>, ecst::signature::system::impl::read<impl<example::c::acceleration>>, ecst::inner_parallelism::strategy::split_evenly_fn::impl::parameters<ecst::inner_parallelism::strategy::split_evenly_fn::v_cores_getter>, std::tuple<ecst::signature::system::impl::tag<impl<example::c::velocity>>, ecst::signature::system::impl::tag<impl<example::c::acceleration>>
```



«Compile-time» ECS – clang++ errors

```
pres_code.cpp:116:70: error: no member named '_x' in 'example::c::position'
    const auto& p0 = data.get(ct::position, eid)._x;
                                         ^~~~~~
/home/vittorioromeo/01Workspace/ecs_thesis/project/include/ecst/context./system./instance/instance.inl:112:32: note: in ins
'example::s::render_colored_circle::process(ecst::context::system::impl::execute_data
```

Customizable settings/policy-based design

- Behavior and storage layouts can be chosen at compile-time.
- Library users can create their own settings easily, including:
 - Scheduling algorithms.
 - Parallelization strategies.
 - Entity and component data structures.
- Focus on **composability**.

```
template <typename TComponent>
class hash_map
{
public:
    using entity_id = ecst::impl::entity_id;
    using component_type = TComponent;

    struct metadata
    {
    };

private:
    std::unordered_map<sz_t, TComponent> _data;

    auto valid_index(sz_t i) const noexcept
    {
        return _data.count(i) > 0;
    }

    auto entity_id_to_index(entity_id eid) const noexcept
    {
        return vrmc::to_sz_t(eid);
    }

    template <typename TSelf>
    decltype(auto) get_impl(
        TSelf&& self, entity_id eid, const metadata&) noexcept
    {
        auto i = self.entity_id_to_index(eid);
        return _data[i];
    }
};
```



Syntax: independent from settings

- Chosen settings and strategies do not affect the syntax of the application logic.

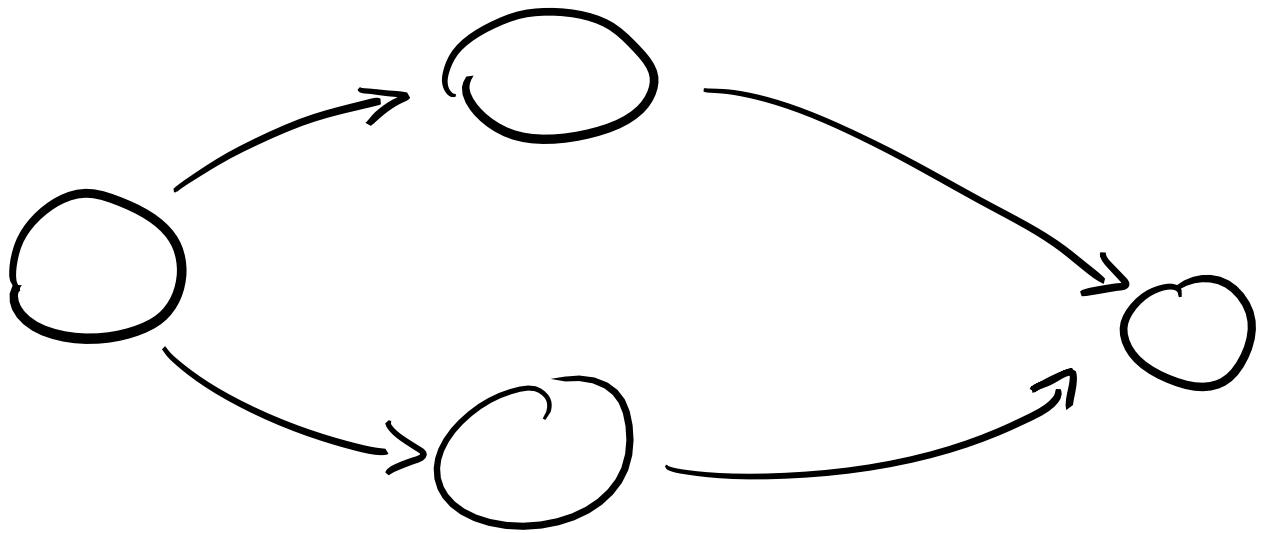
```
struct acceleration
{
    template <typename TData>
    void process(TData& data)
    {
        data.for_entities([&](auto eid)
        {
            auto& v = ecst::get<c::velocity>(data, eid);
            const auto& a = ecst::get<c::acceleration>(data, eid);
            actions::accelerate(v, a);
        });
    }
};
```

```
proxy.execute_systems_overload( // .
    [](&s)::acceleration& s, auto& data)
{
    s.process(data);
},
    [](&s)::velocity& s, auto& data)
{
    s.process(data);
},
    [&window](&s)::rendering& s, auto& data)
{
    s.process(window, data);
});
```



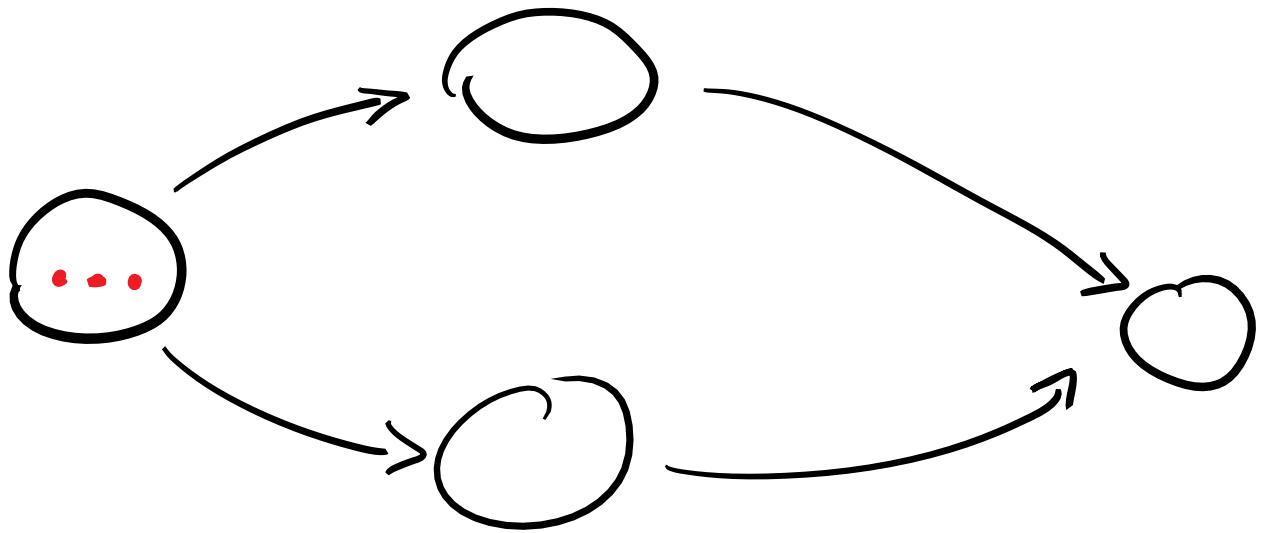
Multithreading support

- Enabled by default, customizable and can be optionally disabled.
- Two levels of parallelism:
 - **Outer**: independent system chains can run in separate parallel tasks.
 - **Inner**: system logic can be independently performed on entity subsets in separate tasks.



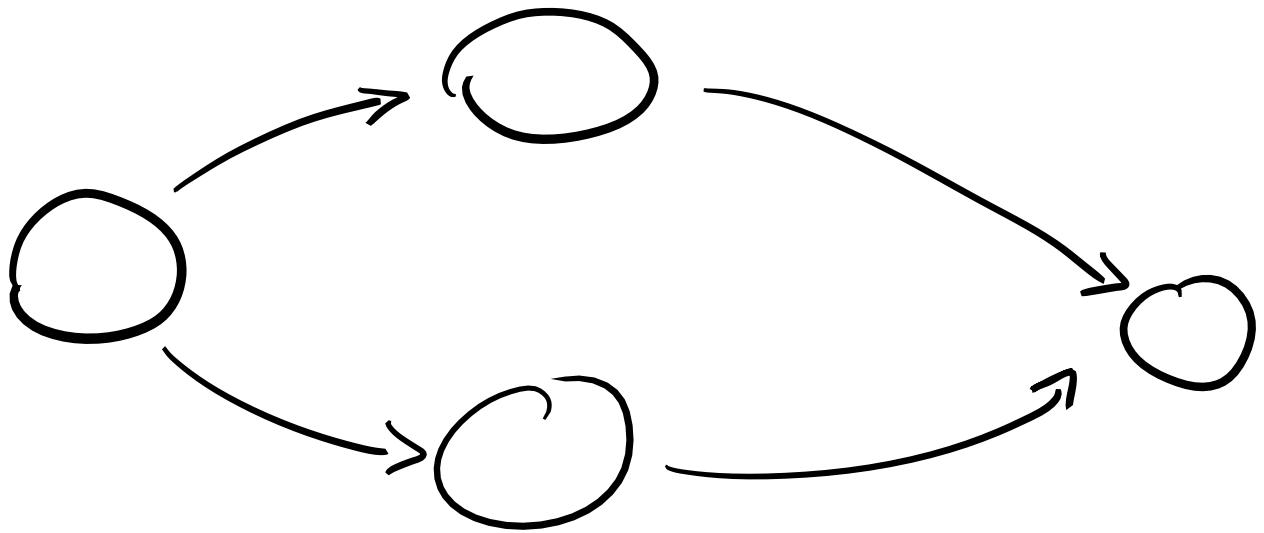
Multithreading support

- Enabled by default, customizable and can be optionally disabled.
- Two levels of parallelism:
 - **Outer**: independent system chains can run in separate parallel tasks.
 - **Inner**: system logic can be independently performed on entity subsets in separate tasks.



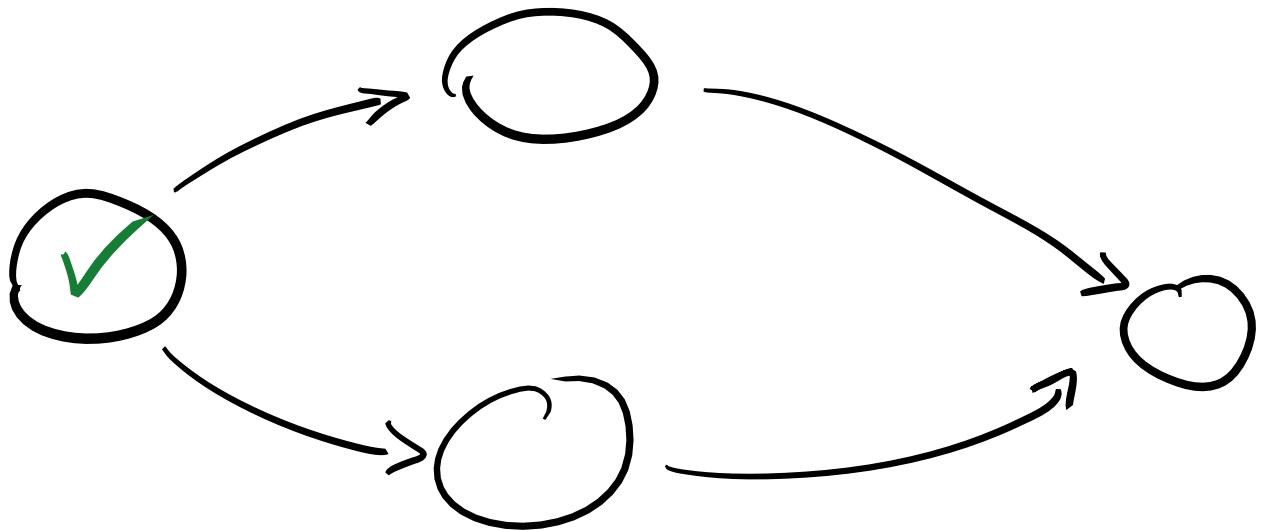
Multithreading support

- Enabled by default, customizable and can be optionally disabled.
- Two levels of parallelism:
 - **Outer**: independent system chains can run in separate parallel tasks.
 - **Inner**: system logic can be independently performed on entity subsets in separate tasks.



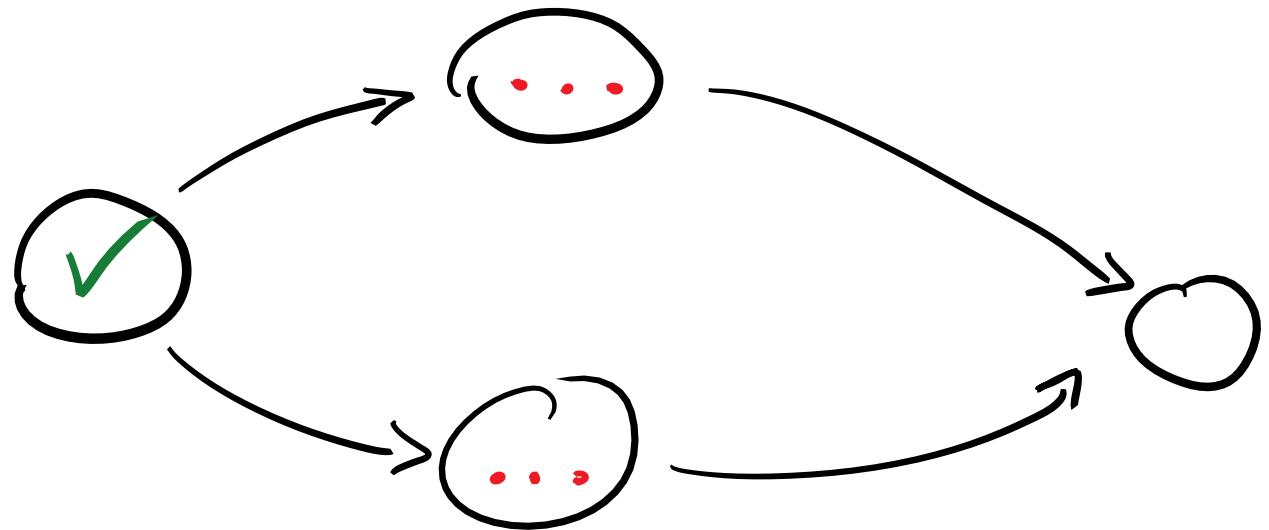
Multithreading support

- Enabled by default, customizable and can be optionally disabled.
- Two levels of parallelism:
 - **Outer**: independent system chains can run in separate parallel tasks.
 - **Inner**: system logic can be independently performed on entity subsets in separate tasks.



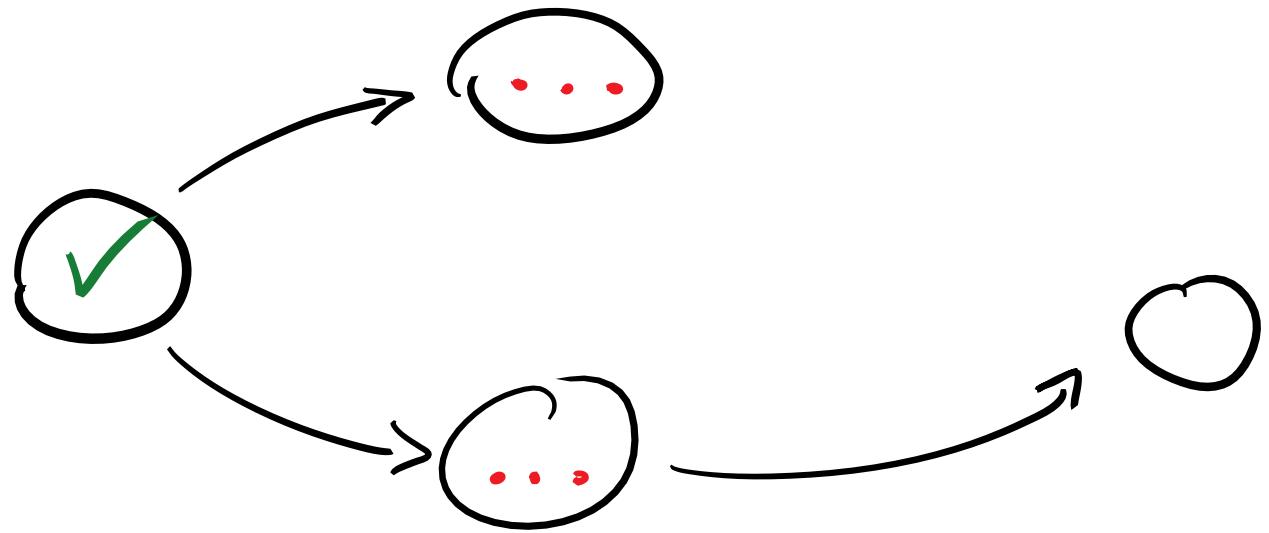
Multithreading support

- Enabled by default, customizable and can be optionally disabled.
- Two levels of parallelism:
 - **Outer**: independent system chains can run in separate parallel tasks.
 - **Inner**: system logic can be independently performed on entity subsets in separate tasks.



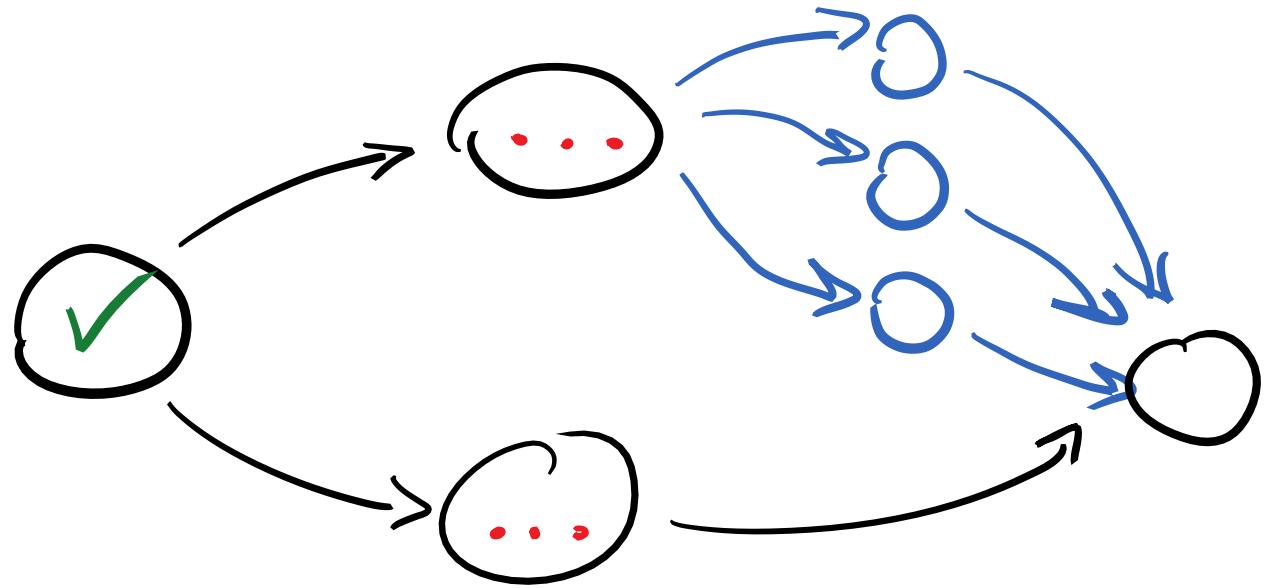
Multithreading support

- Enabled by default, customizable and can be optionally disabled.
- Two levels of parallelism:
 - **Outer**: independent system chains can run in separate parallel tasks.
 - **Inner**: system logic can be independently performed on entity subsets in separate tasks.



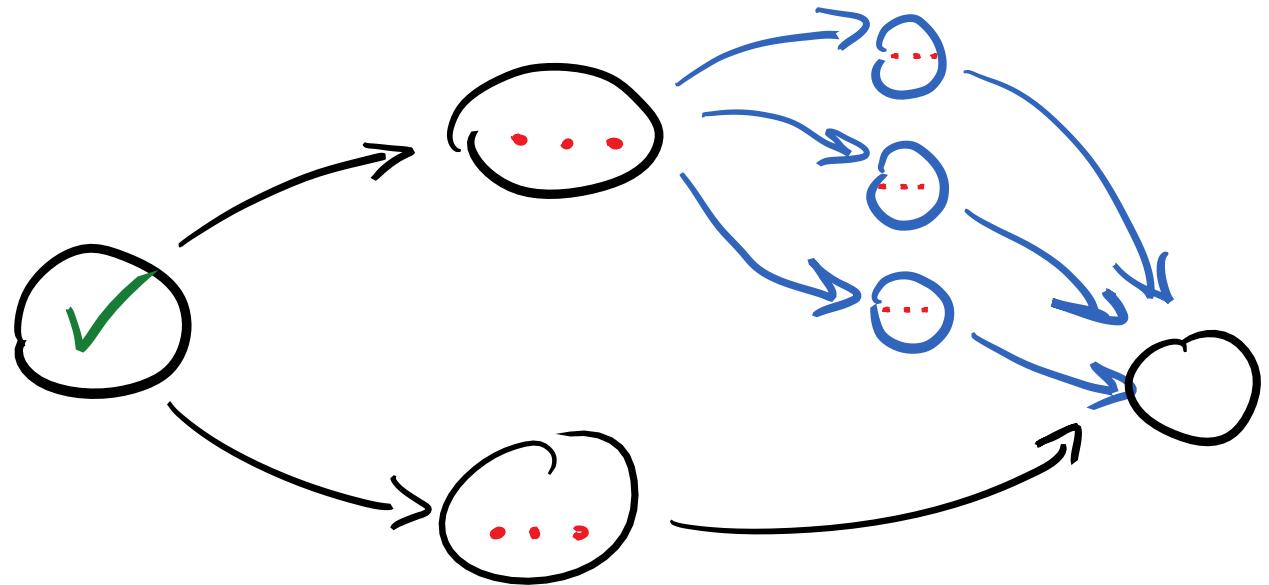
Multithreading support

- Enabled by default, customizable and can be optionally disabled.
- Two levels of parallelism:
 - **Outer**: independent system chains can run in separate parallel tasks.
 - **Inner**: system logic can be independently performed on entity subsets in separate tasks.



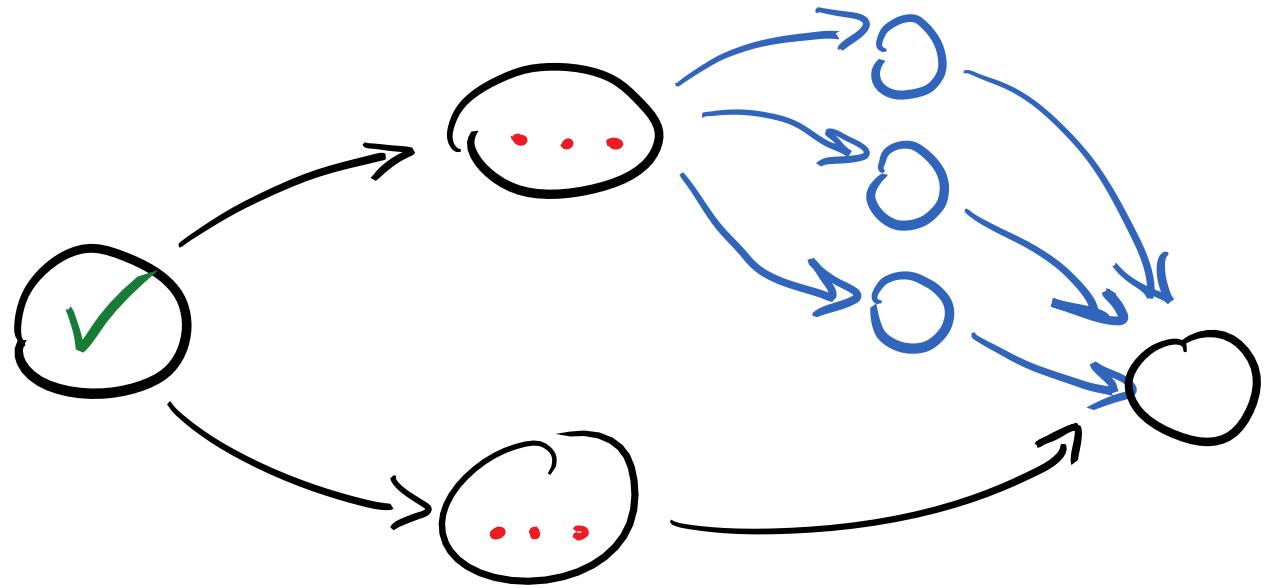
Multithreading support

- Enabled by default, customizable and can be optionally disabled.
- Two levels of parallelism:
 - **Outer**: independent system chains can run in separate parallel tasks.
 - **Inner**: system logic can be independently performed on entity subsets in separate tasks.



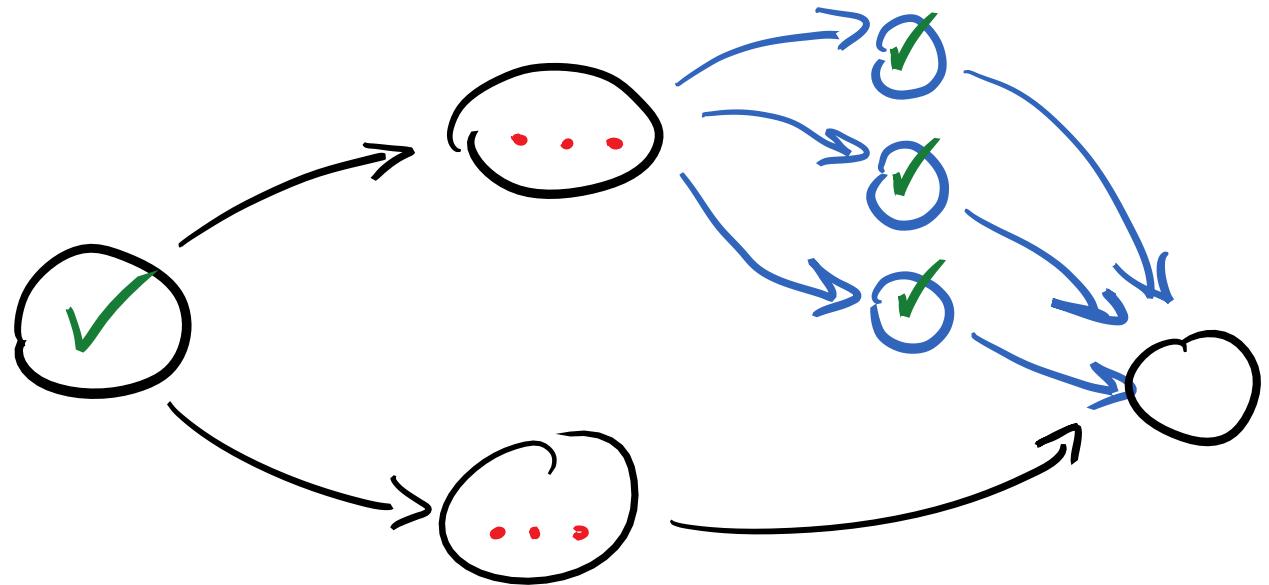
Multithreading support

- Enabled by default, customizable and can be optionally disabled.
- Two levels of parallelism:
 - **Outer**: independent system chains can run in separate parallel tasks.
 - **Inner**: system logic can be independently performed on entity subsets in separate tasks.



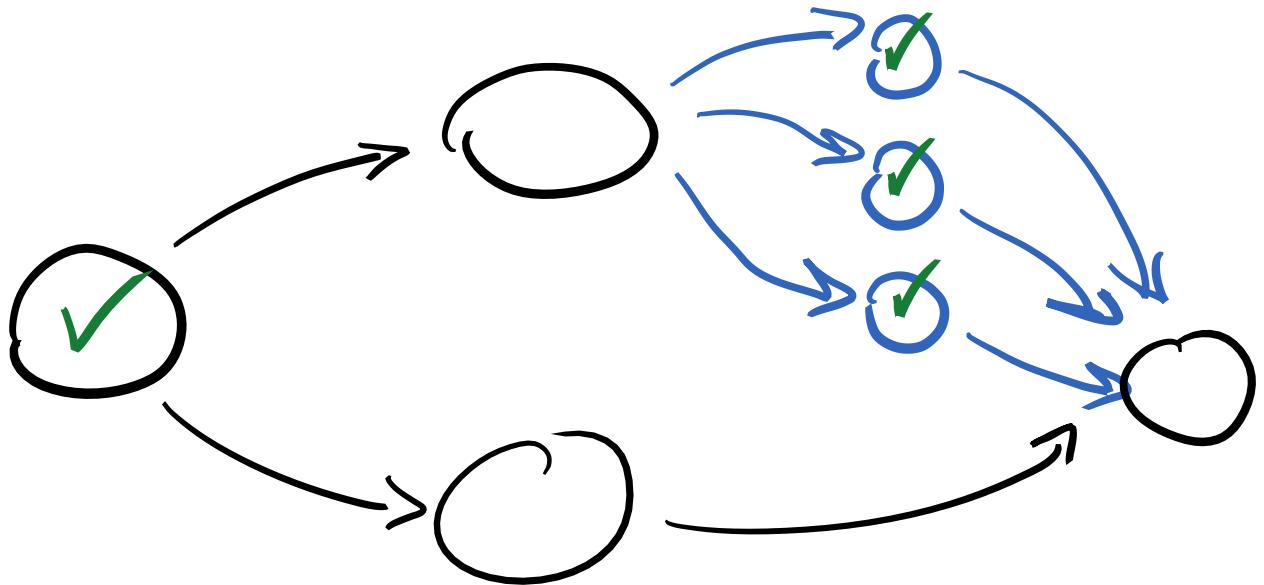
Multithreading support

- Enabled by default, customizable and can be optionally disabled.
- Two levels of parallelism:
 - **Outer**: independent system chains can run in separate parallel tasks.
 - **Inner**: system logic can be independently performed on entity subsets in separate tasks.



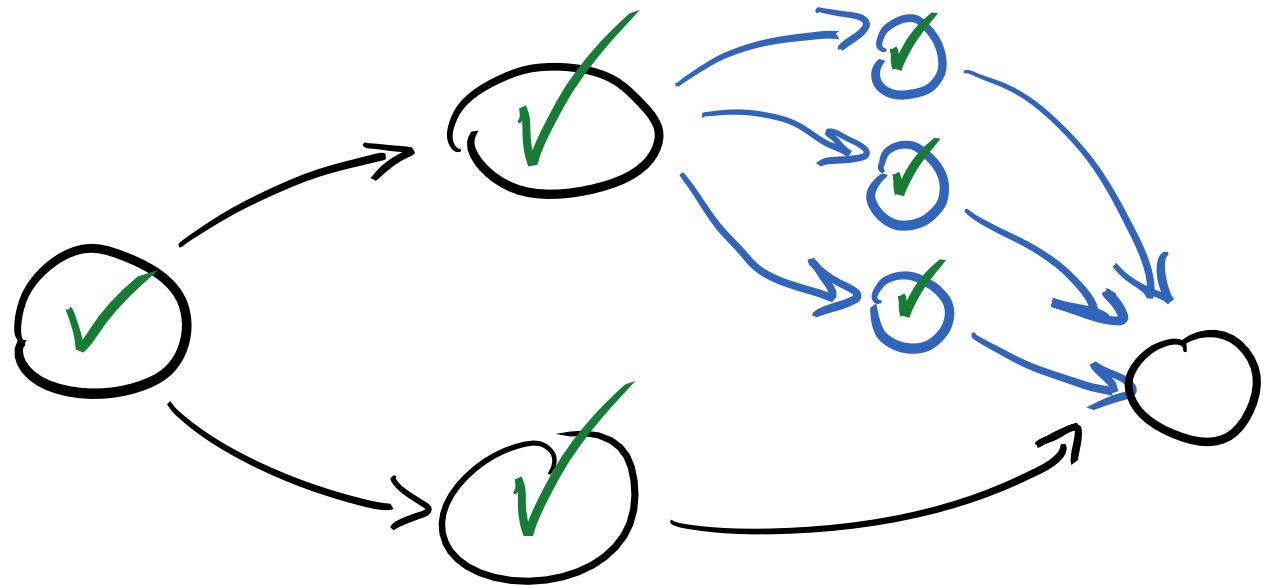
Multithreading support

- Enabled by default, customizable and can be optionally disabled.
- Two levels of parallelism:
 - **Outer**: independent system chains can run in separate parallel tasks.
 - **Inner**: system logic can be independently performed on entity subsets in separate tasks.



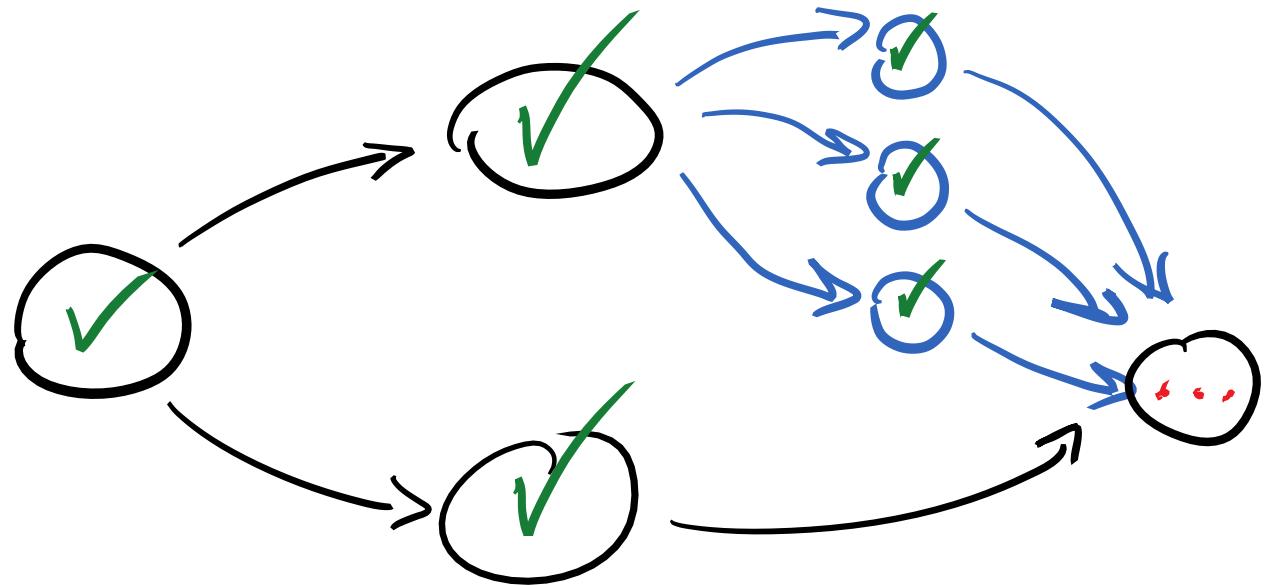
Multithreading support

- Enabled by default, customizable and can be optionally disabled.
- Two levels of parallelism:
 - **Outer**: independent system chains can run in separate parallel tasks.
 - **Inner**: system logic can be independently performed on entity subsets in separate tasks.



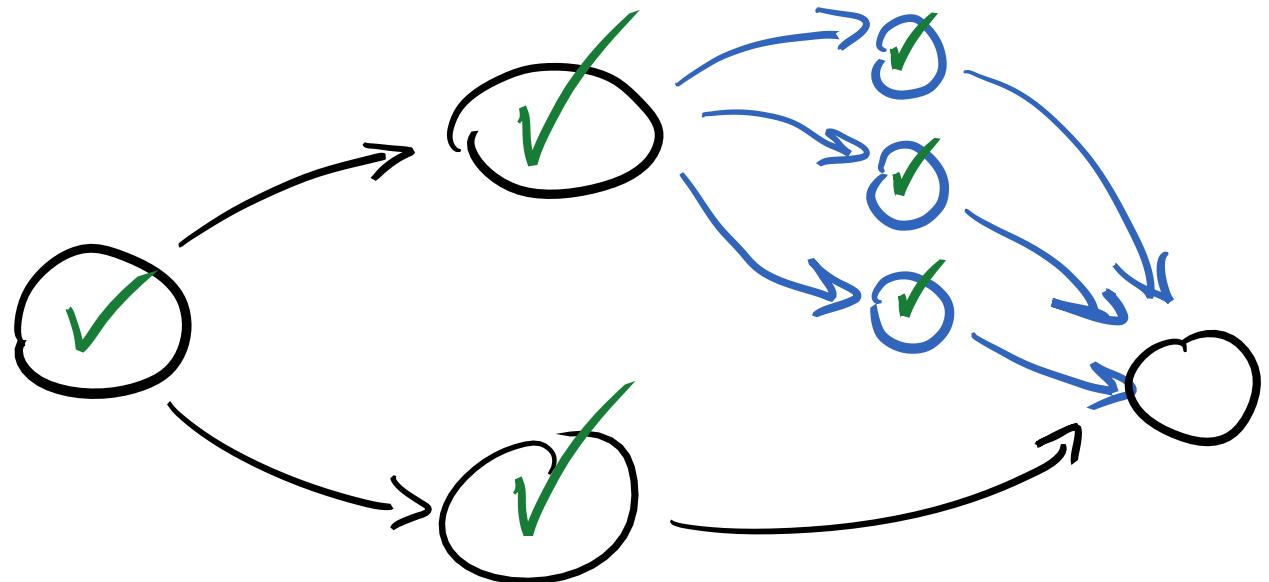
Multithreading support

- Enabled by default, customizable and can be optionally disabled.
- Two levels of parallelism:
 - **Outer**: independent system chains can run in separate parallel tasks.
 - **Inner**: system logic can be independently performed on entity subsets in separate tasks.



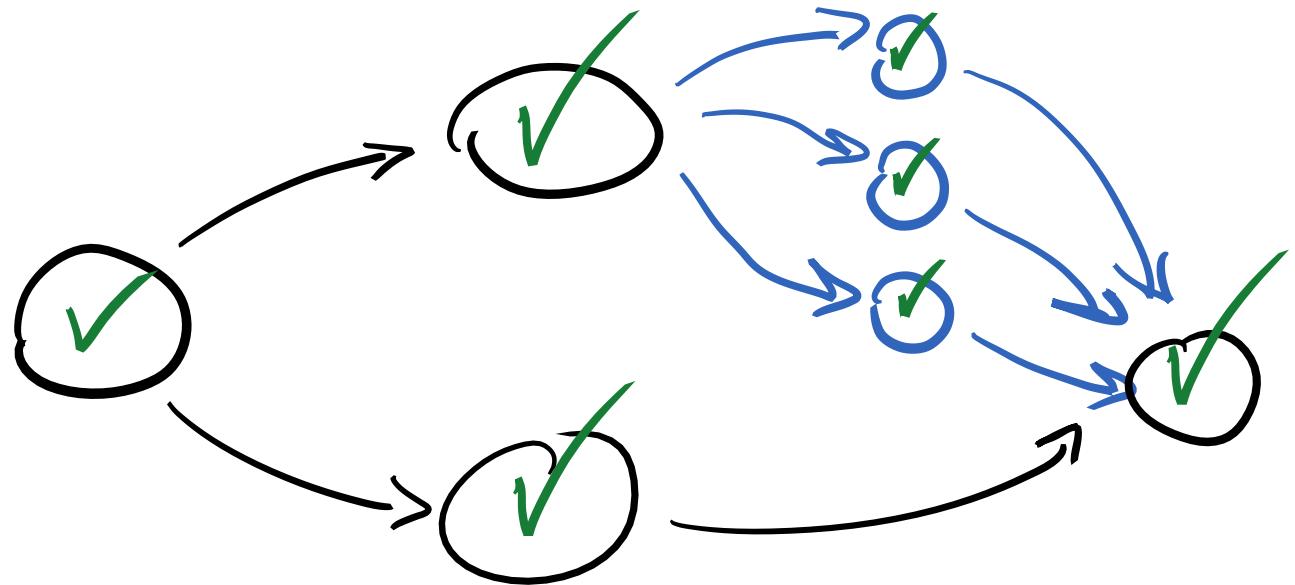
Multithreading support

- Enabled by default, customizable and can be optionally disabled.
- Two levels of parallelism:
 - **Outer**: independent system chains can run in separate parallel tasks.
 - **Inner**: system logic can be independently performed on entity subsets in separate tasks.



Multithreading support

- Enabled by default, customizable and can be optionally disabled.
- Two levels of parallelism:
 - **Outer**: independent system chains can run in separate parallel tasks.
 - **Inner**: system logic can be independently performed on entity subsets in separate tasks.



«Type-value encoding»

- Every value has a **unique type**.
 - Think about `std::integral_constant`.
- Simplifies metaprogramming and compile-time computations.

```
constexpr auto parallelism_policy =
    ipc::none_below_threshold::v(sz_v<10000>,
        ips::split_evenly_fn::v_cores()
    );
```

```
constexpr auto settings =
    ecst::settings::make(
        cs::multithreaded(
            cs::allow_inner_parallelism
        ),
        cs::fixed(sz_v<20000>),
        slc::v<
            c::position, c::velocity, c::acceleration
        >,
        sls::make(
            ssig_acceleration, ssig_velocity
        )
    )  
    cs::scheduler<ss::s_atomic_counter>
);  
  
auto ctx = ecst::context::make(settings);
```



«Type-value encoding»

- Every value has a **unique type**.
 - Think about `std::integral_constant`.
- Simplifies metaprogramming and compile-time computations.

```
constexpr auto parallelism_policy =
    ipc::none_below_threshold::v(sz_v<10000>,
        ips::split_evenly_fn::v_cores()
    );
```

```
constexpr auto settings =
    ecst::settings::make(
        cs::multithreaded(
            cs::allow_inner_parallelism
        ),
        cs::fixed(sz_v<20000>),
        slc::v<
            c::position, c::velocity, c::acceleration
        >,
        sls::make(
            ssig_acceleration, ssig_velocity
        )
    )  
    cs::scheduler<ss::s_atomic_counter>
);

auto ctx = ecst::context::make(settings);
```



«Type-value encoding»

```
/// @brief Given a List of system tags, returns a list of system IDs.  
template <typename TSystemSignatureList, typename TSystemTagList>  
constexpr auto tag_list_to_id_list(  
    TSystemSignatureList ssl, TSystemTagList stl)  
{  
    return mp::list::transform(  
        [ssl](auto x)  
        {  
            return signature_list::system::id_by_tag(ssl, x);  
        },  
        stl);  
}
```

«Type-value encoding»

```
/// @brief Given a List of system tags, returns a list of system IDs.  
template <typename TSystemSignatureList, typename TSystemTagList>  
constexpr auto tag_list_to_id_list(  
    TSystemSignatureList ssl, TSystemTagList stl)  
{  
    return mp::list::transform(  
        [ssl](auto x)  
        {  
            return signature_list::system::id_by_tag(ssl, x);  
        },  
        stl);  
}
```



«Type-value encoding»

```
/// @brief Given a List of system tags, returns a list of system IDs.  
template <typename TSystemSignatureList, typename TSystemTagList>  
constexpr auto tag_list_to_id_list(  
    TSystemSignatureList ssl, TSystemTagList stl) by value  
{  
    return mp::list::transform(  
        [ssl](auto x)  
        {  
            return signature_list::system::id_by_tag(ssl, x);  
        },  
        stl);  
}
```

«Type-value encoding»

```
/// @brief Given a List of system tags, returns a list of system IDs.  
template <typename TSystemSignatureList, typename TSystemTagList>  
constexpr auto tag_list_to_id_list(  
    TSystemSignatureList ssl, TSystemTagList stl) by value  
{  
    return mp::list::transform(  
        [ssl](auto x)  
    {  
        return signature_list::system::id_by_tag(ssl, x);  
    },  
    stl);  
}
```

lambda in compile-time computation

Code example.

Simple particle simulation implemented using ECST.

Architecture of ECST.

“Components” of the entity component system.

CONTEXT

MAIN STORAGE

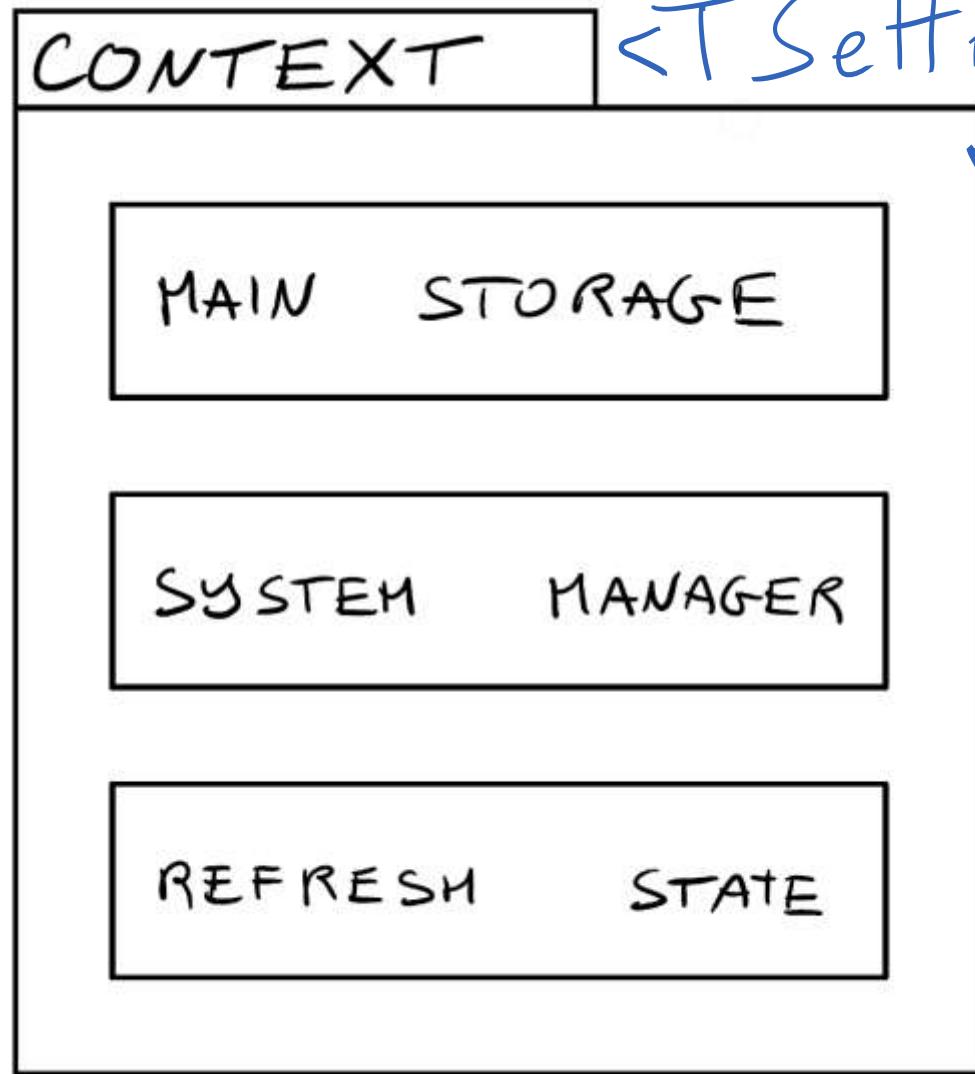
SYSTEM MANAGER

REFRESH STATE

```
constexpr auto component_signatures = setup_components();
constexpr auto system_signatures = setup_systems();
constexpr auto context_settings =
    setup_settings(component_signatures, system_signatures);

auto ctx = ecst::context::make(context_settings);
```

- Movable class that aggregates the main components of **ECST**.

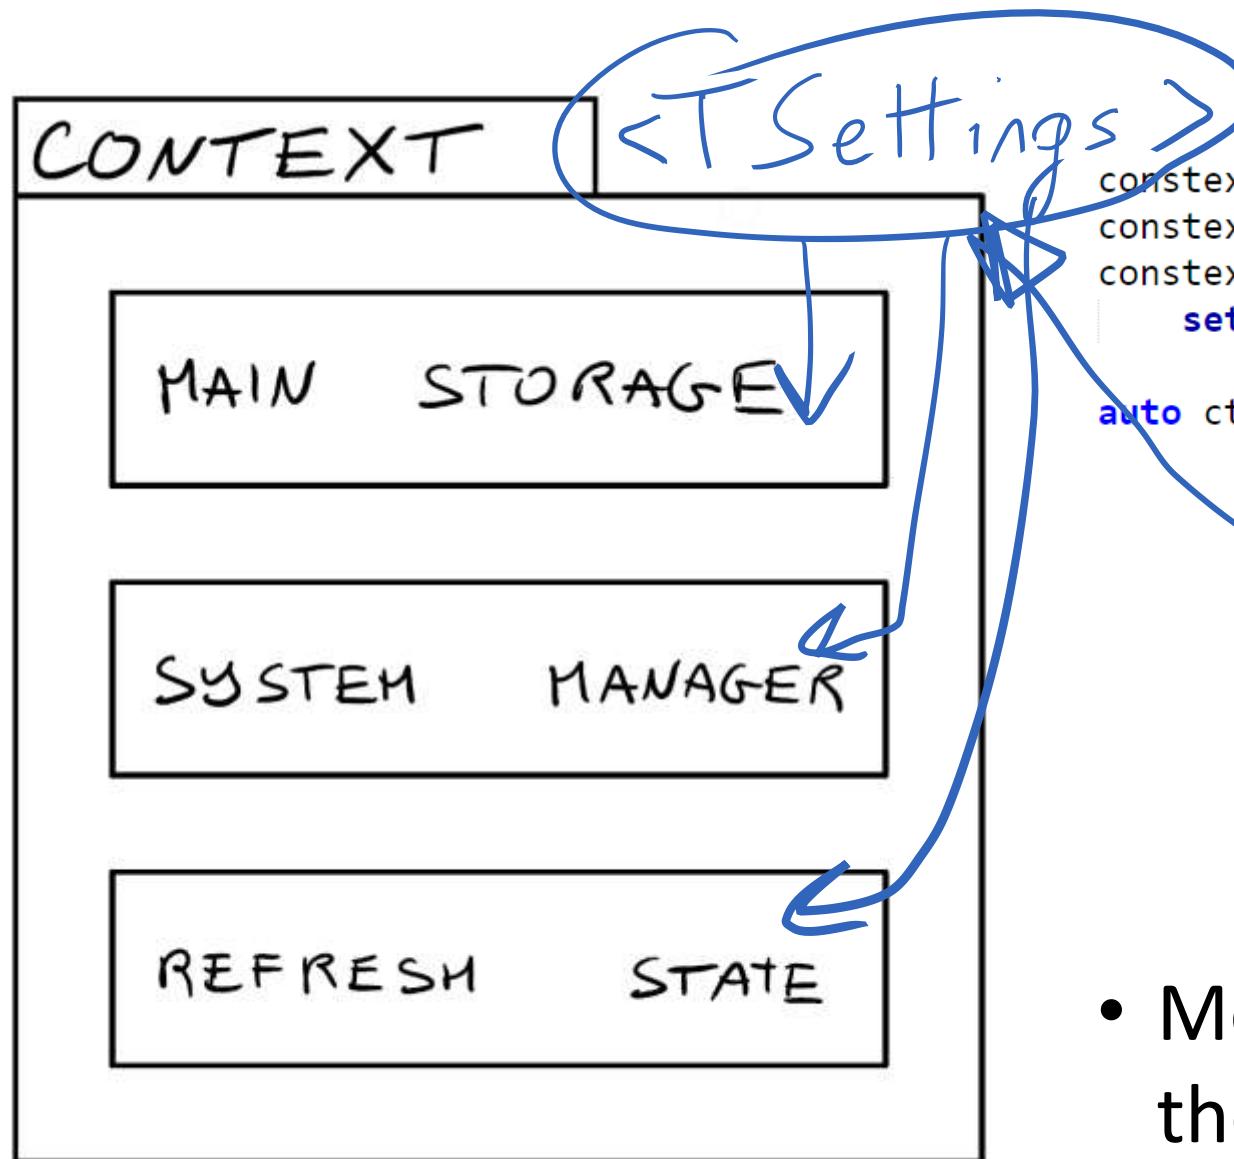


```

constexpr auto component_signatures = setup_components();
constexpr auto system_signatures = setup_systems();
constexpr auto context_settings =
    setup_settings(component_signatures, system_signatures);

auto ctx = ecst::context::make(context_settings);
  
```

- Movable class that aggregates the main components of **ECST**.



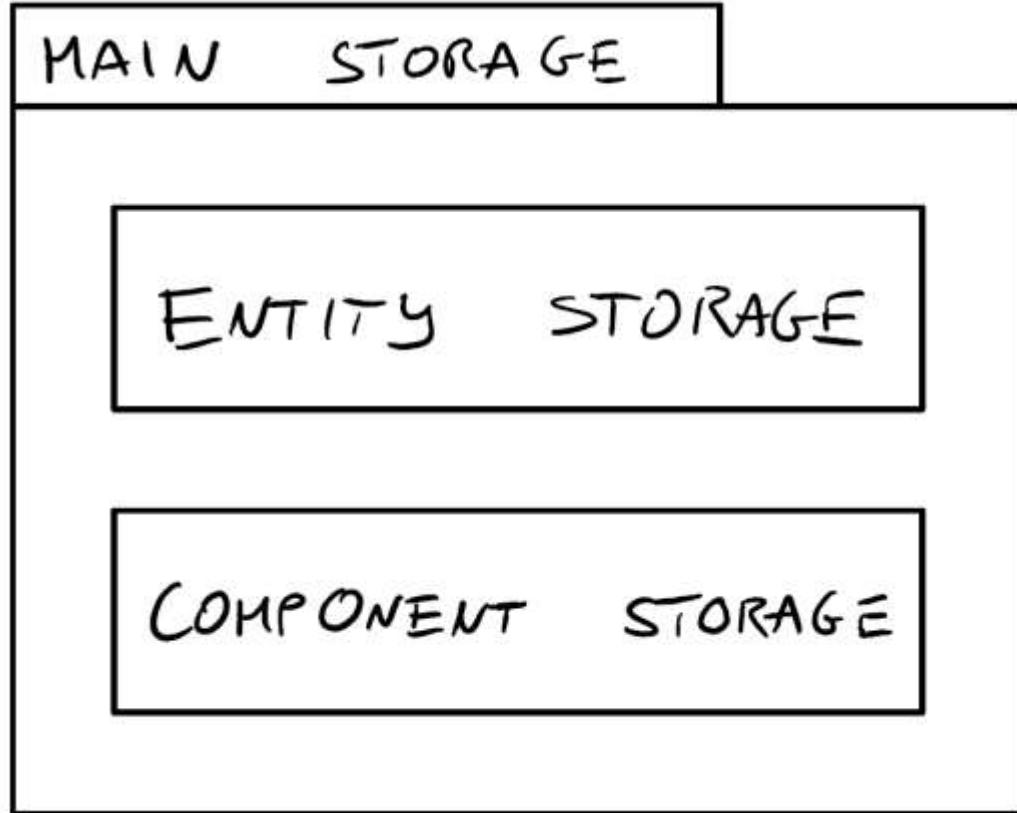
```

constexpr auto component_signatures = setup_components();
constexpr auto system_signatures = setup_systems();
constexpr auto context_settings =
    setup_settings(component_signatures, system_signatures);

auto ctx = ecst::context::make(context_settings);
  
```

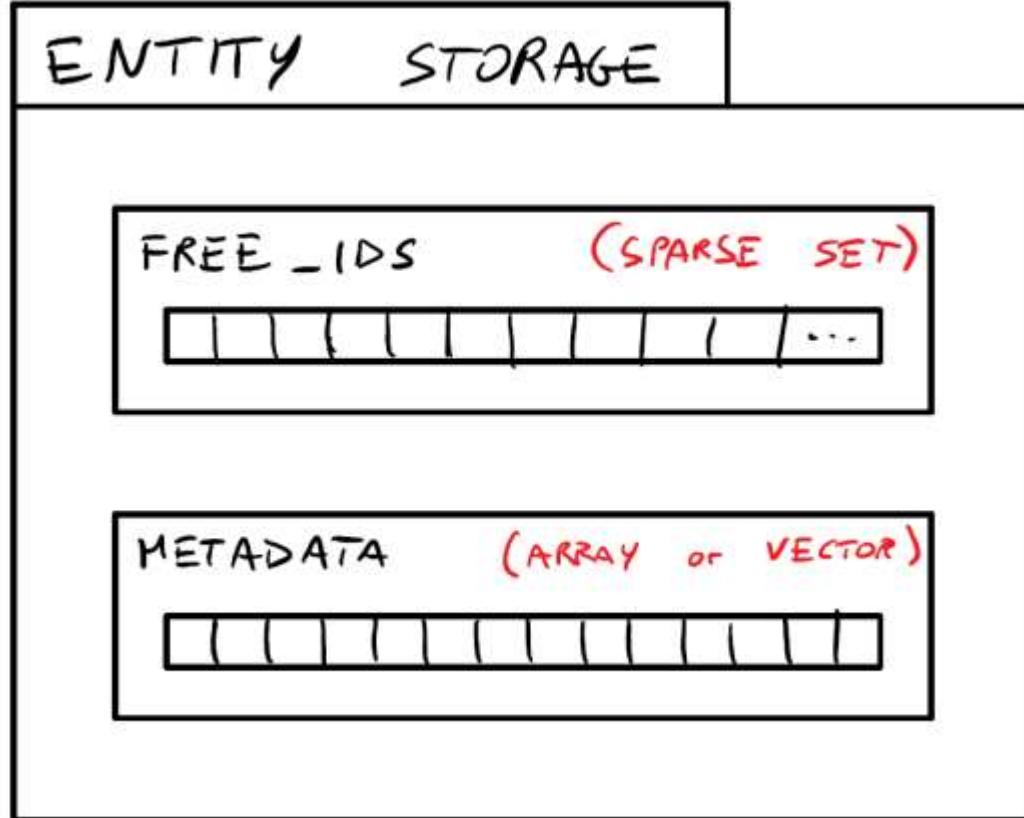
- Movable class that aggregates the main components of **ECST**.

Context -> Main Storage



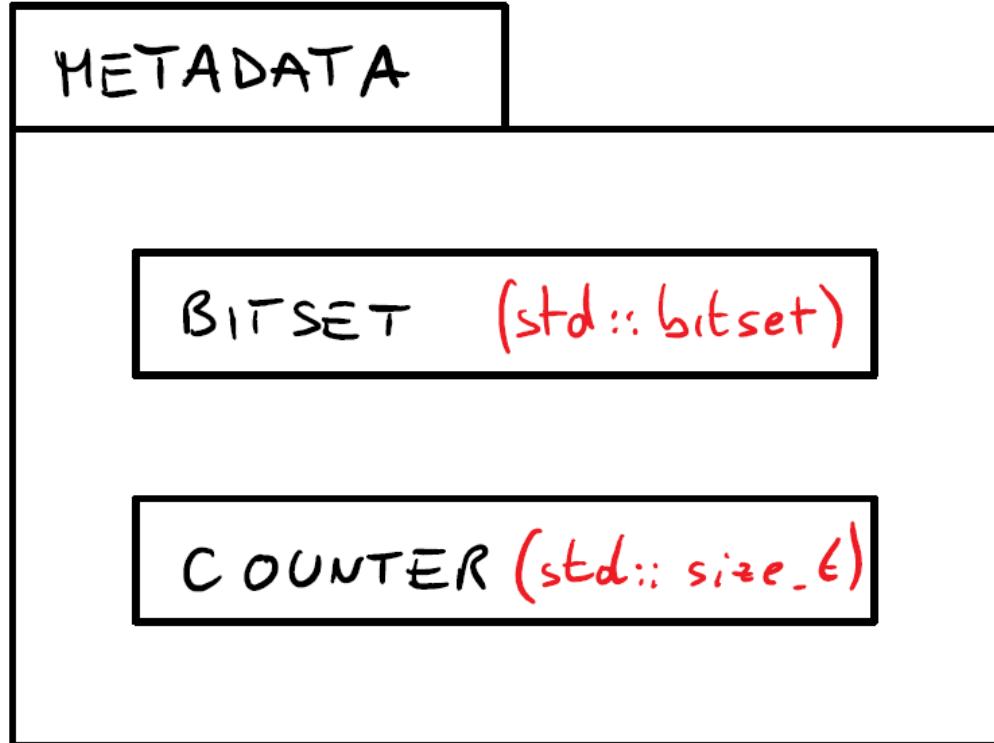
- Encapsulates **entity storage**, which deals with entity IDs and metadata.
- Encapsulates **component storage**, which deals with component data and component storage strategies.

Context -> Main Storage -> Entity Storage



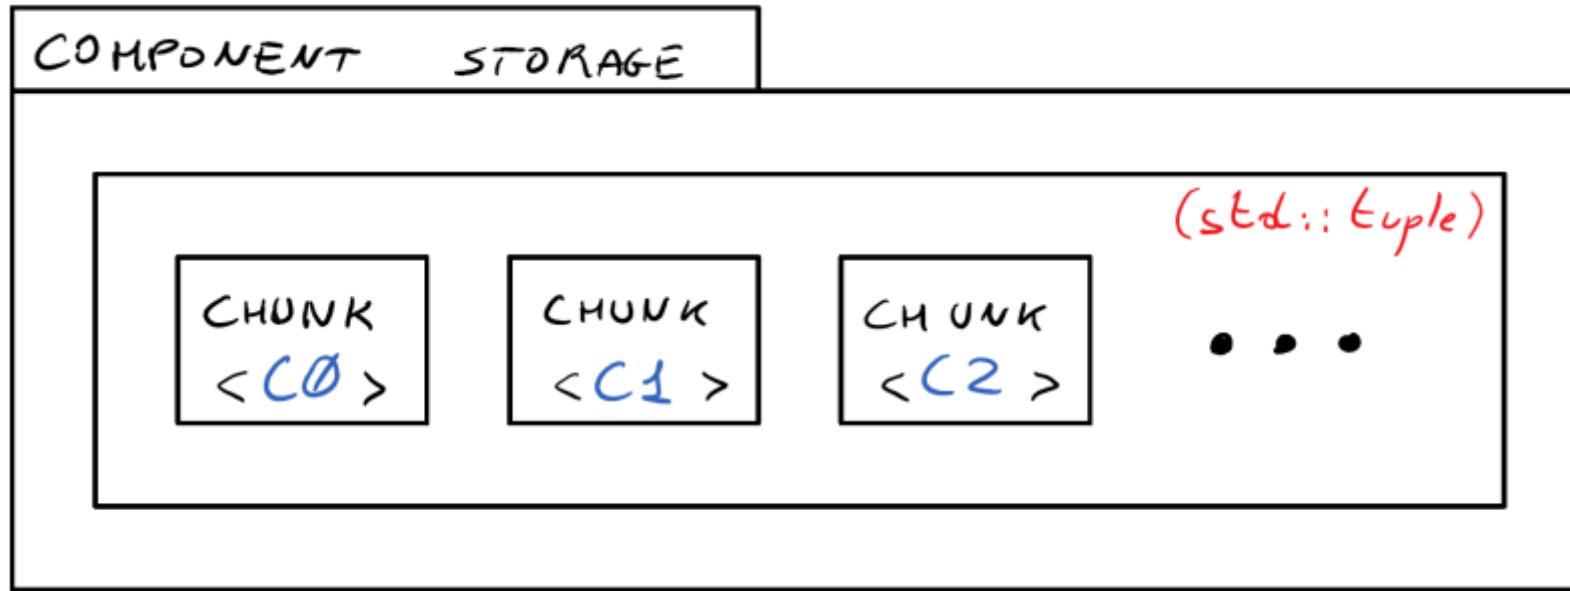
- Keeps track of used and unused **entity IDs**.
- Manages **entity metadata**.
 - Entity metadata may be enriched by component-container-specific metadata at compile-time.

Context -> Main Storage -> Entity Storage -> Metadata



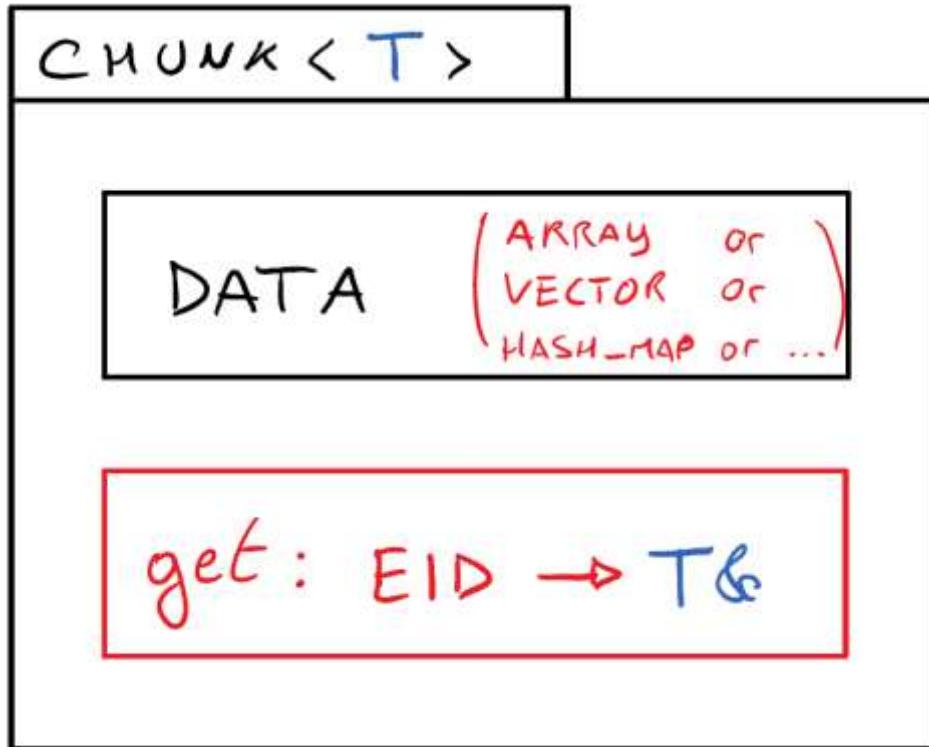
- **Component bitset:** keeps track of current components.
- **Validity counter:** differentiates entity instances which re-use a previous ID. **Handles** check the counter to prevent access to re-used entities.
- Eventual component storage metadata.

Context -> Main Storage -> Component Storage



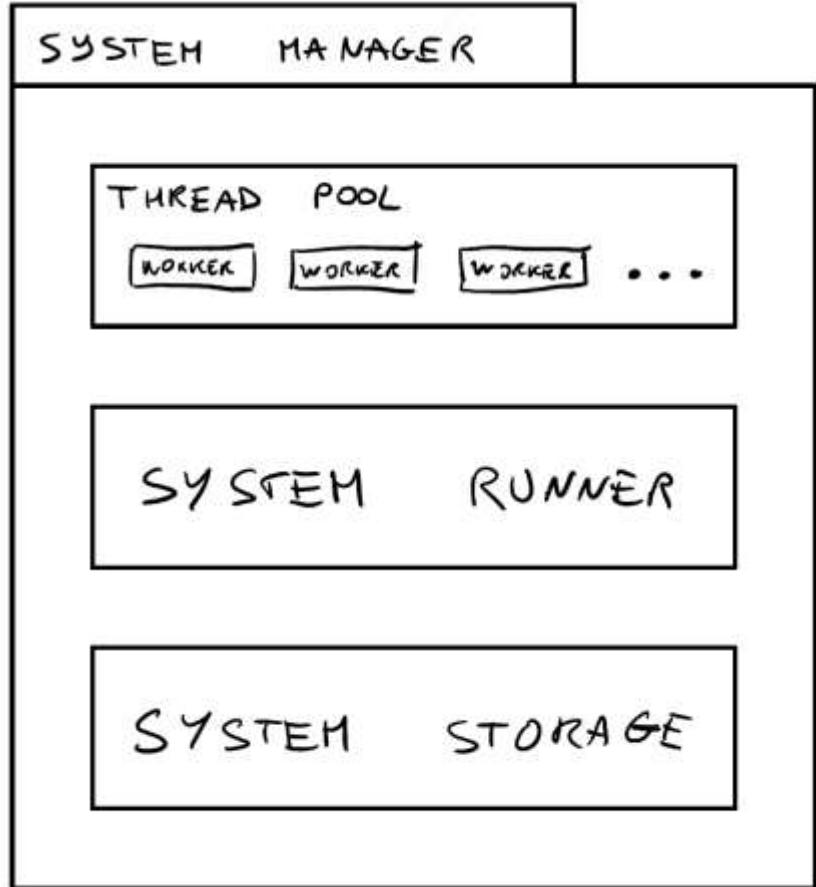
- Stores a **tuple** of component **storage chunks**.
 - Every chunk stores a specific component type in a specific data structure.

Context -> Main Storage -> Component Storage -> Chunk



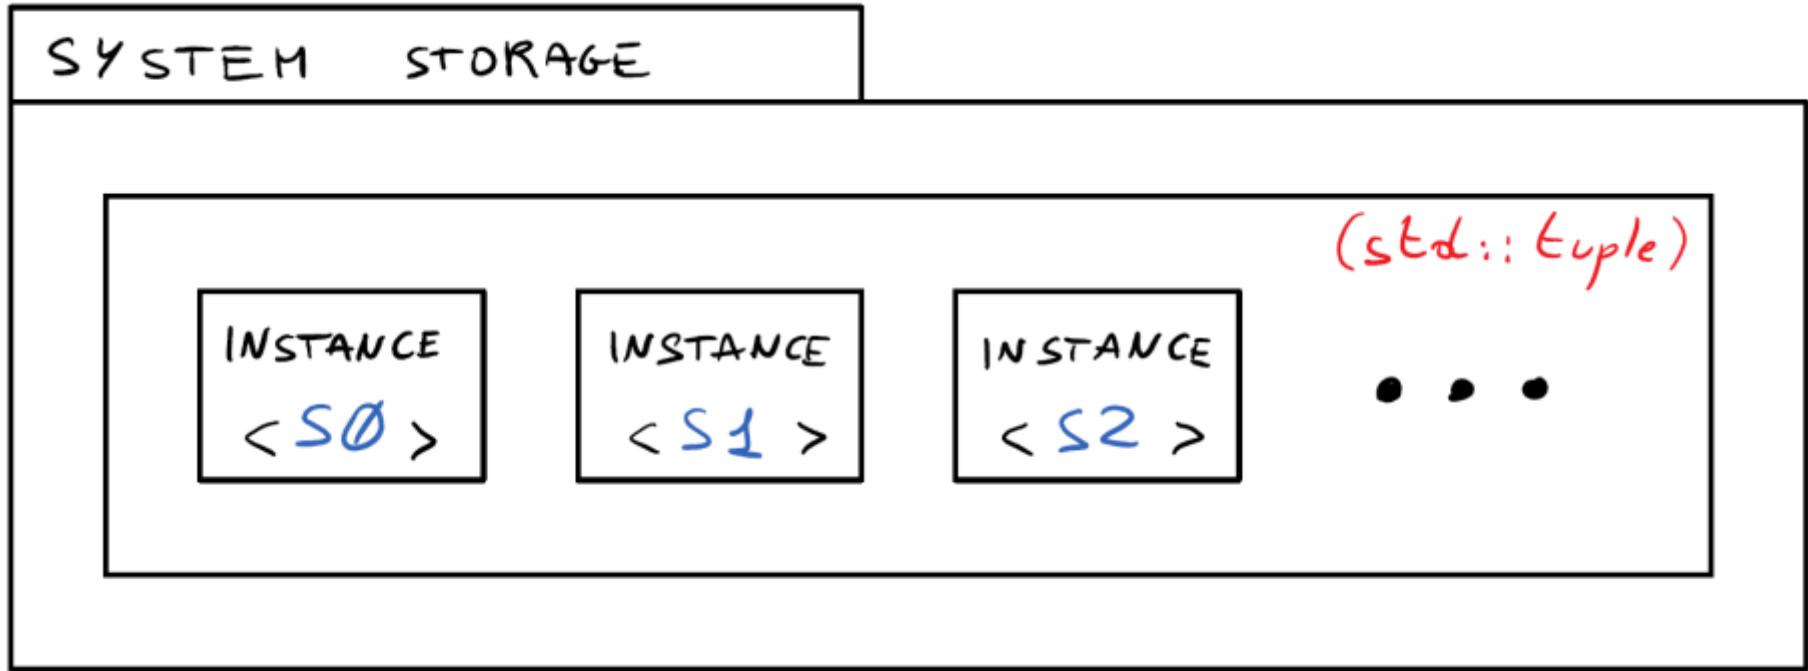
- Manages component data for a particular component type.
- Provides a way to retrieve component data by entity ID.

Context -> System Manager



- Contains a **thread pool**, used to execute system logic.
- Contains a **system runner** which provides a nice interface for a **system scheduler**.
- Stores **system instances** and relative metadata in the **system storage**.

Context -> System Manager -> System Storage



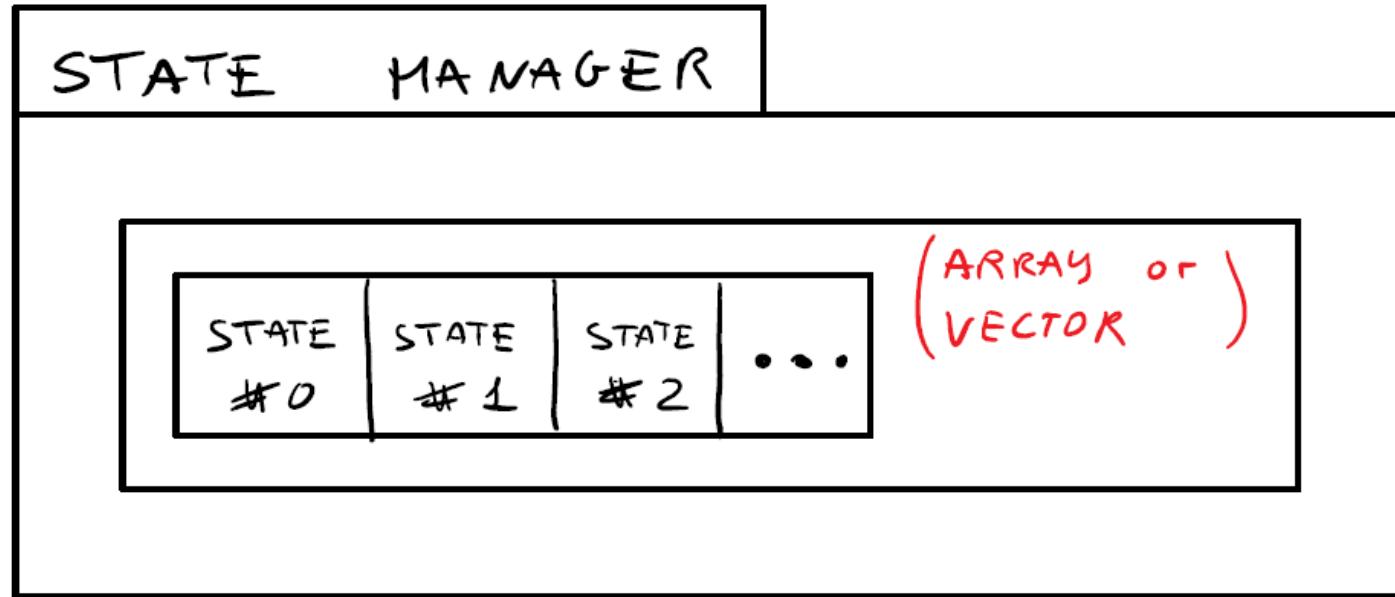
- Stores a **tuple of system instances**, one per system type.

Context -> System Manager -> System Storage -> Instance



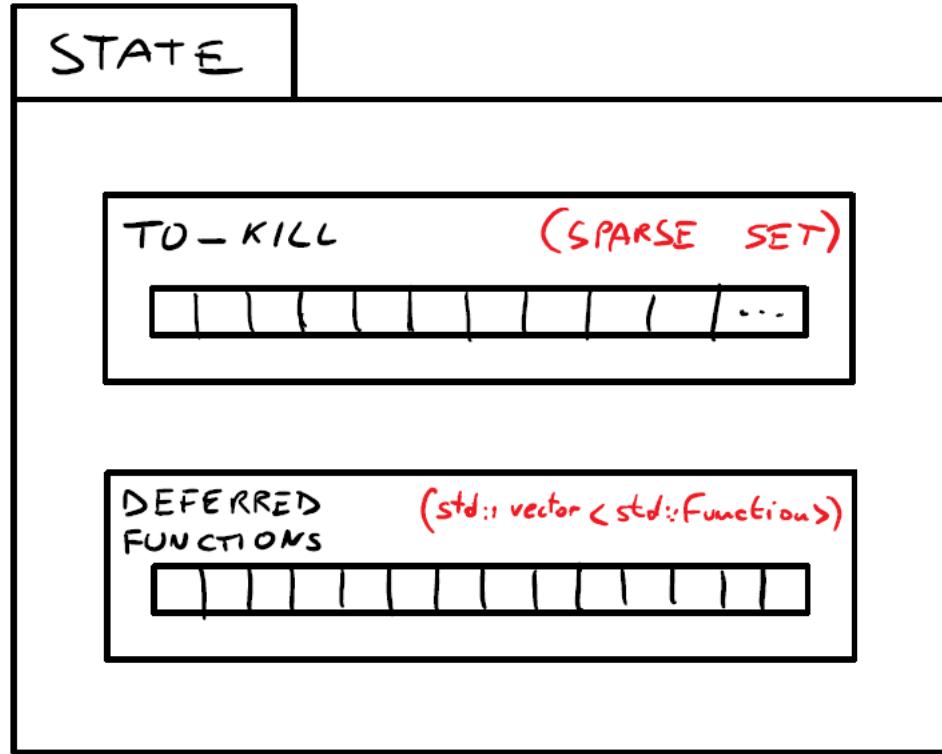
- Contains an instance of the “real” system type.
- Manages N subtasks states through the **state manager**.
- Keeps track of subscribed entities.
- Contains a **inner parallelism executor** instance.
- Contains a **bitset** representing the required components.

Context -> System Manager -> ... -> Instance -> State Manager



- Manages a number of **states**, that can change at run-time.
 - Every state is connected to a **subtask**.
 - Subtasks are executed in separate threads.

Context -> System Manager -> ... -> State Manager -> State



- Contains a set of entity IDs that will be killed during the next **refresh step**.
- Contains a vector of **deferred functions** that will be executed sequentially during the next **refresh step**.
- Optionally contains user-defined **system output data**.

Implementation of ECST.

Challenges, execution flow, multithreading, and metaprogramming.



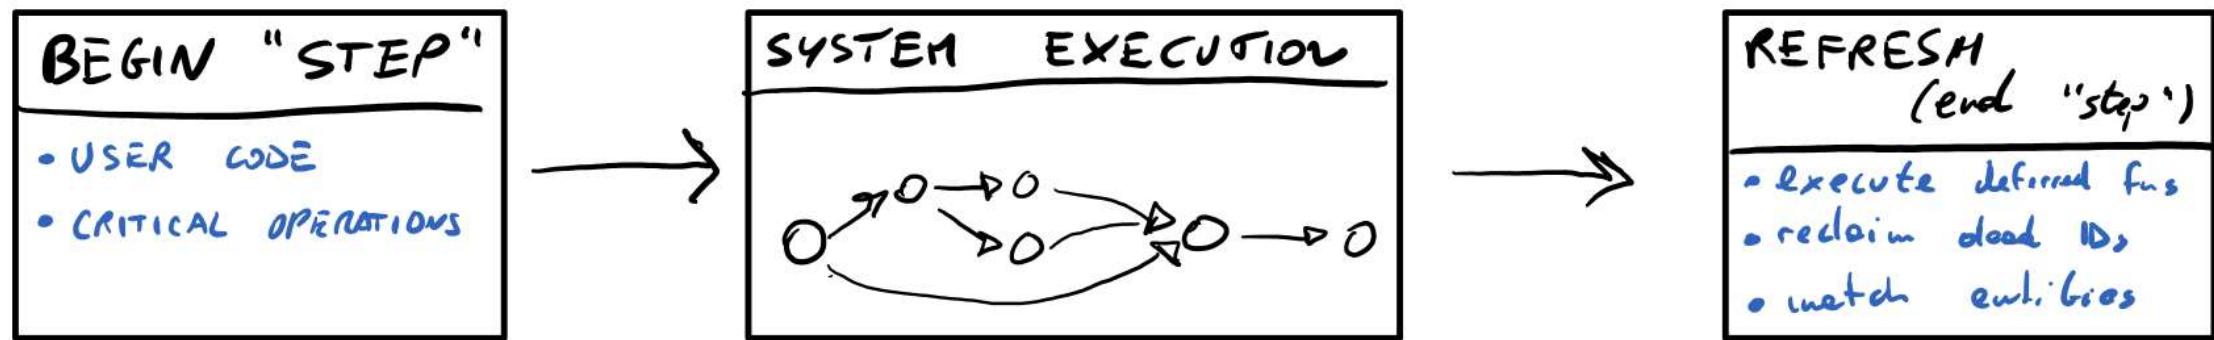
<http://vittorioromeo.info>
vittorio.romeo@outlook.com

<http://github.com/SuperV1234/cppnow2016>

Implementation details overview

- **Execution flow and critical operations.**
- Implementation **challenges**.
- **System processing.**
 - Multithreading details.
 - System scheduling (*outer parallelism*).
 - System inner parallelism.
- Proxies.
- Metaprogramming module examples.

Flow - overview



- System execution is preceded by a **step**, where **critical operations** can occur.
- After system execution, a **refresh** takes place, during which:
 - **Deferred critical operations** are executed.
 - **Dead entity IDs** are **reclaimed**.
 - New and modified entities are **matched to systems**.

Flow - critical operations

- Operations that need to be executed **sequentially** in a **separate final step** are referred to as “**critical**”.
 - Critical operations can be queued and **deferred** to a later step during system logic execution.
- Critical operations:
 - Creating or destroying an entity.
 - Adding or removing a component to/from an entity.
 - Running the system scheduler.
- Non-critical operations:
 - Analyzing/using a dependency’s output system data.
 - Marking an entity as dead.
 - Accessing and mutating existing component data.

```

struct collision_effects
{
    template <typename TData>
    void process(TData& data)
    {
        data.for_entities([&](auto eid)
        {
            const auto& contacts =
                data.get_previous_output<s::collision_detection>();

            if(contacts.was_hit(eid))
            {
                auto& h = ecst::get<c::health>(data, eid);
                h.damage();

                if(h.dead()) data.kill_entity(eid);

                data.defer([&](auto& proxy)
                {
                    auto e = proxy.create_entity();
                    auto& ep = e.add_component(ct::effect_particle);
                    ep.set_effect(effects::explosion);
                });
            }
        });
    };
};

```

```

struct collision_effects
{
    template <typename TData>
    void process(TData& data)
    {
        data.for_entities([&](auto eid)
        {
            const auto& contacts =
                data.get_previous_output<s::collision_detection>();

            if(contacts.was_hit(eid))
            {
                auto& h = ecst::get<c::health>(data, eid);
                h.damage();

                if(h.dead()) data.kill_entity(eid);

                data.defer([&](auto& proxy)
                {
                    auto e = proxy.create_entity();
                    auto& ep = e.add_component(ct::effect_particle);
                    ep.set_effect(effects::explosion);
                });
            }
        });
    };
};

```

non-critical

```

struct collision_effects
{
    template <typename TData>
    void process(TData& data)
    {
        data.for_entities([&](auto eid)
        {
            const auto& contacts =
                data.get_previous_output<s::collision_detection>();

            if(contacts.was_hit(eid))
            {
                auto& h = ecst::get<c::health>(data, eid);
                h.damage();

                if(h.dead()) data.kill_entity(eid);

                data.defer([&](auto& proxy)
                {
                    auto e = proxy.create_entity();
                    auto& ep = e.add_component(ct::effect_particle);
                    ep.set_effect(effects::explosion);
                });
            }
        });
    };
};

```

non-critical

non-critical

```

struct collision_effects
{
    template <typename TData>
    void process(TData& data)
    {
        data.for_entities([&](auto eid)
        {
            const auto& contacts =
                data.get_previous_output<s::collision_detection>();

            if(contacts.was_hit(eid))
            {
                auto& h = ecst::get<c::health>(data, eid);
                h.damage();

                if(h.dead()) data.kill_entity(eid);
            }
            data.defer([&](auto& proxy)
            {
                auto e = proxy.create_entity();
                auto& ep = e.add_component(ct::effect_particle);
                ep.set_effect(effects::explosion);
            });
        });
    });
};

```

non-critical

non-critical

non-crit.col

```

struct collision_effects
{
    template <typename TData>
    void process(TData& data)
    {
        data.for_entities([&](auto eid)
        {
            const auto& contacts =
                data.get_previous_output<s::collision_detection>();

            if(contacts.was_hit(eid))
            {
                auto& h = ecst::get<c::health>(data, eid);
                h.damage();
                if(h.dead()) data.kill_entity(eid);
            }

            data.defer([&](auto& proxy)
            {
                auto e = proxy.create_entity();
                auto& ep = e.add_component(ct::effect_particle);
                ep.set_effect(effects::explosion);
            });
        });
    });
};

```

non-critical

non-critical

non-crit.col

critical

```

struct collision_effects
{
    template <typename TData>
    void process(TData& data)
    {
        data.for_entities([&](auto eid)
        {
            const auto& contacts =
                data.get_previous_output<s::collision_detection>();

            if(contacts.was_hit(eid))
            {
                auto& h = ecst::get<c::health>(data, eid);
                h.damage();
                if(h.dead()) data.kill_entity(eid);
            }

            data.defer([&](auto& proxy)
            {
                auto e = proxy.create_entity();
                auto& ep = e.add_component(ct::effect_particle);
                ep.set_effect(effects::explosion);
            });
        });
    });
};

```

non-critical

non-critical

non-crit.col

critical

Flow – user code

```
context.step([&](auto& proxy)
{
    proxy.system(st::render).prepare();

    proxy.execute_systems_overload(
        [dt](s::physics& s, auto& data){ s.process(dt, data); },
        [](s::render& s, auto& data){ s.process(data); });

    proxy.for_system_outputs(st::render,
        [&window](auto& s, auto& va)
    {
        window.draw(va.data(), va.size(),
                    PrimitiveType::Triangles, RenderStates::Default);
    });
});
```

Flow – user code

```
context.step(&)(auto& proxy)
{
    proxy.system(st::render).prepare();

    proxy.execute_systems_overload(
        [dt](s::physics& s, auto& data){ s.process(dt, data); },
        [](s::render& s, auto& data){ s.process(data); });

    proxy.for_system_outputs(st::render,
        [&window](auto& s, auto& va)
    {
        window.draw(va.data(), va.size(),
                    PrimitiveType::Triangles, RenderStates::Default);
    });
});
```



Flow – user code

```
context.step(&)(auto& proxy)
{
    proxy.system(st::render).prepare();

    proxy.execute_systems_overload(
        [dt](s::physics& s, auto& data){ s.process(dt, data); },
        [](s::render& s, auto& data){ s.process(data); });

    proxy.for_system_outputs(st::render,
        [&window](auto& s, auto& va)
    {
        window.draw(va.data(), va.size(),
                    PrimitiveType::Triangles, RenderStates::Default);
    });
});
```



Flow – user code

• critical operations

```
context.step(&)(auto& proxy)
{
    proxy.system(st::render).prepare();

    proxy.execute_systems_overload(
        [dt](s::physics& s, auto& data){ s.process(dt, data); },
        [](s::render& s, auto& data){ s.process(data); });

    proxy.for_system_outputs(st::render,
        [&window](auto& s, auto& va)
    {
        window.draw(va.data(), va.size(),
                    PrimitiveType::Triangles, RenderStates::Default);
    });
});
```



Flow – user code

```
context.step(&)(auto& proxy)
{
    proxy.system(st::render).prepare();

    proxy.execute_systems_overload(
        [dt](s::physics& s, auto& data){ s.process(dt, data); },
        [](s::render& s, auto& data){ s.process(data); });

    proxy.for_system_outputs(st::render,
        [&window](auto& s, auto& va)
    {
        window.draw(va.data(), va.size(),
                    PrimitiveType::Triangles, RenderStates::Default);
    });
});
```

- critical operations
- system execution



Flow – user code

```
context.step(&)(auto& proxy)
{
    proxy.system(st::render).prepare();

    proxy.execute_systems_overload(
        [dt](s::physics& s, auto& data){ s.process(dt, data); },
        [](s::render& s, auto& data){ s.process(data); });

    proxy.for_system_outputs(st::render,
        [&window](auto& s, auto& va)
    {
        window.draw(va.data(), va.size(),
                    PrimitiveType::Triangles, RenderStates::Default);
    });
}
```

- critical operations
 - system execution
- "user" code

"user"
code



Flow – user code

```
context.step(&)(auto& proxy){  
    proxy.system(st::render).prepare();  
  
    proxy.execute_systems_overload(  
        [dt](s::physics& s, auto& data){ s.process(dt, data); },  
        [](s::render& s, auto& data){ s.process(data); });  
  
    proxy.for_system_outputs(st::render,  
        [&window](auto& s, auto& va)  
    {  
        window.draw(va.data(), va.size(),  
                    PrimitiveType::Triangles, RenderStates::Default);  
    });  
};
```

- critical operations execution
- system "user" code

- system execution

"user" code



Flow – user code

```
context.step(&)(auto& proxy){  
    proxy.system(st::render).prepare();  
  
    proxy.execute_systems_overload(  
        [dt](s::physics& s, auto& data){ s.process(dt, data); },  
        [](s::render& s, auto& data){ s.process(data); });  
  
    proxy.for_system_outputs(st::render,  
        [&window](auto& s, auto& va)  
    {  
        window.draw(va.data(), va.size(),  
                    PrimitiveType::Triangles, RenderStates::Default);  
    });  
};
```

- critical operations execution
- system "user" code

- system execution

"user" code

});
Trigger REFRESH



Challenges

- Efficient management of **entity IDs**.
- Exploiting **compile-time knowledge** to increase performance and safety.
- Executing systems **respecting dependencies** between them and using **parallelism** when possible.
- Processing entities subsets of the same system in different threads.
- Dealing with entity/component addition/removal during system execution.
- Providing a clean and safe interface to the user.

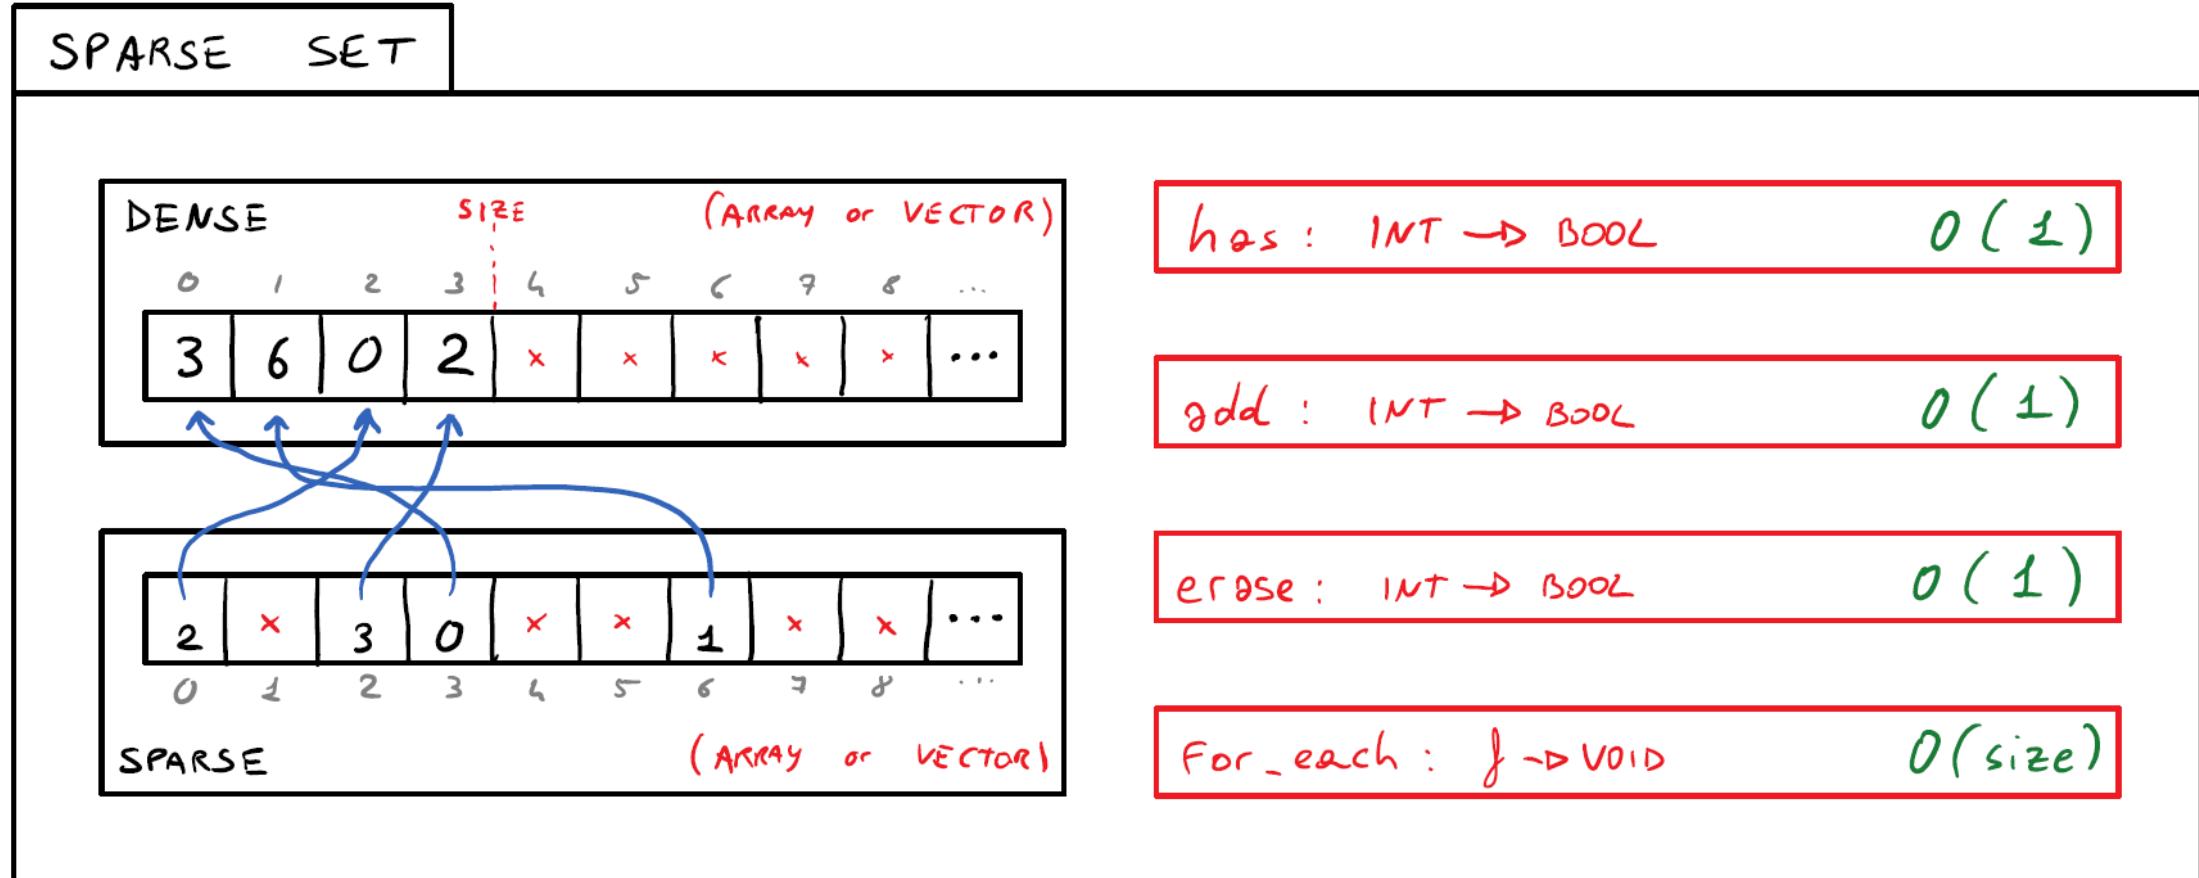
Sparse Integer Sets

- Extremely useful data structure to deal with entity ID management.
- Conceptually represents a set of unsigned integers.
- Allows:
 - **O(1)** test, insertion and removal.
 - **O(k)** iteration, where **k** is the number of integers in the set.
- Used in **system instances** (*subscribed entities*) and in **entity storage** (*available IDs*).

```
/// @brief Sparse integer set, with fixed array storage.  
template <typename T, sz_t TCapacity>  
using fixed_array_sparse_set = impl::base_sparse_set<  
    impl::sparse_set_settings<  
        T,  
        impl::sparse_set_storage::fixed_array<T, TCapacity>  
    >  
>;
```



Sparse Integer Sets



Data structure static dispatching

- Users can choose settings at compile-time that can restrict the functionality of ECST in order to achieve superior performance.
- Example: **max entity limit**.

```
constexpr auto settings =
    ecst::settings::make(
        cs::multithreaded_default,
        cs::fixed(sz_v<20000>),
        component_signature_list,
        system_signature_list,
        cs::scheduler<ss::s_atomic_counter>
    );
```

```
constexpr auto settings =
    ecst::settings::make(
        cs::multithreaded_default,
        cs::dynamic,
        component_signature_list,
        system_signature_list,
        cs::scheduler<ss::s_atomic_counter>
    );
```

Data structure static dispatching

- Users can choose settings at compile-time that can restrict the functionality of ECST in order to achieve superior performance.
- Example: **max entity limit**.

```
constexpr auto settings =
    ecst::settings::make(
        cs::multithreaded_default,
        cs::fixed(sz_v<20000>),
        component_signature_list,
        system_signature_list,
        cs::scheduler<ss::s_atomic_counter>
    );
```

```
constexpr auto settings =
    ecst::settings::make(
        cs::multithreaded_default,
        cs::dynamic,
        component_signature_list,
        system_signature_list,
        cs::scheduler<ss::s_atomic_counter>
    );
```

Data structure static dispatching

```
template <typename TSettings>
auto dispatch_set_type(TSettings s) noexcept
{
    return static_if(settings::has_fixed_entity_storage<TSettings>)
        .then([](auto ts)
        {
            return mp::type_v<fixed_set<settings::fixed_capacity(ts)>>;
        })
        .else_([](auto)
        {
            return mp::type_v<dynamic_set>;
        })(s);
}

template <typename TSettings>
using dispatch_set = mp::unwrap<decltype(dispatch_set_type(TSettings{}))>;
```

Data structure static dispatching

```
template <typename TSettings>
auto dispatch_set_type(TSettings s) noexcept
{
    return static_if(settings::has_fixed_entity_storage<TSettings>)
        .then([](auto ts)
        {
            return mp::type_v<fixed_set<settings::fixed_capacity(ts)>>;
        })
        .else_([](auto)
        {
            return mp::type_v<dynamic_set>;
        })(s);
}

template <typename TSettings>
using dispatch_set = mp::unwrap<decltype(dispatch_set_type(TSettings{}))>;
```

Data structure static dispatching

```
template <typename TSettings>
auto dispatch_set_type(TSettings s) noexcept
{
    return static_if(settings::has_fixed_entity_storage<TSettings>)
        .then([](auto ts)
        {
            return mp::type_v<fixed_set<settings::fixed_capacity(ts)>>;
        })
        .else_([](auto)
        {
            return mp::type_v<dynamic_set>;
        })(s);
}

template <typename TSettings>
using dispatch_set = mp::unwrap<decltype(dispatch_set_type(TSettings{}))>;
```

Data structure static dispatching

```
template <typename TSettings>
auto dispatch(TSettings s)
{
    return static_if(settings::has_fixed_entity_storage<TSettings>)
        .then([](auto ts)
    {
        return mp::type_v<
            impl::fixed_entity_storage<
                context::entity::dispatch<decltype(ts)>, // .
                settings::fixed_capacity(ts) // .
            >;
    })
    .else_([](auto ts)
    {
        return mp::type_v<
            impl::dynamic_entity_storage<TSettings, // .
            context::entity::dispatch<decltype(ts)> // .
        >
    });
})(s);
}

template <typename TSettings>
using dispatch = mp::unwrap<decltype(impl::dispatch(TSettings{}))>;
```

Data structure static dispatching

```
template <typename TSettings>
auto dispatch(TSettings s)
{
    return static_if(settings::has_fixed_entity_storage<TSettings>)
        .then([](auto ts)
    {
        return mp::type_v<
            impl::fixed_entity_storage<
                context::entity::dispatch<decltype(ts)>, // .
                settings::fixed_capacity(ts) // .
            >;
    })
    .else_([](auto ts)
    {
        return mp::type_v<
            impl::dynamic_entity_storage<TSettings, // .
            context::entity::dispatch<decltype(ts)> // .
        >
    });
}(s);
}

template <typename TSettings>
using dispatch = mp::unwrap<decltype(impl::dispatch(TSettings{}))>;
```

up
or
growth
allocations



Data structure static dispatching

```
template <typename TSettings>
auto dispatch(TSettings s)
{
    return static_if(settings::has_fixed_entity_storage<TSettings>)
        .then([](auto ts)
    {
        return mp::type_v<
            impl::fixed_entity_storage<
                context::entity::dispatch<decltype(ts)>,
                settings::fixed_capacity(ts)
            >;
    })
    .else_([](auto ts)
    {
        return mp::type_v<
            impl::dynamic_entity_storage<TSettings,
                context::entity::dispatch<decltype(ts)>
            >;
    });
}

template <typename TSettings>
using dispatch = mp::unwrap<decltype(impl::dispatch(TSettings{}))>;
```

up or growth
allocations

branches and
re-allocations

System settings and dependencies

```
constexpr auto ssig_acceleration =
    ss::make<s::acceleration>(
        divide_evenly(sz_v<4>),
        ss::no_dependencies,
        ss::component_use(
            ss::mutate<c::velocity>,
            ss::read<c::acceleration>
        )
    );
    constexpr auto ssig_movement =
        ss::make<s::velocity>(
            divide_evenly(sz_v<4>),
            ss::depends_on<s::acceleration>,
            ss::component_use(
                ss::mutate<c::position>,
                ss::read<c::velocity>
            )
        );
    constexpr auto ssig_render =
        ss::make<s::render>(
            no_parallelism,
            ss::depends_on<s::velocity>,
            ss::component_use(
                ss::read<c::position>,
                ss::read<c::sprite>
            )
        );

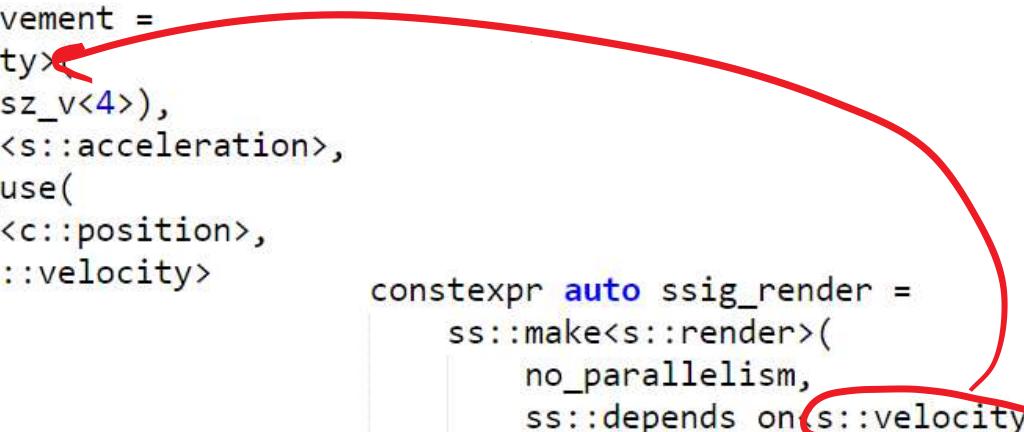
```



System settings and dependencies

```
constexpr auto ssig_acceleration =
    ss::make<s::acceleration>(
        divide_evenly(sz_v<4>),
        ss::no_dependencies,
        ss::component_use(
            ss::mutate<c::velocity>,
            ss::read<c::acceleration>
        )
    );
    constexpr auto ssig_movement =
        ss::make<s::velocity>(
            divide_evenly(sz_v<4>),
            ss::depends_on<s::acceleration>,
            ss::component_use(
                ss::mutate<c::position>,
                ss::read<c::velocity>
            )
        );
    constexpr auto ssig_render =
        ss::make<s::render>(
            no_parallelism,
            ss::depends_on<s::velocity>,
            ss::component_use(
                ss::read<c::position>,
                ss::read<c::sprite>
            )
        );

```



System settings and dependencies

```
constexpr auto ssig_acceleration =  
    ss::make<s::acceleration>(  
        divide_evenly(sz_v<4>),  
        ss::no_dependencies,  
        ss::component_use(  
            ss::mutate<c::velocity>,  
            ss::read<c::acceleration>  
        )  
    );  
  
constexpr auto ssig_movement =  
    ss::make<s::velocity>(  
        divide_evenly(sz_v<4>),  
        ss::depends_on<s::acceleration>,  
        ss::component_use(  
            ss::mutate<c::position>,  
            ss::read<c::velocity>  
        )  
    );  
  
constexpr auto ssig_render =  
    ss::make<s::render>(  
        no_parallelism,  
        ss::depends_on<s::velocity>,  
        ss::component_use(  
            ss::read<c::position>,  
            ss::read<c::sprite>  
        )  
    );
```



System settings and dependencies

```
constexpr auto ssig_acceleration =  
    ss::make<s::acceleration>(  
        divide_evenly(sz_v<4>),  
        ss::no_dependencies,  
        ss::component_use(  
            ss::mutate<c::velocity>,  
            ss::read<c::acceleration>  
        )  
    );  
  
constexpr auto ssig_movement =  
    ss::make<s::velocity>(  
        divide_evenly(sz_v<4>),  
        ss::depends_on<s::acceleration>,  
        ss::component_use(  
            ss::mutate<c::position>,  
            ss::read<c::velocity>  
        )  
    );  
  
constexpr auto ssig_render =  
    ss::make<s::render>(  
        no_parallelism,  
        ss::depends_on<s::velocity>,  
        ss::component_use(  
            ss::read<c::position>,  
            ss::read<c::sprite>  
        )  
    );
```

System settings and dependencies

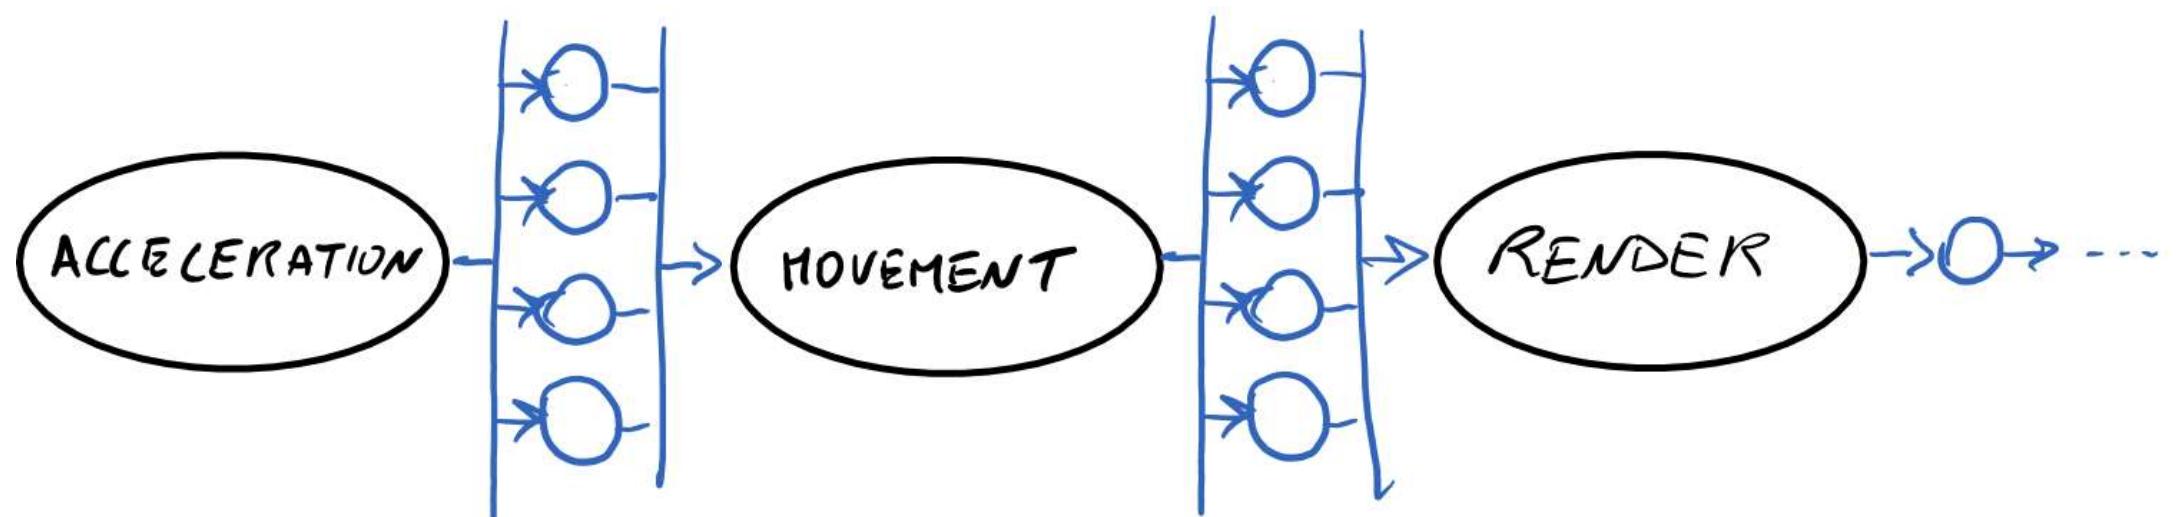
```
constexpr auto ssig_acceleration =  
    ss::make<s::acceleration>(  
        divide_evenly(sz_v<4>),  
        ss::no_dependencies,  
        ss::component_use(  
            ss::mutate<c::velocity>,  
            ss::read<c::acceleration>  
        )  
    );  
  
constexpr auto ssig_movement =  
    ss::make<s::velocity>(  
        divide_evenly(sz_v<4>),  
        ss::depends_on<s::acceleration>,  
        ss::component_use(  
            ss::mutate<c::position>,  
            ss::read<c::velocity>  
        )  
    );  
  
constexpr auto ssig_render =  
    ss::make<s::render>(  
        no_parallelism,  
        ss::depends_on<s::velocity>,  
        ss::component_use(  
            ss::read<c::position>,  
            ss::read<c::sprite>  
        )  
    );
```

System settings and dependencies

```
constexpr auto ssig_acceleration =  
    ss::make<s::acceleration>(  
        divide_evenly(sz_v<4>),  
        ss::no_dependencies,  
        ss::component_use(  
            ss::mutate<c::velocity>,  
            ss::read<c::acceleration>  
        )  
    );  
  
constexpr auto ssig_movement =  
    ss::make<s::velocity>(  
        divide_evenly(sz_v<4>),  
        ss::depends_on<s::acceleration>,  
        ss::component_use(  
            ss::mutate<c::position>,  
            ss::read<c::velocity>  
        )  
    );  
  
constexpr auto ssig_render =  
    ss::make<s::render>(  
        no_parallelism,  
        ss::depends_on<s::velocity>,  
        ss::component_use(  
            ss::read<c::position>,  
            ss::read<c::sprite>  
        )  
    );
```

System settings and dependencies

- An implicit **directed acyclic graph** is built thanks to the compile-time system signatures.



System settings and dependencies

```
struct acceleration
{
    template <typename TData>
    void process(ft dt, TData& data)
    {
        data.for_entities([&](auto eid)
        {
            auto& v = ecst::get<c::velocity>(data, eid);
            const auto& a = ecst::get<c::acceleration>(data, eid);
            actions::accelerate(dt, v, a);
        });
    }
};
```



System settings and dependencies

```
struct acceleration
{
    template <typename TData>
    void process(ft dt, TData& data)
    {
        data.for_entities([&](auto eid)
        {
            auto& v = ecst::get<c::velocity>(data, eid);
            const auto& a = ecst::get<c::acceleration>(data, eid);
            actions::accelerate(dt, v, a);
        });
    }
};
```

.



System settings and dependencies

```
struct acceleration
{
    template <typename TData>
    void process(ft dt, TData& data)
    {
        data.for_entities([&](auto eid)
        {
            auto& v = ecst::get<c::velocity>(data, eid);
            const auto& a = ecst::get<c::acceleration>(data, eid);
            actions::accelerate(dt, v, a);
        });
    }
};
```



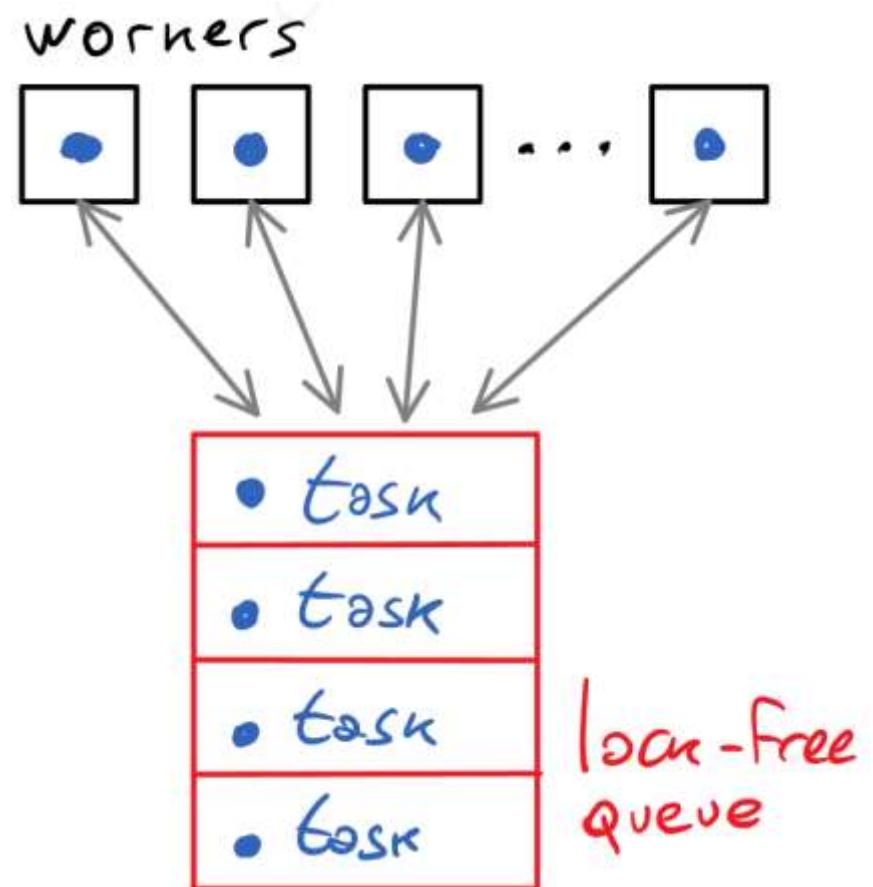
System settings and dependencies

```
struct acceleration
{
    template <typename TData>
    void process(ft dt, TData& data)
    {
        data.for_entities([&](auto eid)
        {
            auto& v = ecst::get<c::velocity>(data, eid);
            const auto& a = ecst::get<c::acceleration>(data, eid);
            actions::accelerate(dt, v, a);
        });
    }
};
```

statically asserted

Multithreading details

- Avoid **busy waiting**, use **condition variables**.
- Use **thread pooling** with a fast **lock-free concurrent queue**.
- Avoid **futures**, use **atomic counters** with **condition variables** to wait for tasks to complete.



Multithreading details – lock-free queue

- `modycamel::BlockingConcurrentQueue`
 - By **Cameron Desrochers** (*cameron314* on GitHub).
 - <https://github.com/cameron314/concurrentqueue> (*Simplified BSD License*)

```
template <typename TF>
void pool::post(TF&& f)
{
    _queue->enqueue(std::move(f));
}

void worker::run()
{
    task t;

    while(_running)
    {
        // Dequeue (blocking) and execute.
        _queue->wait_dequeue(_queue.ctok(), t);
        t();
    }

    _finished = true;
}
```

Multithreading details – synchronization

- **Condition variables** and **atomic counters** are used to wait for multiple tasks to complete.
 - Used both in **system scheduling** and in **inner parallelism**.
- The **counter_blocker** struct aggregates a **mutex**, a **condition variable** and an **atomic counter**.

```
{  
    counter_blocker cb{n};  
  
    execute_and_wait_until_counter_zero(cb, []  
    {  
        run_tasks(n);  
    });  
}
```



Multithreading details – synchronization

```
/// @brief Data structure containing synchronization primitives required for
/// waiting over an atomic counter with an `std::condition_variable` and
/// `std::mutex`.
struct counter_blocker
{
    std::condition_variable _cv;
    std::mutex _mutex;
    std::atomic<sz_t> _counter;

    counter_blocker(sz_t initial_count) noexcept : _counter{initial_count}
    {
    }
};
```

- `counter_blocker` is essentially a simple abstraction to prevent busy waiting.
- Clean «interface» functions that take it as a parameter are defined for convenience and safety.

Multithreading details – synchronization

```
/// @brief Accesses `cv` and `c` through a `lock_guard` on `mutex`, and
/// calls `f(cv, c)`.

template <typename TF>
void access_cv_counter(
    mutex_type& mutex, cv_type& cv, counter_type& c, TF&& f) noexcept
{
    std::lock_guard<std::mutex> l(mutex);
    f(cv, c);

    // Prevent warnings.
    (void)l;
}
```

- Access to the counter is synchronized thanks to a `lock_guard` that locks the `mutex` contained in the counter blocker.
- The passed function will be executed after locking, to prevent mistakes and explicit locking.

Multithreading details – synchronization

```
/// @brief Decrements `c` through `mutex`, and calls `cv.notify_one()`.  
void decrement_cv_counter_and_notify_one(  
    mutex_type& mutex, cv_type& cv, counter_type& c) noexcept  
{  
    access_cv_counter(mutex, cv, c, [](auto& x_cv, auto& x_c)  
    {  
        ECST_ASSERT(x_c > 0);  
        --x_c;  
        x_cv.notify_one();  
    });  
}
```

- This is how the counter is decremented and the waiting threads are notified.

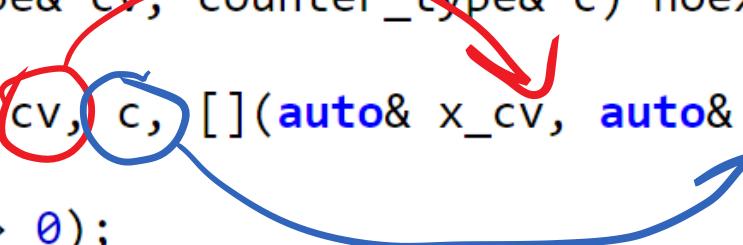
Multithreading details – synchronization

```
/// @brief Decrements `c` through `mutex`, and calls `cv.notify_one()`.  
void decrement_cv_counter_and_notify_one(  
    mutex_type& mutex, cv_type& cv, counter_type& c) noexcept  
{  
    access_cv_counter(mutex, cv, c, [](auto& x_cv, auto& x_c)  
    {  
        ECST_ASSERT(x_c > 0);  
        --x_c;  
        x_cv.notify_one();  
    });  
}
```

- This is how the counter is decremented and the waiting threads are notified.

Multithreading details – synchronization

```
/// @brief Decrements `c` through `mutex`, and calls `cv.notify_one()`.  
void decrement_cv_counter_and_notify_one(  
    mutex_type& mutex, cv_type& cv, counter_type& c) noexcept  
{  
    access_cv_counter(mutex, cv, c, [](auto& x_cv, auto& x_c)  
    {  
        ECST_ASSERT(x_c > 0);  
        --x_c;  
        x_cv.notify_one();  
    });  
}
```



- This is how the counter is decremented and the waiting threads are notified.

Multithreading details – synchronization

```
/// @brief Locks `mutex`, executes `f` and waits until `predicate`  
/// through `cv`.  
template <typename TPredicate, typename TF>  
void execute_and_wait_until(mutex_type& mutex, cv_type& cv,  
    TPredicate&& predicate, TF&& f) noexcept  
{  
    std::unique_lock<std::mutex> l(mutex);  
    f();  
    cv.wait(l, FWD(predicate));  
}
```

- The above function is used to wait for a predicate to become true.
- As previously seen, a function is passed to avoid mistakes, explicit locking and explicit waiting.

Multithreading details – synchronization

- The previously shown functions have “interface versions” that take a counter blocker as a parameter.

```
/// @brief Decrements `cb`'s atomic counter by one, and calls `notify_one`  
/// on the inner condition variable.  
void decrement_cv_counter_and_notify_one(counter_blocker & cb) noexcept  
{  
    impl::decrement_cv_counter_and_notify_one(  
        cb._mutex, cb._cv, cb._counter);  
}  
  
/// @brief Decrements `cb`'s atomic counter by one, and calls `notify_all`  
/// on the inner condition variable.  
void decrement_cv_counter_and_notify_all(counter_blocker & cb) noexcept  
{  
    impl::decrement_cv_counter_and_notify_all(  
        cb._mutex, cb._cv, cb._counter);  
}  
  
/// @brief Executes `f` and waits until the blocker's counter is zero. Uses  
/// the blocker's synchronization primitives.  
template <typename TF>  
void execute_and_wait_until_counter_zero(  
    counter_blocker & cb, TF && f) noexcept  
{  
    impl::execute_and_wait_until_counter_zero(  
        cb._mutex, cb._cv, cb._counter, FWD(f));  
}
```

System scheduling

- Current scheduling policy: “**atomic counter**”.
 - Explored alternative: `when_all(...).then(...)` chain generation.
- The scheduler uses a **remaining systems** `counter_blocker` which keeps track of how many systems have been executed.
 - The caller thread blocks until all systems’ executions have been completed.
- Every system instance uses a **remaining dependencies** atomic counter.
 - In the beginning, all instances with **zero** dependencies are executed.
 - As soon as an instance is done, all dependents’ counters are decremented.
 - If a dependency counter reaches **zero**, that instance is executed.

System scheduling - scheduler

```
template <typename TContext, typename TF>
void atomic_counter::execute(TContext& ctx, TF&& f)
{
    reset();

    // Aggregates the required synchronization objects.
    counter_blocker b{mp::list::size(ssl)};

    // Starts every independent task and waits until the remaining tasks
    // counter reaches zero. We forward `f` into the Lambda here, then
    // refer to it everywhere else.
    execute_and_wait_until_counter_zero(b,
        [ this, &ctx, &b, f = FWD(f) ]
    {
        this->start_execution(ctx, b, f);
    });
}
```

System scheduling - scheduler

```
template <typename TContext, typename TF>
void atomic_counter::execute(TContext& ctx, TF&& f)
{
    reset();

    // Aggregates the required synchronization objects.
    counter_blocker b{mp::list::size(ssl)}; number of systems

    // Starts every independent task and waits until the remaining tasks
    // counter reaches zero. We forward `f` into the Lambda here, then
    // refer to it everywhere else.
    execute_and_wait_until_counter_zero(b,
        [ this, &ctx, &b, f = FWD(f) ]
    {
        this->start_execution(ctx, b, f);
    });
}
```

System scheduling - scheduler

```
template <typename TContext, typename TF>
void atomic_counter::execute(TContext& ctx, TF&& f)
{
    reset();

    // Aggregates the required synchronization objects.
    counter_blocker b{mp::list::size(ssl)}; number of systems

    // Starts every independent task and waits until the remaining tasks
    // counter reaches zero. We forward `f` into the Lambda here, then
    // refer to it everywhere else.
    execute_and_wait_until_counter_zero(b,
        [ this, &ctx, &b, f = FWD(f) ]
    {
        this->start_execution(ctx, b, f);
    });
}
```



System scheduler – system task execution

```
void task<TDependencyData>::run(
    TTaskGroup & tg, TCounterBlocker & b, TID my_id, TContext & ctx, TF && f)
{
    // Execute function on instance and decrement remaining subtasks counter.
    _instance.execute(ctx, f);
    decrement_cv_counter_and_notify_one(b);

    // For every dependent task ID...
    for_dependents([this, &tg, &b, &ctx, &f](auto& dt)
    {
        if(!dt.dependency_data().decrement_and_check())
            return;

        // Recursively run the dependent instance.
        ctx.post_in_thread_pool([this, &dt, &tg, &b, &ctx, id, &f]
        {
            dt.run(tg, b, id, ctx, f);
        });
    });
}
```

System scheduler – system task execution

```
void task<TDependencyData>::run(
    TTaskGroup & tg, TCounterBlocker & b, TID my_id, TContext & ctx, TF && f)
{
    // Execute function on instance and decrement remaining subtasks counter.
    _instance.execute(ctx, f);
    decrement_cv_counter_and_notify_one(b);

    // For every dependent task ID...
    for_dependents([this, &tg, &b, &ctx, &f](auto& dt)
    {
        if(!dt.dependency_data().decrement_and_check())
            return;

        // Recursively run the dependent instance.
        ctx.post_in_thread_pool([this, &dt, &tg, &b, &ctx, id, &f]
        {
            dt.run(tg, b, id, ctx, f);
        });
    });
}
```

prevent execution
if dependencies
are left

System scheduler – system task execution

```
void task<TDependencyData>::run(
    TTaskGroup & tg, TCounterBlocker & b, TID my_id, TContext & ctx, TF && f)
{
    // Execute function on instance and decrement remaining subtasks counter.
    _instance.execute(ctx, f);
    decrement_cv_counter_and_notify_one(b);

    // For every dependent
    for_dependents([this,
    {
        if(!dt.dependent)
            return;
    }]);

    context.step([&](auto& proxy)
    {
        proxy.system(st::render).prepare();

        proxy.execute_systems_overload(
            [dt](s::physics& s, auto& data){ s.process(dt, data); },
            [](){ s::render& s, auto& data){ s.process(data); });
    });

    // Recursively run the dependent instance.
    ctx.post_in_thread_pool([this, &dt, &tg, &b, &ctx, id, &f]
    {
        dt.run(tg, b, id, ctx, f);
    });
});
```

... next execution
dependencies
left

System inner parallelism

- Every running system instance can be split in **subtasks**.
 - The instance has a counter of **running subtasks**.
 - As soon as a subtask is finished, the counter is decremented.
 - A **counter blocker** is used to wait for the counter to become **zero**.
- Splitting strategies can be **composed** at compile-time and extended by the users.
 - Strategy executors can be stateful.
 - Motivating example: self-profiling executor that chooses between single-threaded and multi-threaded system execution during application run-time.

System inner parallelism – parallel execution

```
template <typename TSettings, typename TSystemSignature>
template <typename TContext, typename TF>
void instance<TSettings, TSystemSignature>::execute_in_parallel(
    TContext & ctx, TF && f
)
{
    // Aggregates synchronization primitives.
    _parallel_executor.execute(*this, ctx, f);
}
```

- System inner parallel execution is done through a **parallel executor**, which is stored inside the system instance and encapsulates an execution strategy (*or a composition of strategies*).
- Synchronization is handled inside the executor.

System inner parallelism – parallel execution

```
template <typename TSettings, typename TSystemSignature>
template <typename TContext, typename TF>
void instance<TSettings, TSystemSignature>::execute_in_parallel(
    TContext & ctx, TF && f
)
{
    // Aggregates synchronization primitives.
    _parallel_executor.execute(*this, ctx, f);
}
```

- System inner parallel execution is done through a **parallel executor**, which is stored inside the system instance and encapsulates an execution strategy (*or a composition of strategies*).
- Synchronization is handled inside the executor.

System inner parallelism – executor example

```
template <typename TInstance, typename TContext, typename TF>
void split_evenly::execute(TInstance& inst, TContext& ctx, TF&& f)
{
    constexpr auto split_count = sz_v<parameters::subtask_count>();

    auto per_split = inst.subscribed_count() / split_count;
    inst.prepare_and_wait_n_subtasks(split_count, [&](auto& b)
    {
        auto run_subtask = [&inst, &b, &ctx, &f](
            auto split_idx, auto xi_begin, auto xi_end)
        {
            // Create looping execution function.
            auto bse = inst.make_bound_slice_executor(
                b, ctx, split_idx, xi_begin, xi_end, f);

            inst.run_subtask_in_thread_pool(ctx, std::move(bse));
        };

        // Builds and runs the subtasks.
        utils::execute_split_compile_time(inst.subscribed_count(),
            per_split, split_count, run_subtask);
    });
}
```

System inner parallelism – executor example

```
template <typename TInstance, typename TContext, typename TF>
void split_evenly::execute(TInstance& inst, TContext& ctx, TF&& f)
{
    constexpr auto split_count = sz_v<parameters::subtask_count()>; — number of subtasks

    auto per_split = inst.subscribed_count() / split_count;
    inst.prepare_and_wait_n_subtasks(split_count, [&](auto& b)
    {
        auto run_subtask = [&inst, &b, &ctx, &f](
            auto split_idx, auto xi_begin, auto xi_end)
        {
            // Create looping execution function.
            auto bse = inst.make_bound_slice_executor(
                b, ctx, split_idx, xi_begin, xi_end, f);

            inst.run_subtask_in_thread_pool(ctx, std::move(bse));
        };

        // Builds and runs the subtasks.
        utils::execute_split_compile_time(inst.subscribed_count(),
            per_split, split_count, run_subtask);
    });
}
```

System inner parallelism – executor example

```
template <typename TInstance, typename TContext, typename TF>
void split_evenly::execute(TInstance& inst, TContext& ctx, TF&& f)
{
    constexpr auto split_count = sz_v<parameters::subtask_count>(); — number of subtasks
    auto per_split = inst.subscribed_count() / split_count; — entities per subtask
    inst.prepare_and_wait_n_subtasks(split_count, [&](auto& b)
    {
        auto run_subtask = [&inst, &b, &ctx, &f](
            auto split_idx, auto xi_begin, auto xi_end)
        {
            // Create looping execution function.
            auto bse = inst.make_bound_slice_executor(
                b, ctx, split_idx, xi_begin, xi_end, f);

            inst.run_subtask_in_thread_pool(ctx, std::move(bse));
        };

        // Builds and runs the subtasks.
        utils::execute_split_compile_time(inst.subscribed_count(),
            per_split, split_count, run_subtask);
    });
}
```

System inner parallelism – executor example

```
template <typename TInstance, typename TContext, typename TF>
void split_evenly::execute(TInstance& inst, TContext& ctx, TF&& f)
{
    constexpr auto split_count = sz_v<parameters::subtask_count()>; — number of subtasks
    auto per_split = inst.subscribed_count() / split_count; — entities per subtask
    inst.prepare_and_wait_n_subtasks(split_count, [&](auto& b)
    {
        auto run_sub: template <typename TSettings, typename TSystemSignature>
        auto spl: template <typename TF>
        { void instance<TSettings, TSystemSignature>::prepare_and_wait_n_subtasks(
            sz_t n, TF && f)
            _sm.clear_and_prepare(n);
            counter_blocker b{n};

            // Runs the parallel executor and waits until the remaining subtasks
            // counter is zero.
            execute_and_wait_until_counter_zero(b, [this, &b, &f]
            {
                f(b);
            });
        }
    });
}
```

System inner parallelism – executor example

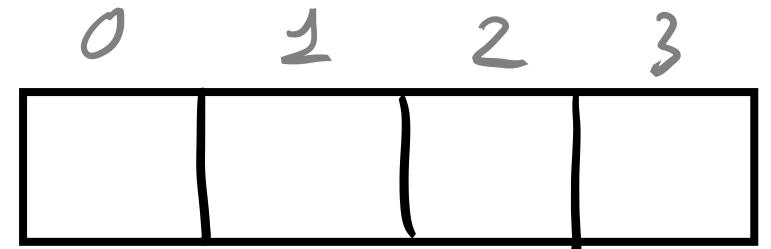
```
template <typename TInstance, typename TContext, typename TF>
void split_evenly::execute(TInstance& inst, TContext& ctx, TF&& f)
{
    constexpr auto split_count = sz_v<parameters::subtask_count>(); — number of subtasks
    auto per_split = inst.subscribed_count() / split_count; — entities per subtask
    inst.prepare_and_wait_n_subtasks(split_count, [&](auto& b)
    {
        auto run_subtask = [&inst, &b, &ctx, &f](
            auto split_idx, auto xi_begin, auto xi_end)
        {
            // Create looping execution function.
            auto bse = inst.make_bound_slice_executor(
                b, ctx, split_idx, xi_begin, xi_end, f);

            inst.run_subtask_in_thread_pool(ctx, std::move(bse));
        };

        // Builds and runs the subtasks.
        utils::execute_split_compile_time(inst.subscribed_count(),
            per_split, split_count, run_subtask);
    });
}
```

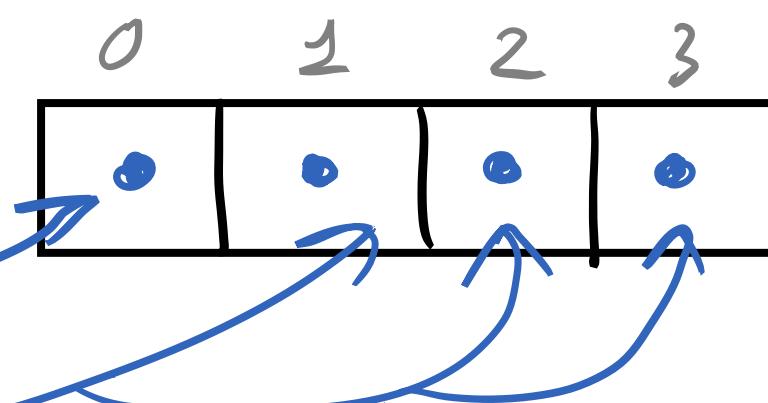
System inner parallelism – executor example

```
template <typename TInstance, typename TContext, typename TF>
void split_evenly::execute(TInstance& inst, TContext& ctx, TF&& f)
{
    constexpr auto split_count = sz_v<parameters::subtask_count(); — number of subtasks
    auto per_split = inst.subscribed_count() / split_count; — entities per subtask
    inst.prepare_and_wait_n_subtasks(split_count, [&](auto& b)
    {
        auto run_subtask = [&inst, &b, &ctx, &f](
            auto split_idx, auto xi_begin, auto xi_end)
        {
            // Create looping execution function.
            auto bse = inst.make_bound_slice_executor(
                b, ctx, split_idx, xi_begin, xi_end, f);
            inst.run_subtask_in_thread_pool(ctx, std::move(bse));
        };
        // Builds and runs the subtasks.
        utils::execute_split_compile_time(inst.subscribed_count(),
            per_split, split_count, run_subtask);
    });
}
```



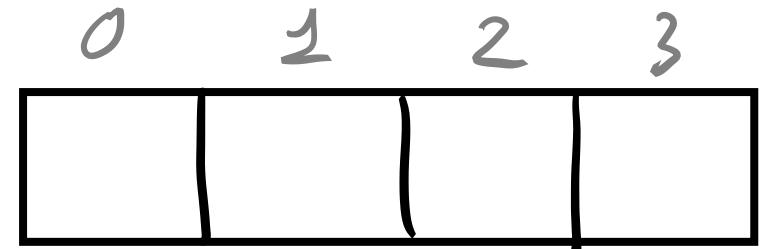
System inner parallelism – executor example

```
template <typename TInstance, typename TContext, typename TF>
void split_evenly::execute(TInstance& inst, TContext& ctx, TF& f)
{
    constexpr auto split_count = sz_v<parameters::subtask_count(); — number of subtasks
    auto per_split = inst.subscribed_count() / split_count; — entities per subtask
    inst.prepare_and_wait_n_subtasks(split_count, [&](auto& b)
    {
        auto run_subtask = [&inst, &b, &ctx, &f](
            auto split_idx, auto xi_begin, auto xi_end)
        {
            // Create looping execution function.
            auto bse = inst.make_bound_slice_executor(
                b, ctx, split_idx, xi_begin, xi_end, f);
            inst.run_subtask_in_thread_pool(ctx, std::move(bse));
        };
        // Builds and runs the subtasks.
        utils::execute_split_compile_time(inst.subscribed_count(),
            per_split, split_count, run_subtask);
    });
}
```



System inner parallelism – executor example

```
template <typename TInstance, typename TContext, typename TF>
void split_evenly::execute(TInstance& inst, TContext& ctx, TF&& f)
{
    constexpr auto split_count = sz_v<parameters::subtask_count(); — number of subtasks
    auto per_split = inst.subscribed_count() / split_count; — entities per subtask
    inst.prepare_and_wait_n_subtasks(split_count, [&](auto& b)
    {
        auto run_subtask = [&inst, &b, &ctx, &f](
            auto split_idx, auto xi_begin, auto xi_end)
        {
            // Create looping execution function.
            auto bse = inst.make_bound_slice_executor(
                b, ctx, split_idx, xi_begin, xi_end, f);
            inst.run_subtask_in_thread_pool(ctx, std::move(bse));
        };
        // Builds and runs the subtasks.
        utils::execute_split_compile_time(inst.subscribed_count(),
            per_split, split_count, run_subtask);
    });
}
```



System inner parallelism – executor example

```
template <typename TInstance, typename TContext, typename TF>
void split_evenly::execute(TInstance& inst, TContext& ctx, TF&& f)
{
    constexpr auto split_count = sz_v<parameters::subtask_count>(); — number of subtasks
    auto per_split = inst.subscribed_count() / split_count; — entities per subtask
    inst.prepare_and_wait_n_subtasks(split_count, [&](auto& b)
    {
        auto run_subtask = [&inst, &b, &ctx, &f](
            auto split_idx, auto xi_begin, auto xi_end)
        {
            // Create looping execution function.
            auto bse = inst.make_bound_slice_executor(
                b, ctx, split_idx, xi_begin, xi_end, f);

            inst.run_subtask_in_thread_pool(ctx, std::move(bse));
        };

        // Builds and runs the subtasks.
        utils::execute_split_compile_time(inst.subscribed_count(),
            per_split, split_count, run_subtask);
    });
}
```

System inner parallelism - slicing

```
template <typename TSettings, typename TSystemSignature>
template <typename TCounterBlocker, typename TContext, typename TF>
auto instance<TSettings, TSystemSignature>::make_bound_slice_executor(
    TCounterBlocker & cb, TContext & ctx, sz_t state_idx, sz_t i_begin,
    sz_t i_end, TF && f) noexcept
{
    return [this, &cb, &ctx, &f, state_idx, i_begin, i_end]
    {
        this->make_slice_executor(cb, ctx, state_idx, i_begin, i_end)(f);
    };
}
```

- The above function binds a range of entities (*subset of subscribed entities*) to an execution callable object.

System inner parallelism - slicing

```
template <typename TSettings, typename TSystemSignature>
template <typename TCounterBlocker, typename TContext, typename TF>
auto instance<TSettings, TSystemSignature>::make_bound_slice_executor(
    TCounterBlocker & cb, TContext & ctx, sz_t state_idx, sz_t i_begin,
    sz_t i_end, TF && f) noexcept
{
    return [this, &cb, &ctx, &f, state_idx, i_begin, i_end]
    {
        this->make_slice_executor(cb, ctx, state_idx, i_begin, i_end)(f);
    };
}
```

- The above function binds a range of entities (*subset of subscribed entities*) to an execution callable object.

System inner parallelism - slicing

```
template <typename TSettings, typename TSystemSignature>
template <typename TCounterBlocker, typename TContext>
auto ECST_PURE_FN instance<TSettings,
    TSystemSignature>::make_slice_executor(TCounterBlocker & cb,
    TContext & ctx, sz_t state_idx, sz_t i_begin, sz_t i_end) noexcept
{
    return [this, &cb, &ctx, state_idx, i_begin, i_end](auto&& f)
    {
        auto data =
            this->make_entity_range_data(ctx, state_idx, i_begin, i_end);

        // Executes the processing function over the slice of entities.
        this->execute_subtask_and_decrement_counter(cb, data, f);
    };
}
```

- The above function creates a **data proxy** and runs the parallel subtask function passing it as one of the parameters.

System inner parallelism - slicing

```
template <typename TSettings, typename TSystemSignature>
template <typename TCounterBlocker, typename TContext>
auto ECST_PURE_FN instance<TSettings,
    TSystemSignature>::make_slice_executor(TCounterBlocker & cb,
    TContext & ctx, sz_t state_idx, sz_t i_begin, sz_t i_end) noexcept
{
    return [this, &cb, &ctx, state_idx, i_begin, i_end](auto&& f)
    {
        auto data =
            this->make_entity_range_data(ctx, state_idx, i_begin, i_end);
// A red oval surrounds the line 'auto data = ...'; above it, the text 'data proxy' is written in red.
data proxy
        // Executes the processing function over the slice of entities.
        this->execute_subtask_and_decrement_counter(cb, data, f);
    };
}
```

- The above function creates a **data proxy** and runs the parallel subtask function passing it as one of the parameters.

System inner parallelism - slicing

```
template <typename TSettings, typename TSystemSignature>
template <typename TCounterBlocker, typename TData, typename TF>
void instance<TSettings,
TSystemSignature>::execute_subtask_and_decrement_counter( // .
TCounterBlocker & cb, TData & data, TF && f           // .
)
{
    f(_system, data);
    decrement_cv_counter_and_notify_all(cb);
}

// Executes the processing function over the slice of entities.
this->execute_subtask_and_decrement_counter(cb, data, f);
};

auto ECST_PURE_FN instance<TSettings, TSystemSignature>
instantiate<TSettings, TSystemSignature>( // .
TCounterBlocker & cb, TData & data, TF && f           // .
)
{
    return [this, &cb] {
        auto data = this->make_data_proxy();
        execute_subtask_and_decrement_counter(cb, data, f);
    };
}
```

- The above function creates a **data proxy** and runs the parallel subtask function passing it as one of the parameters.

System inner parallelism - slicing

```
template <typename TSettings, typename TSystemSignature>
template <typename TCounterBlocker, typename TContext>
auto ECST_PURE_FN instance<TSettings,
    TSystemSignature>::make_slice_executor(TCounterBlocker & cb,
    TContext & ctx, sz_t state_idx, sz_t i_begin, sz_t i_end) noexcept
{
    return [this, &cb, &ctx, state_idx, i_begin, i_end](auto&& f)
    {
        auto data =
            this->make_entity_range_data(ctx, state_idx, i_begin, i_end);
// A red oval surrounds the line 'auto data = ...'; above it, the text 'data proxy' is written in red.
data proxy
        // Executes the processing function over the slice of entities.
        this->execute_subtask_and_decrement_counter(cb, data, f);
    };
}
```

- The above function creates a **data proxy** and runs the parallel subtask function passing it as one of the parameters.

Proxies

- **Proxies** are used to limit the scope in which **critical operations** can be executed.
- **Data proxies** are used during system execution to provide component data access and system output access.
- **Defer proxies** are used in `defer(...)` calls to allow the user to queue up critical operations that will be executed in the **refresh step**. **Step proxies** derive from defer proxies, and allow the execution of system schedulers.

Proxies – data proxy

```
struct acceleration
{
    template <typename TData>
    void process(ft dt, TData& data)
    {
        data.for_entities([&](auto eid)
        {
            auto& v = ecst::get<c::velocity>(data, eid);
            const auto& a = ecst::get<c::acceleration>(data, eid);
            actions::accelerate(dt, v, a);
        });
    }
};
```

Proxies – data proxy

```
struct acceleration
{
    template <typename TData>
    void process(ft dt, TData& data)
    {
        data.for_entities([&](auto eid)
        {
            auto& v = ecst::get<c::velocity>(data, eid);
            const auto& a = ecst::get<c::acceleration>(data, eid);
            actions::accelerate(dt, v, a);
        });
    }
};
```



Proxies – defer proxy

```
data.for_entities([&](auto eid)
{
    const auto& contacts =
        data.get_previous_output<s::collision_detection>();

    if(contacts.was_hit(eid))
    {
        auto& h = ecst::get<c::health>(data, eid);
        h.damage();

        if(h.dead()) data.kill_entity(eid);

        data.defer([&](auto& proxy)
        {
            auto e = proxy.create_entity();
            auto& ep = e.add_component(ct::effect_particle);
            ep.set_effect(effects::explosion);
        });
    }
});
```

Proxies – defer proxy

```
data.for_entities([&](auto eid)
{
    const auto& contacts =
        data.get_previous_output<s::collision_detection>();

    if(contacts.was_hit(eid))
    {
        auto& h = ecst::get<c::health>(data, eid);
        h.damage();

        if(h.dead()) data.kill_entity(eid);

        data.defer([&](auto& proxy)
        {
            auto e = proxy.create_entity();
            auto& ep = e.add_component(ct::effect_particle);
            ep.set_effect(effects::explosion);
        });
    }
});
```



Proxies – defer proxy

```
data.for_entities([&](auto eid)
{
    const auto& contacts =
        data.get_previous_output<s::collision_detection>();

    if(contacts.was_hit(eid))
    {
        auto& h = ecst::get<c::health>(data, eid);
        h.damage();

        if(h.dead()) data.kill_entity(eid);

        data.defer([&](auto& proxy)
        {
            auto e = proxy.create_entity();
            auto& ep = e.add_component(ct::effect_particle);
            ep.set_effect(effects::explosion);
        });
    }
});
```



Proxies – step proxy

```
context.step([&](auto& proxy)
{
    proxy.system(st::render).prepare();

    proxy.execute_systems_overload(
        [dt](s::physics& s, auto& data){ s.process(dt, data); },
        [](s::render& s, auto& data){ s.process(data); });

    proxy.for_system_outputs(st::render,
        [&window](auto& s, auto& va)
    {
        window.draw(va.data(), va.size(),
                    PrimitiveType::Triangles, RenderStates::Default);
    });
});
```



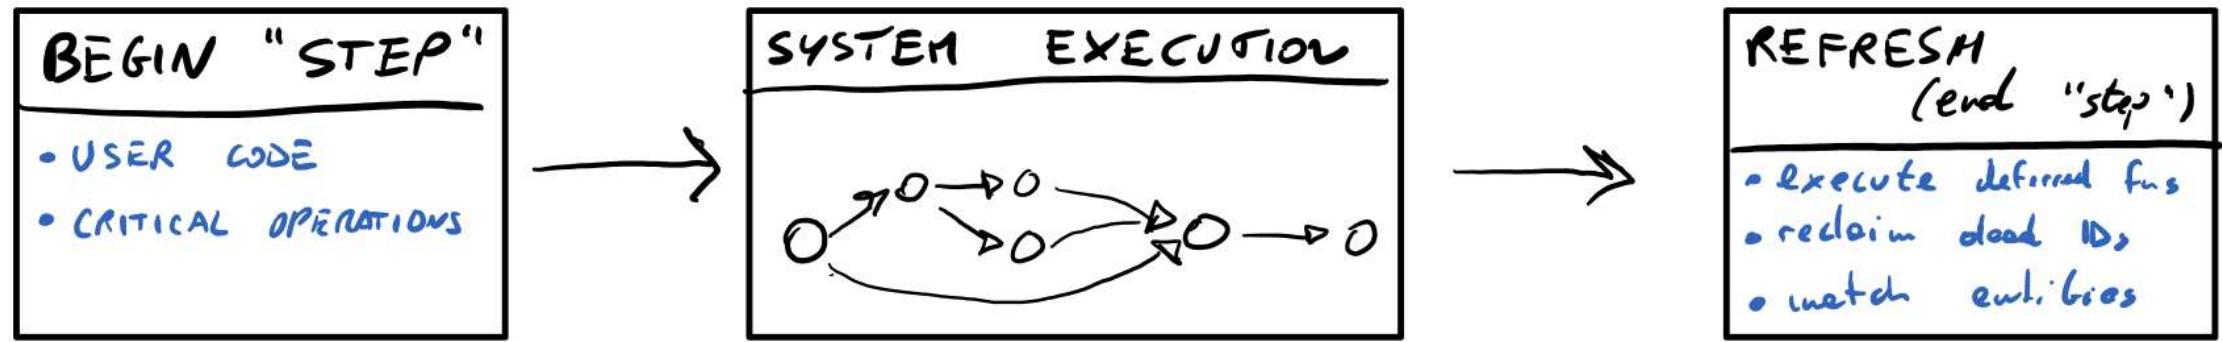
Proxies – step proxy

```
context.step([&](auto& proxy)
{
    proxy.system(st::render).prepare();

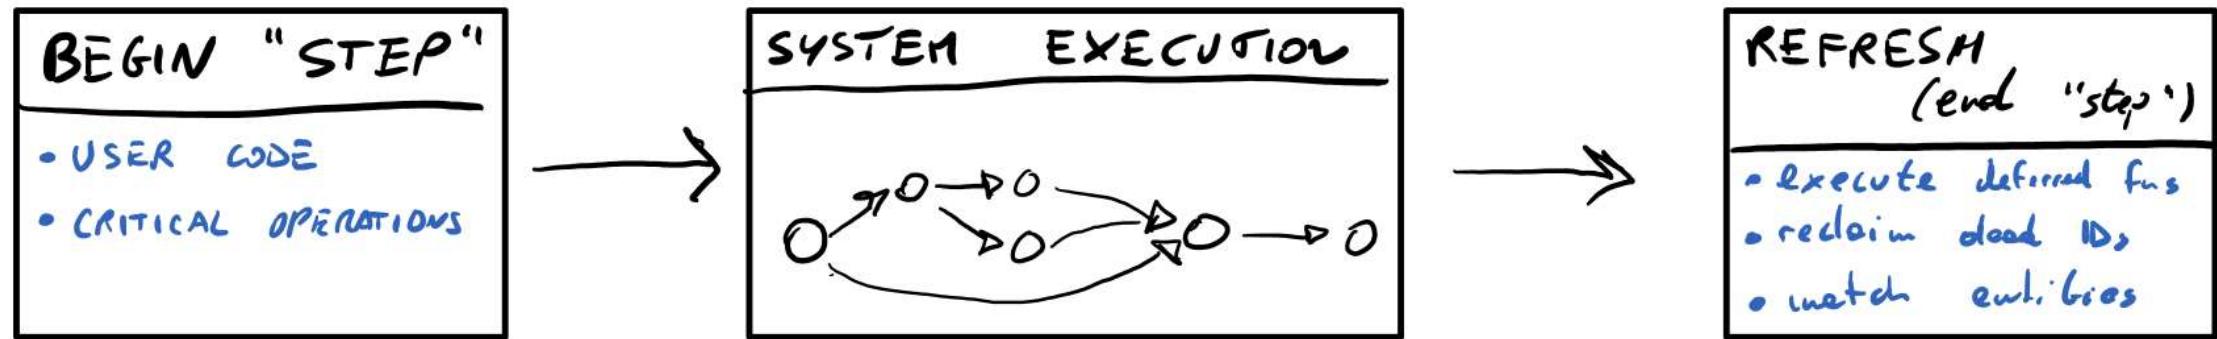
    proxy.execute_systems_overload(
        [dt](s::physics& s, auto& data){ s.process(dt, data); },
        [](s::render& s, auto& data){ s.process(data); });

    proxy.for_system_outputs(st::render,
        [&window](auto& s, auto& va)
    {
        window.draw(va.data(), va.size(),
                    PrimitiveType::Triangles, RenderStates::Default);
    });
});
```

Flow – implementation overview (2)

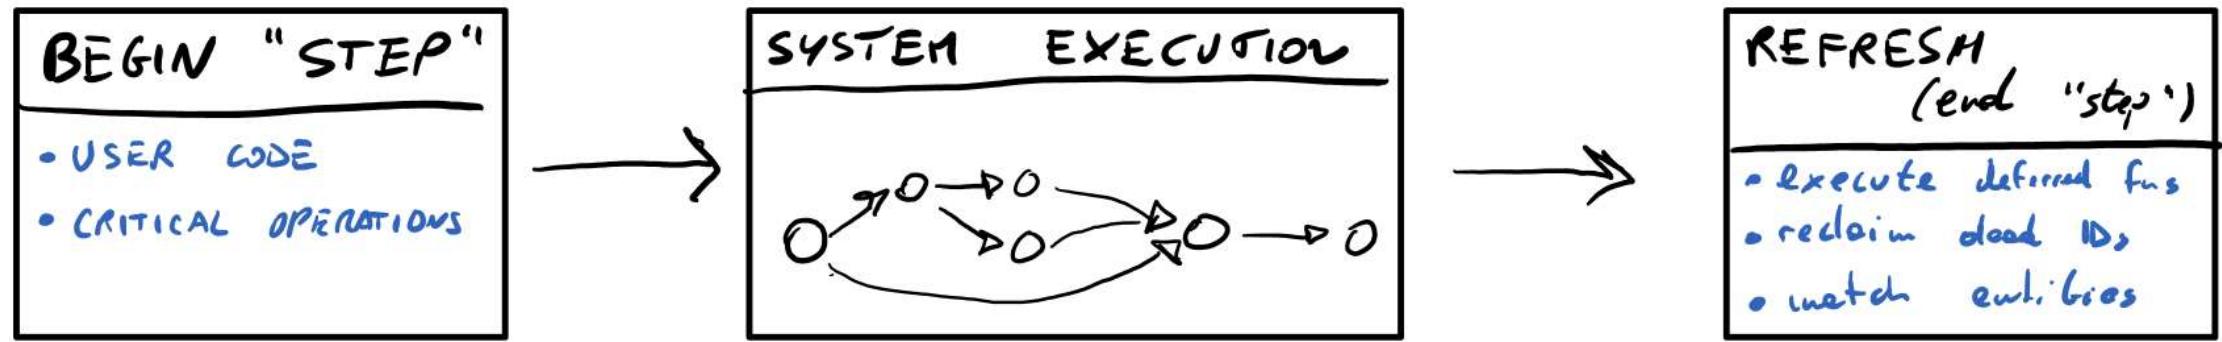


Flow – implementation overview (2)



```
_ctx.step([this, dt](auto& proxy)
{
    /* ... */
    proxy.execute_systems(/* ... */);
    /* ... */
};
```

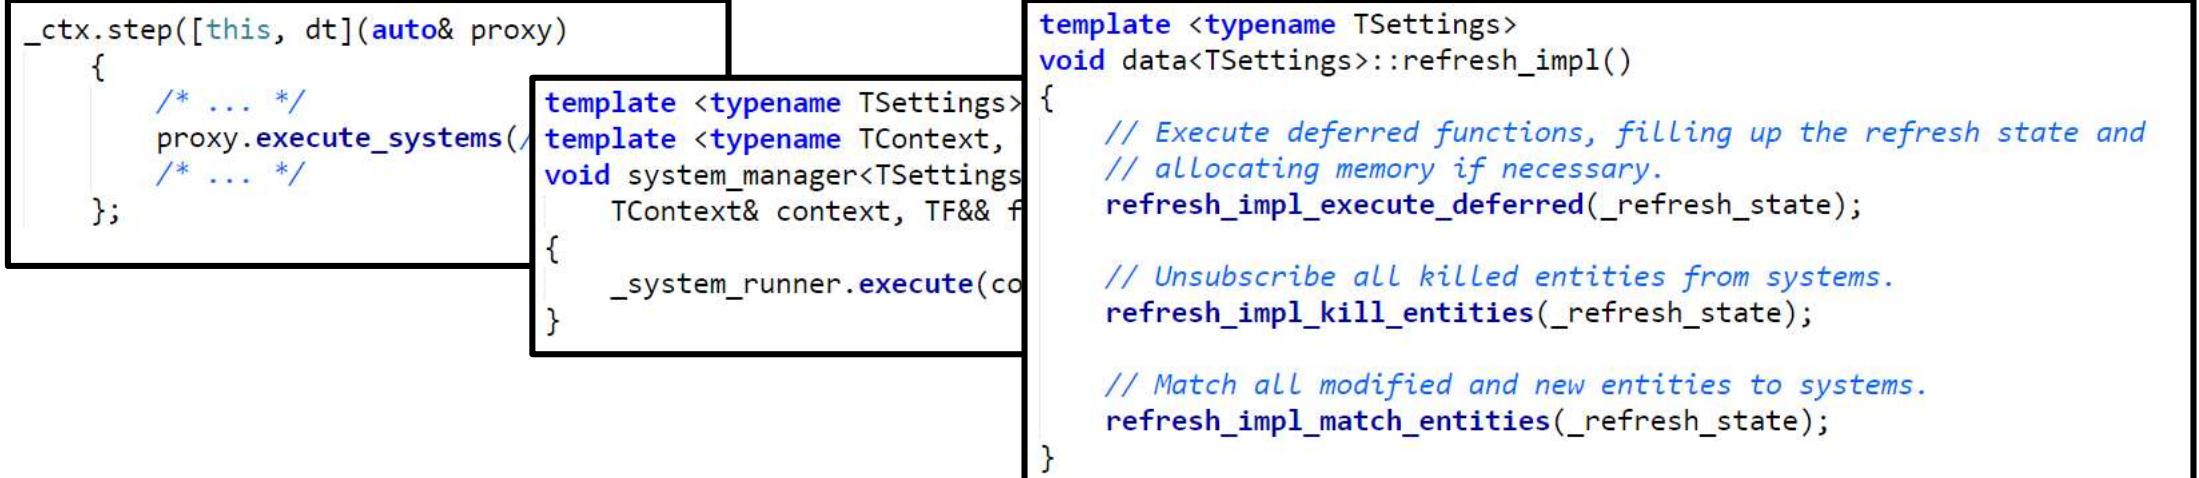
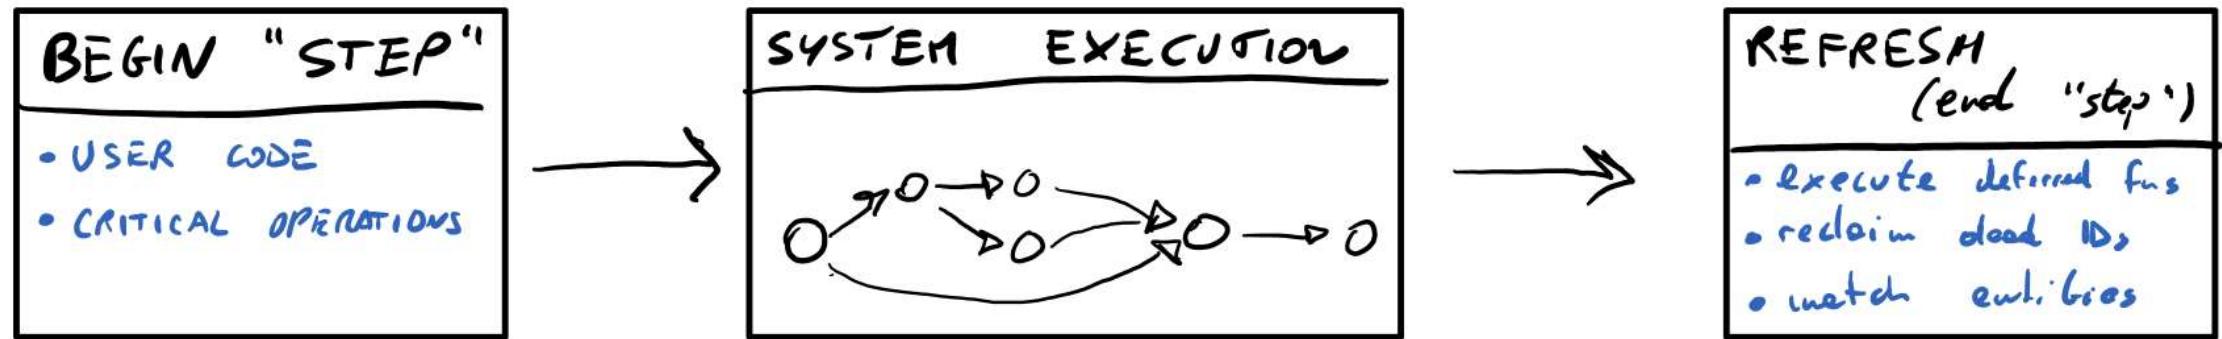
Flow – implementation overview (2)



```
_ctx.step([this, dt](auto& proxy)
{
    /* ... */
    proxy.execute_systems(
    /* ... */
});
```

```
template <typename TSettings>
template <typename TContext, typename TF>
void system_manager<TSettings>::execute_systems(
    TContext& context, TF&& f)
{
    _system_runner.execute(context, FWD(f));
}
```

Flow – implementation overview (2)



Flow – “begin” step

```
template <typename TSettings>
template <typename TF>
auto context::data<TSettings>::step(TF&& f)
{
    // Clear refresh state.
    _refresh_state.clear();

    // Build context step proxy.
    step_proxy_type step_proxy{*this, _refresh_state};

    // Execute user-defined step.
    f(step_proxy);

    // Refresh context.
    refresh();
}
```



Flow – “begin” step

```
template <typename TSettings>
template <typename TF>
auto context::data<TSettings>::step(TF&& f)
{
    // Clear refresh state.
    _refresh_state.clear();

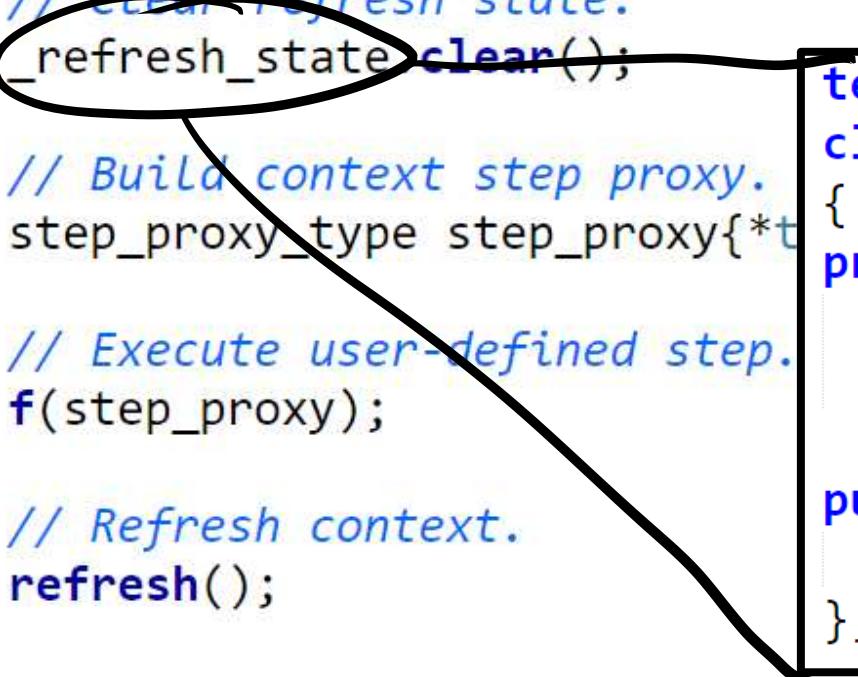
    // Build context step proxy.
    step_proxy_type step_proxy{*this};

    // Execute user-defined step.
    f(step_proxy);

    // Refresh context.
    refresh();
}
```

```
template <typename TSettings>
class refresh_state
{
private:
    using set_type = dispatch_set<TSettings>;
    set_type _to_match_ids, _to_kill_ids;

public:
    // ...
};
```



Flow – “begin” step

```
template <typename TSettings>
template <typename TF>
auto context::data<TSettings>::step(TF&& f)
{
    // Clear refresh state.
    _refresh_state.clear();

    // Build context step proxy.
    step_proxy_type step_proxy{*this, _refresh_state};

    // Execute user-defined step.
    f(step_proxy);

    // Refresh context.
    refresh();
}
```



Flow – “begin” step

```
template <typename TSettings>
template <typename TF>
auto context::data<TSettings>::step(TF&& f)
{
    // Clear refresh state.
    _refresh_state.clear();

    // Build context step proxy.
    step_proxy_type step_proxy{*this, _refresh_state};

    // Execute user-defined step.
    f(step_proxy);

    // Refresh context.
    refresh();
}
```



Flow – refresh step

```
template <typename TSettings>
void context::data<TSettings>::refresh()
{
    // Execute deferred functions, filling up the refresh state and
    // allocating memory if necessary.
    refresh_impl_execute_deferred(_refresh_state);

    // Unsubscribe all killed entities from systems.
    refresh_impl_kill_entities(_refresh_state);

    // Match all modified and new entities to systems.
    refresh_impl_match_entities(_refresh_state);
}
```

Flow – refresh step – execute deferred functions

```
template <typename TSettings>
template <typename TRefreshState>
void context::data<TSettings>::refresh_impl_execute_deferred(TRefreshState& rs)
{
    defer_proxy_type defer_proxy{*this, rs};

    for_systems_sequential([&defer_proxy](auto& system)
    {
        system.for_states([&defer_proxy](auto& state)
        {
            // The execution of deferred functions fills the
            // refresh state and alters the context state.
            state.execute_deferred_fns(defer_proxy);
        });
    });
}
```



Flow – refresh step – execute deferred functions

```
template <typename TSettings>
template <typename TRefreshState>
void context::data<TSettings>::refresh_impl_execute_deferred(TRefreshState& rs)
{
    defer_proxy_type defer_proxy{*this, rs};

    for_systems_sequential([&defer_proxy](auto& system)
    {
        system.for_states([&defer_proxy](auto& state)
        {
            // The execution of deferred functions fills the
            // refresh state and alters the context state.
            state.execute_deferred_fns(defer_proxy);
        });
    });
}
```



Flow – refresh step – execute deferred functions

```
template <typename TSettings>
template <typename TRefreshState>
void context::data<TSettings>::refresh_impl_execute_deferred(TRefreshState& rs)
{
    defer_proxy_type defer_proxy{*this, rs};

    for_systems_sequential([&defer_proxy](auto& system)
    {
        system.for_states([&defer_proxy](auto& state)
        {
            // The execution of deferred functions fills the
            // refresh state and alters the context state.
            state.execute_deferred_fns(defer_proxy);
        });
    });
}
```



Flow – refresh step – reclaim dead entities

```
template <typename TSettings>
template <typename TRefreshState>
void context::data<TSettings>::refresh_impl_kill_entities(TRefreshState& rs)
{
    for_systems_sequential([&rs](auto& system)
    {
        system.for_states([&rs](auto& state)
        {
            state.for_to_kill([&](auto eid){ rs.add_to_kill(eid); });
        });
    });

    for_systems_parallel ([this, &rs](auto& system)
    {
        rs.for_to_kill([&system](auto eid){ system.unsubscribe(eid); });
    });

    rs.for_to_kill([this](auto eid){ this->reclaim(eid); });
}
```

Flow – refresh step – reclaim dead entities

```
template <typename TSettings>
template <typename TRefreshState>
void context::data<TSettings>::refresh_impl_kill_entities(TRefreshState& rs)
{
    for_systems_sequential([&rs](auto& system)
    {
        system.for_states([&rs](auto& state)
        {
            state.for_to_kill([&](auto eid){ rs.add_to_kill(eid), });
        });
    });

    for_systems_parallel ([this, &rs](auto& system)
    {
        rs.for_to_kill([&system](auto eid){ system.unsubscribe(eid); });
    });

    rs.for_to_kill([this](auto eid){ this->reclaim(eid); });
}
```

kill RS

Flow – refresh step – reclaim dead entities

```
template <typename TSettings>
template <typename TRefreshState>
void context::data<TSettings>::refresh_impl_kill_entities(TRefreshState& rs)
{
    for_systems_sequential([&rs](auto& system)
    {
        system.for_states([&rs](auto& state)
        {
            state.for_to_kill([&](auto eid){ rs.add_to_kill(eid), });
        });
    });

    for_systems_parallel ([this, &rs](auto& system)
    {
        rs.for_to_kill([&system](auto eid){ system.unsubscribe(eid); });
    });

    rs.for_to_kill([this](auto eid){ this->reclaim(eid); });
}
```

kill RS

Flow – refresh step – reclaim dead entities

```
template <typename TSettings>
template <typename TRefreshState>
void context::data<TSettings>::refresh_impl_kill_entities(TRefreshState& rs)
{
    for_systems_sequential([&rs](auto& system)
    {
        system.for_states([&rs](auto& state)
        {
            state.for_to_kill([&](auto eid){ rs.add_to_kill(eid), });
        });
    });

    for_systems_parallel ([this, &rs](auto& system)
    {
        rs.for_to_kill([&system](auto eid){ system.unsubscribe(eid); });
    });

    rs.for_to_kill([this](auto eid){ this->reclaim(eid); });
}
```

kill RS

Flow – refresh step – match entities to systems

```
template <typename TSettings>
template <typename TRefreshState>
void context::data<TSettings>::refresh_impl_match_entities(TRefreshState& rs)
{
    for_systems_parallel ([this, &rs](auto& system)
    {
        rs.for_to_match([this, &system](auto eid)
        {
            // Get entity metadata.
            auto& em(this->metadata(eid));

            // Check if the bitset matches the system.
            if(system.matches_bitset(em.bitset()))
                system.subscribe(eid);
            else
                system.unsubscribe(eid);
        });
    });
}
```

Flow – refresh step – match entities to systems

```
template <typename TSettings>
template <typename TRefreshState>
void context::data<TSettings>::refresh_impl_match_entities(TRefreshState& rs)
{
    for_systems_parallel ([this, &rs](auto& system)
    {
        rs.for_to_match([this, &system](auto eid)
        {
            // Get entity metadata.
            auto& em(this->metadata(eid));

            // Check if the bitset matches the system.
            if(system.matches_bitset(em.bitset()))
                system.subscribe(eid);
            else
                system.unsubscribe(eid);
        });
    });
}
```

Filled by creating entities or adding/removing components

Flow – refresh step – match entities to systems

```
template <typename TSettings>
template <typename TRefreshState>
void context::data<TSettings>::refresh_impl_match_entities(TRefreshState& rs)
{
    for_systems_parallel ([this, &rs](auto& system)
    {
        rs.for_to_match([this, &system](auto eid)
        {
            // Get entity metadata.
            auto& em(this->metadata(eid));

            // Check if the bitset matches the system.
            if(system.matches_bitset(em.bitset()))
                system.subscribe(eid);
            else
                system.unsubscribe(eid);
        });
    });
}
```

Filled by creating entities or adding/removing components

$(s_{_bs} \& e_{_bs}) == s_{_bs}$

Metaprogramming

- “Type-value encoding” oriented data structures:
 - List, pair, map, graph.
- Heavy usage of **variable templates**, **constexpr** and **generic lambdas**.
- Control flow implemented using generic lambdas and “**static if**”.
- Loops/recursion implemented using **Y-combinators**.
- Namespace-oriented usage syntax.

Metaprogramming – list – overview

- Implemented with `std::tuple`.
- Heavily used throughout **ECST**.
- Supports **left folding, slicing, filtering**, etc...

```
constexpr auto l0 = mp::list::make(iv<0>, iv<10>);  
static_assert(mp::list::size(l0) == 2);
```

```
constexpr auto l1 = mp::list::append(l0, iv<20>);  
static_assert(mp::list::size(l1) == 3);
```

```
constexpr auto l0 = mp::list::make(  
    iv<0>, iv<1>, iv<2>, iv<3>, iv<4>);  
  
auto l1 = mp::list::filter(l0, [](auto x)  
{  
    return bool_v<(x % 2 == 0)>;  
});  
  
constexpr auto l1_c = decltype(l1){};
```

Metaprogramming – list – implementation

```
namespace list
{
    template <typename... Ts>
    constexpr std::tuple<Ts...> v{};

    template <typename... Ts>
    constexpr auto make(Ts...)
    {
        return v<Ts...>;
    }
}

auto l0 = list::v<int, float, double>;
auto l1 = list::make(sz_v<0>, sz_v<10>, sz_v<20>);
```

Metaprogramming – list – implementation

```
namespace list
{
    template <typename... Ts>
    constexpr std::tuple<Ts...> v{};

    template <typename... Ts>
    constexpr auto make(Ts...)
    {
        return v<Ts...>;
    }
}

auto l0 = list::v<int, float, double>;
auto l1 = list::make(sz_v<0>, sz_v<10>, sz_v<20>);
```

Metaprogramming – list – use cases

```
namespace signature_list::component
{
    template <typename... TComponentSignatures>
    constexpr auto make(TComponentSignatures... css)
    {
        static_assert(signature::component::are_signatures(css...));
        return mp::list::make(css...);
    }

    template <typename... TComponents>
    constexpr auto v = make(signature::component::make<TComponents>()...);
}

auto cl = signature_list::component::v<c::position, c::velocity>;
```

- Component and system **signature lists** that context settings require are implemented as lists.

Metaprogramming – forcing constexpr

```
namespace impl
{
    template <typename TList, typename TComparer>
    auto find_if_impl(TList l, TComparer c)
    {
        auto res = find_first_index_of_matching(l, c);
        return static_if(is_null(res))
            .then([](auto){ return null_v; })
            .else_([=](auto y){ return at(l, y); })(res);
    }
}

template <typename TList, typename TComparer>
constexpr auto find_if(TList l, TComparer c)
{
    return decltype(impl::find_if_impl(l, c))();
}
```

Metaprogramming – forcing constexpr

```
namespace impl
{
    non-constexpr
    template <typename TList, typename TComparer>
    auto find_if_impl(TList l, TComparer c)
    {
        auto res = find_first_index_of_matching(l, c);
        return static_if(is_null(res))
            .then([](auto){ return null_v; })
            .else_([=](auto y){ return at(l, y); })(res);
    }
}

template <typename TList, typename TComparer>
constexpr auto find_if(TList l, TComparer c)
{
    return decltype(impl::find_if_impl(l, c))();
}
```

Metaprogramming – forcing constexpr

```
namespace impl
{
    non-constexpr
    template <typename TList, typename TComparer>
    auto find_if_impl(TList l, TComparer c)
    {
        auto res = find_first_index_of_matching(l, c);
        return static_if(is_null(res))
            .then([](auto){ return null_v; })
            .else_([=](auto y){ return at(l, y); })(res);
    }
}

template <typename TList, typename TComparer>
constexpr auto find_if(TList l, TComparer c)
{
    return decltype(impl::find_if_impl(l, c))();
}
```

Metaprogramming – forcing constexpr

```
namespace impl
{
    non-constexpr
    template <typename TList, typename TComparer>
    auto find_if_impl(TList l, TComparer c)
    {
        auto res = find_first_index_of_matching(l, c);
        return static_if(is_null(res))
            .then([](auto){ return null_v; })
            .else_([=](auto y){ return at(l, y); })(res);
    }
}

template <typename TList, typename TComparer>
constexpr auto find_if(TList l, TComparer c)
{
    return decltype(impl::find_if_impl(l, c))();
}
```

Conclusion

Resources and future ideas.



<http://vittorioromeo.info>
vittorio.romeo@outlook.com

<http://github.com/SuperV1234/cppnow2016>

Resources – (1)

- <http://t-machine.org>
 - Articles on **data structures, multithreading and networking.**
 - Wiki with **ES approaches** and existing implementations.
- <http://stackoverflow.com/questions/1901251>
 - In-depth analysis of **component-based engine design.**
- <http://bitsquid.blogspot.it>
 - Articles on **contiguous component data** allocation strategies.
- <http://gameprogrammingpatterns.com/component>
 - Covers **component-based design** and **entity communication** techniques.
- <http://randygaul.net>
 - Articles on **component-based design**, covering **communication** and **allocation.**

Resources – (2)

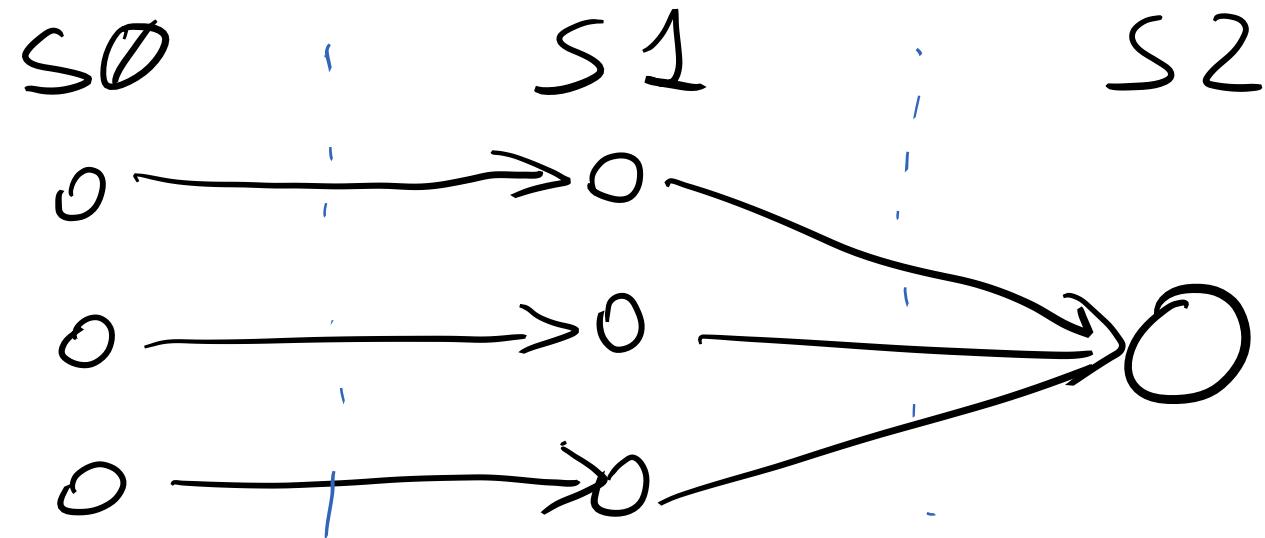
- Existing C++11/14 ECS libraries:
 - <https://github.com/grandstack/Pronto> (*uses compile-time entities*).
 - <https://github.com/discoLoda/Diana> (*uses sparse sets heavily*).
 - <https://github.com/alecthomas/entityx> (*uses compile-time components and events*).
- <https://maikklein.github.io/post/2016-01-14-Entity-Component-System/>
 - «The general design of my flawed compile-time entity component system in C++14.»
- <https://www.reddit.com/r/gamedev/comments/3nv8uz>
 - Discussion on my previous CppCon 2015 ECS talk.

Resources – (3)

- <http://stackoverflow.com/questions/35778864/building-asynchronous-future-callback-chain-from-compile-time-dependency-graph>
 - Building asynchronous `future` callback chain from compile-time dependency graph (DAG)
 - StackOverflow question I asked regarding the generation of a callback chain for system execution.
- <http://cs.stackexchange.com/questions/2524/getting-parallel-items-in-dependency-resolution>
 - Executing tasks with dependencies in parallel.

Future ideas – (1)

- How to deal with SIMD operations?
 - Possibly mark certain components/systems as SIMD and use special storage strategies.
- Match system subtasks to next system subtasks to use partial outputs without merging.



Future ideas – (2)

- Find an alternative to deferred functions.
 - Avoid std::function overhead.
 - Proxy object + command queues?
 - Move-only SBO functions?
- Generic “settings/signature” library.
 - Define “settings” schemas at compile-time.
 - Focus on constraints and composability.
 - Statically verify settings with nice error messages.
- Allow the refresh step to be customized.
 - Subscribe/unsubscribe callbacks.
 - Matching/killing callbacks.

Future ideas – (3)

- Allow multiple separate system execution steps.
 - Allow systems to easily have different “frequencies”.

```
// Enable access to critical operations and system execution.
context.step([](auto& proxy)
{
    // Most general version.
    proxy.execute_systems_overload_detailed(ss::start_from(st::physics),
        [](s::physics& s, auto& executor)
    {
        // Only once, even with multiple subtasks.
        s.prepare();

        executor.for_subtasks([&s](auto& subtask, auto& data)
        {
            s.process(data);
            auto out = subtask.output();
            do_something(out);
        });
    });
});
```

Questions?

<http://vittorioromeo.info>

vittorio.romeo@outlook.com

<http://github.com/SuperV1234>

Thank you for attending!